

# Feature Matching and Robust Fit

IACV  
December 10th 2021  
Giacomo Boracchi

<https://boracchi.faculty.polimi.it/>

Slide credits Luca Magri

# Matching of Computer Vision Features

## Estimating Image Correspondences

👉 Extract features from each image

- Keypoint detection
- Descriptor Computation
- Match features between images
- Prune matches and then perform triangulation  
detect objects / stitching ...



# Matching of Computer Vision Features

## Estimating Image Correspondences

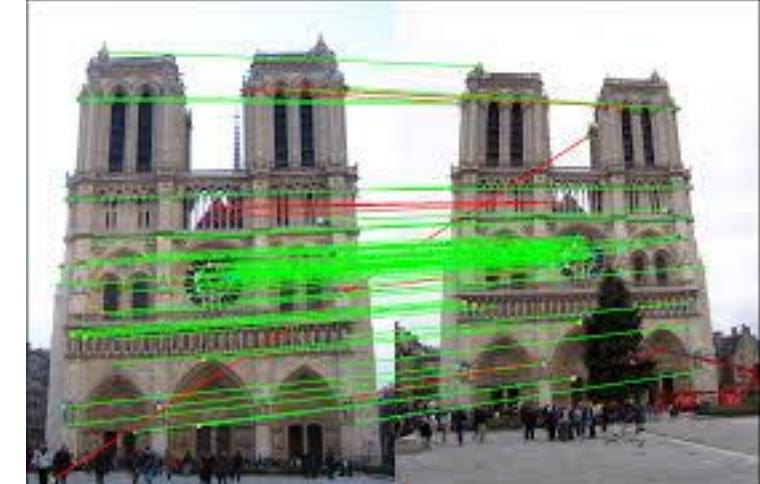
- Extract features from each image

- Keypoint detection
- Descriptor Computation

👉 Match features between images

- Prune matches and then perform triangulation  
detect objects / stitching ...

**Matching:**  
*identify in a pair of images  
the image points that correspond  
to the same 3D point in the scene*



# Matching of Computer Vision Features

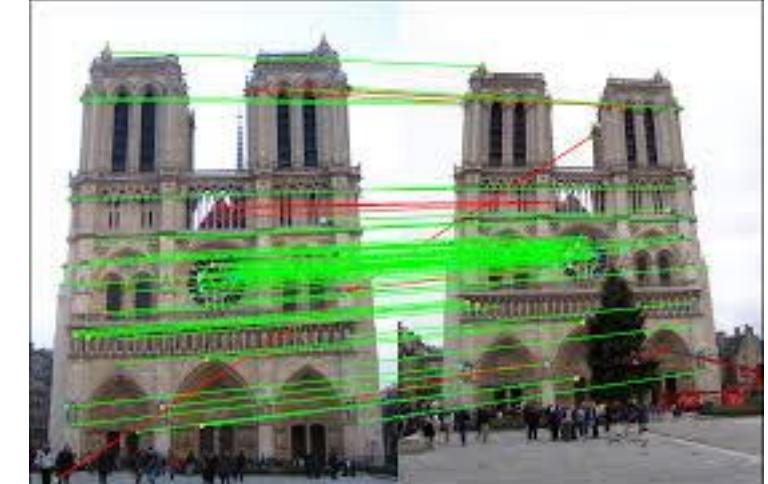
## Estimating Image Correspondences

- Extract features from each image
    - Keypoint detection
    - Descriptor Computation
  - Match features between images
- 👉 Prune matches and then perform triangulation  
detect objects / stitching. Geometry into play!

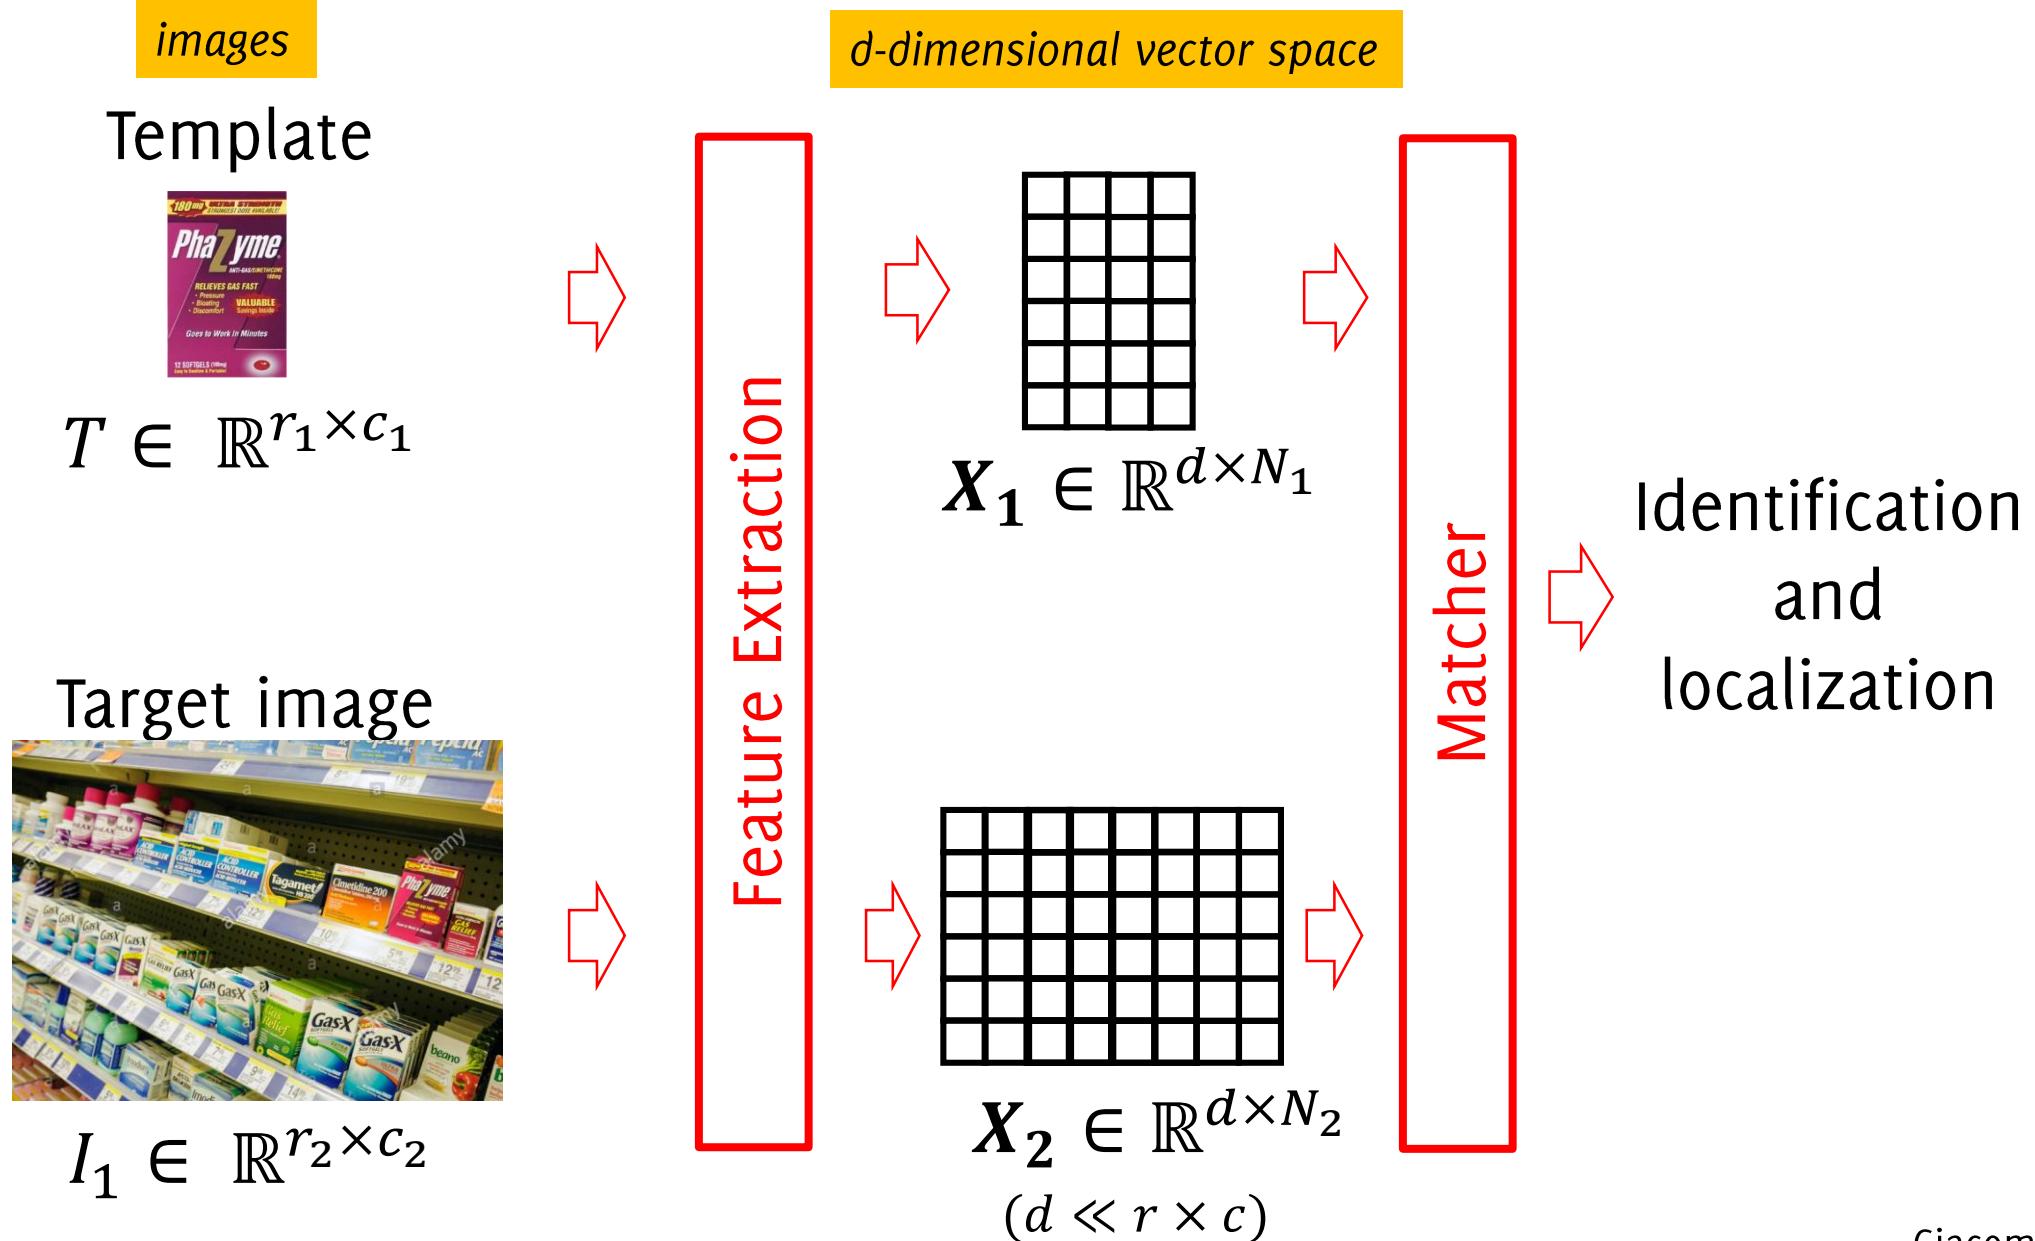
**Robust fit:**

*fit a geometric model to noisy data  
corrupted by outliers*

**Matching:**  
*identify in a pair of images  
the image points that correspond  
to the same 3D point in the scene*



# An application: template matching



# Feature Matching

# The General Approach to Feature Matching

When comparing two images  $I_1$  and  $I_2$  one typically:

- Extract SIFT features (keypoint + descriptor) from independently from  $I_1$  and  $I_2$
- For each descriptor in  $I_1$  we look for the most similar, descriptors in  $I_2$ , i.e., we compute its nearest neighborhood in  $I_2$  in terms of the Euclidean distance
- Matches are confirmed when the distance between the two descriptors is below a threshold. A matched feature connects both keypoint location and scale.
- Match clustering / Outlier removal: Geometric criteria to discard wrong matches can be implemented.

# Feature Matching

Input:  $X_1$  and  $X_2$ , features extracted from  $I_1$  and  $I_2$

Template  
features

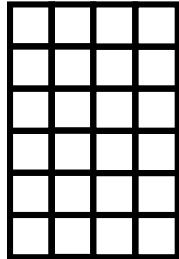
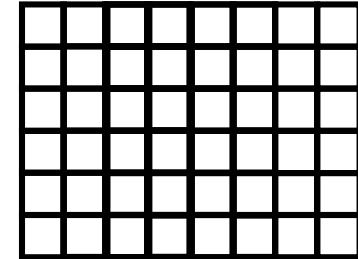
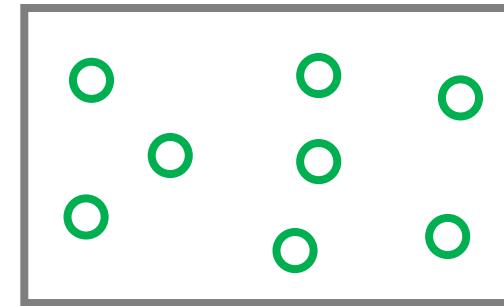
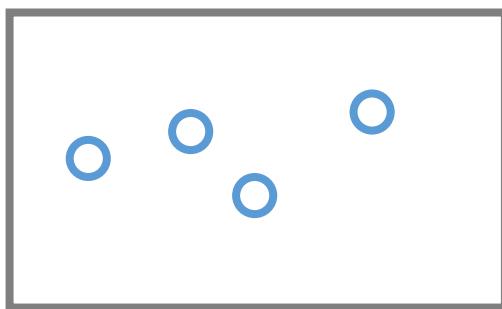

$$X_1 \in \mathbb{R}^{d \times N_1}$$

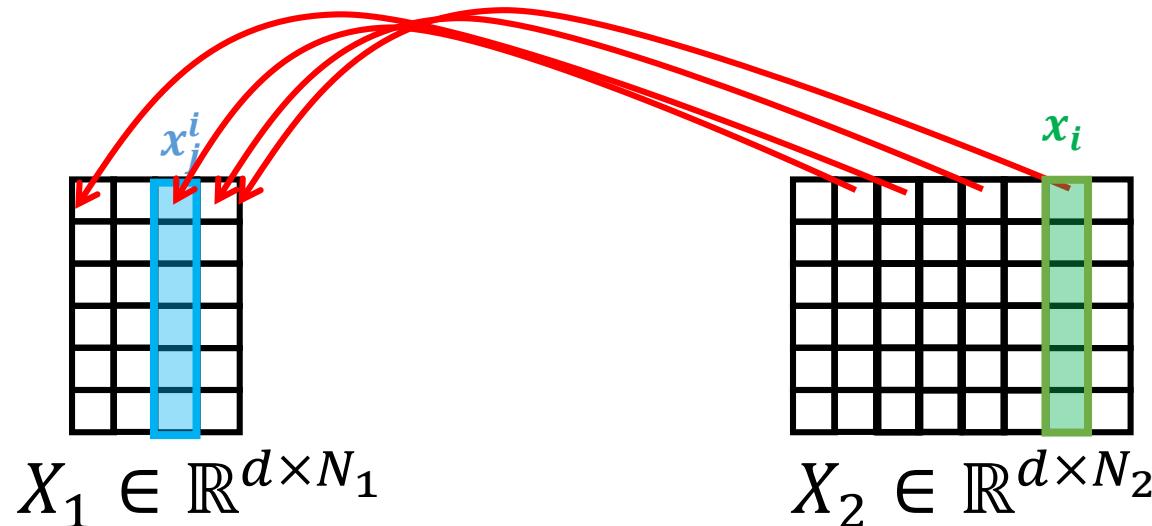
Image  
features


$$X_2 \in \mathbb{R}^{d \times N_2}$$



# Feature Matching

Input:  $X_1$  and  $X_2$ , features extracted from  $I_1$  and  $I_2$



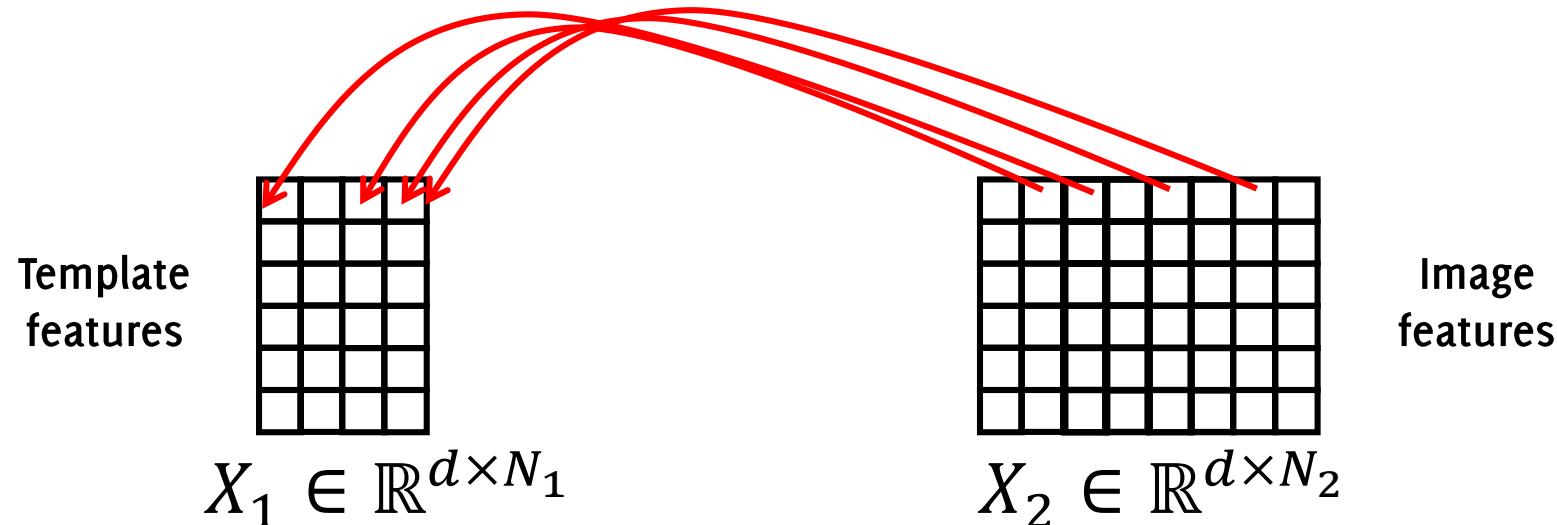
Goal: match image features on the second image to the first image features.

Identify, for each feature  $x_i \in X_2$ , the feature  $x_j^i \in X_1$  minimizing the Euclidean distance

$$x_j^i = \operatorname{argmin}_{x_j \in X_1} \left( \|x_i - x_j\|_2 \right)$$

# Feature Matching

Input:  $X_1$  and  $X_2$ , features extracted from  $I_1$  and  $I_2$



Goal: match image features on the second image to the first image features.

Identify, for  
distance

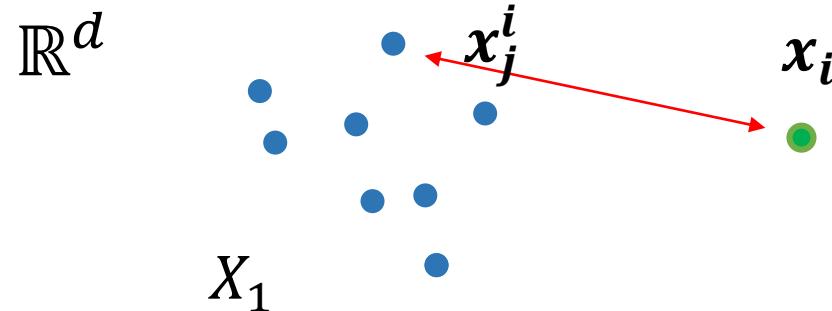
This is the **Nearest-Neighbor Matching Problem**  
(on high-dimensional data)

A central problem in CV, ML, document retrieval,  
data analysis, bioinformatics, compression

$$x_j \in \mathbb{R}^d$$

# Feature Matching

One option is **linear search**: exhaustively looking for the closest point in a loop

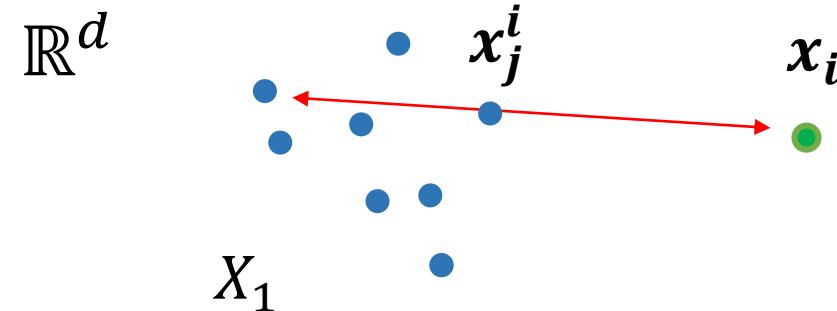


Finding nearest neighbor matches to high dimensional vectors of the training set is **one of the most computationally expensive part of CL and ML algorithms**, it is  $O(n)$  but training set are typically very large

This is computationally infeasible in cases of practical interest, and **more efficient solutions are needed**

# Feature Matching

One option is **linear search**: exhaustively looking for the closest point in a loop

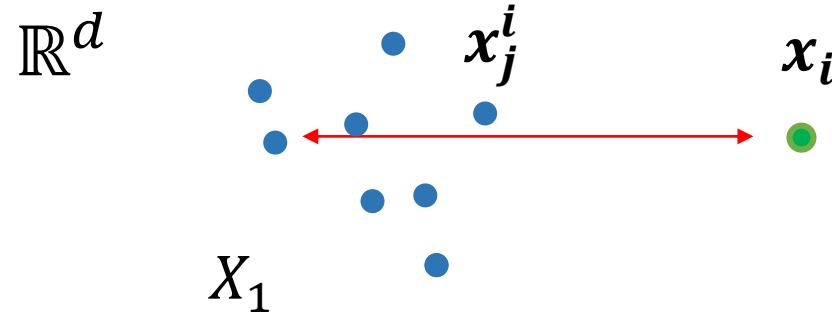


Finding nearest neighbor matches to high dimensional vectors of the training set is **one of the most computationally expensive part of CL and ML algorithms**, it is  $O(n)$  but training set are typically very large

This is computationally infeasible in cases of practical interest, and **more efficient solutions are needed**

# Feature Matching

One option is **linear search**: exhaustively looking for the closest point in a loop

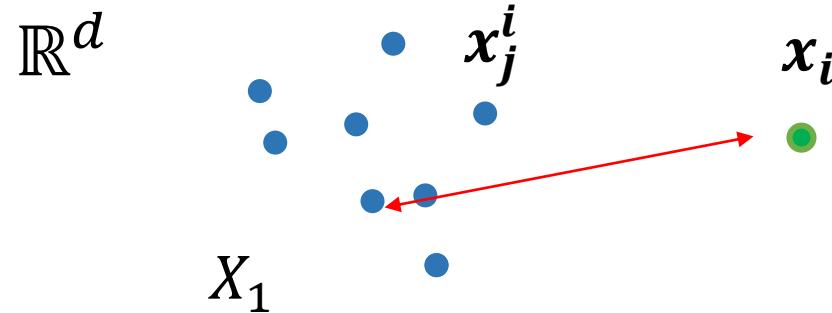


Finding nearest neighbor matches to high dimensional vectors of the training set is **one of the most computationally expensive part of CL and ML algorithms**, it is  $O(n)$  but training set are typically very large

This is computationally infeasible in cases of practical interest, and **more efficient solutions are needed**

# Feature Matching

One option is **linear search**: exhaustively looking for the closest point in a loop

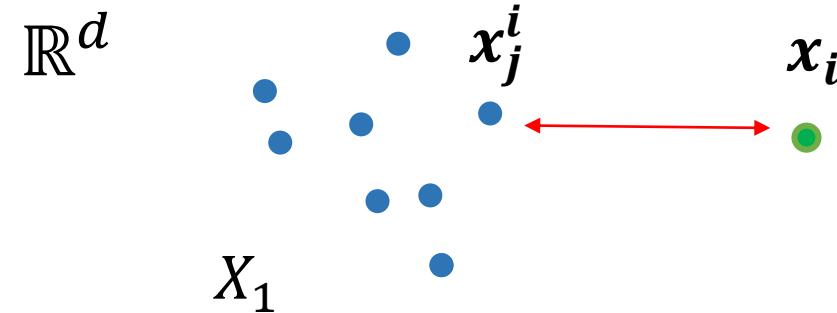


Finding nearest neighbor matches to high dimensional vectors of the training set is **one of the most computationally expensive part of CL and ML algorithms**, it is  $O(n)$  but training set are typically very large

This is computationally infeasible in cases of practical interest, and **more efficient solutions are needed**

# Feature Matching

One option is **linear search**: exhaustively looking for the closest point in a loop



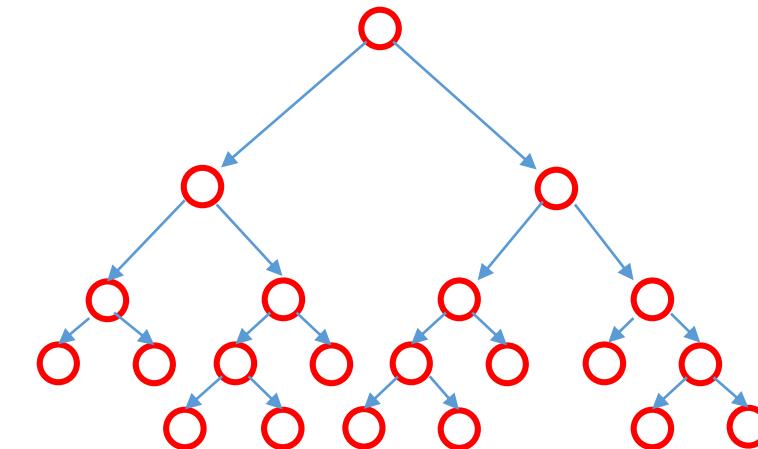
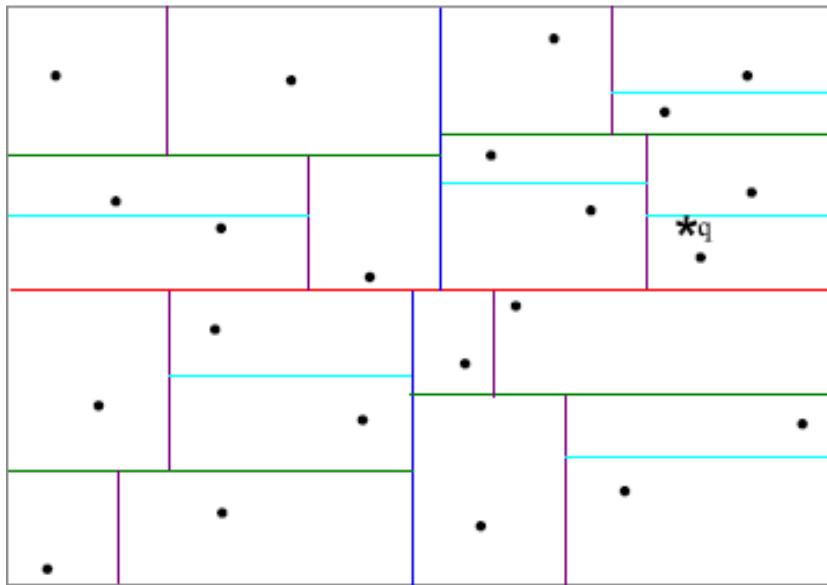
Finding nearest neighbor matches to high dimensional vectors of the training set is **one of the most computationally expensive part of CL and ML algorithms**, it is  $O(n)$  but training set are typically very large

This is computationally infeasible in cases of practical interest, and **more efficient solutions are needed**

# K-d trees: rationale

We can do better than linear search, if we organize the descriptors in proper data-structures as K-d trees.

**K-d trees** are constructed by **recursive binary splits** and allow to efficiently search for nearest neighbor (average complexity  $O(\log(n))$ )



# FLANN: Fast Library for Approximated Nearest Neighborhood

A library which implements advanced approximated searching methods based on trees, including new algorithms:

Priority search k-means tree algorithms (which are not constructed as splits along the axis)

Hierarchical Clustering Tree (meant for binary features)

# Pruning Matches

# Feature matching overview

## Estimating Image Correspondences

- Extract features from each image
  - Keypoint detection
  - Descriptor Computation
- Match features between images
- Prune matches and then perform triangulation
- detect objects / stitching:
  - Ratio test
  - Geometric verification

# Remove matches that are not good enough

The major issue:

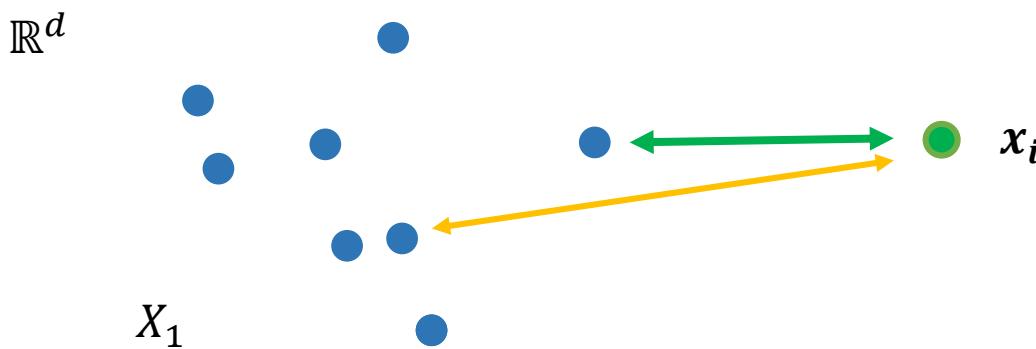
- There does not exist a reference value for a descriptor distance for good/wrong matches
- Descriptor distance can vary a lot from scene to scene, and feature to feature

Matches have to be pruned by looking at their **relative** distance

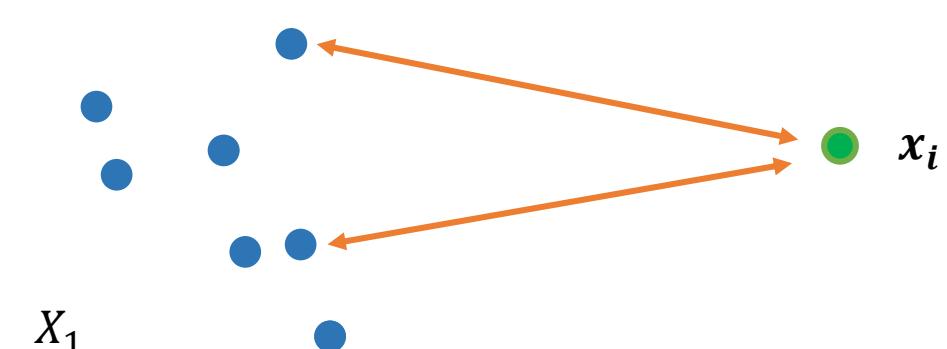
# Ratio Test

By nearest neighbor search we get for each image feature  $x_i \in X_2$ , the closest template feature  $x_j^i \in X_1$

Wrong matches  $(x_i, x_j^i)$  need still to be rejected.



✓ There is a clear closest match



✗ Matches are ambiguous, the distance to the best match is similar to the one with the second-best match

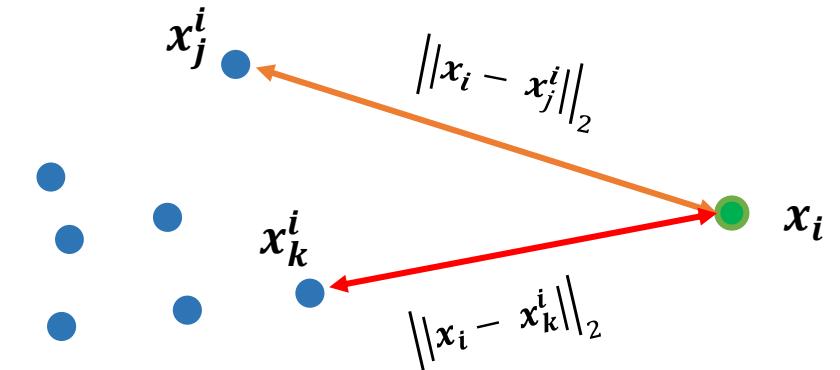
# Ratio Test

During matches, retrieve the 2-NN neighbor of each  $x_i$ , i.e.

$$(x_i, x_j^i) \text{ and } (x_i, x_k^i)$$

Discard keypoints where

$$\frac{\|x_i - x_k^i\|_2}{\|x_i - x_j^i\|_2} > 0.8$$



This criteria discards matches where the second nearest neighbor is very close to the first. These are:

- Ambiguous matches
- False matches arising from background clutter

# Application: template matching



# Application: template matching



# Application: template matching



# Application: template matching

Repeated structures



a alamy stock photo

J3TXBT  
www.alamy.com



template

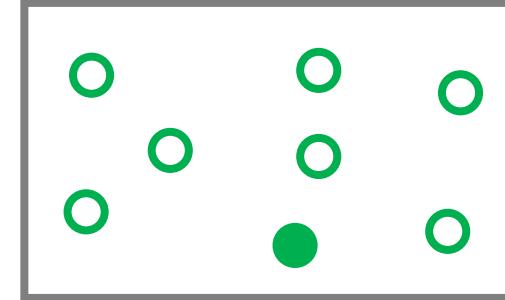
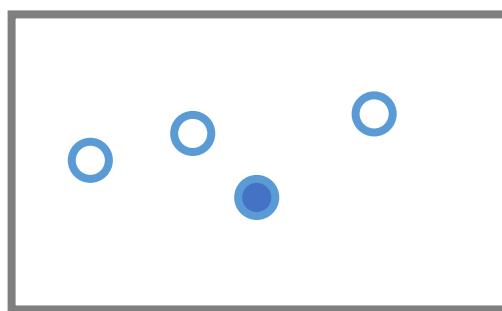
Giacomo Boracchi

# Matches preserved by the ratio test



# Geometric verification

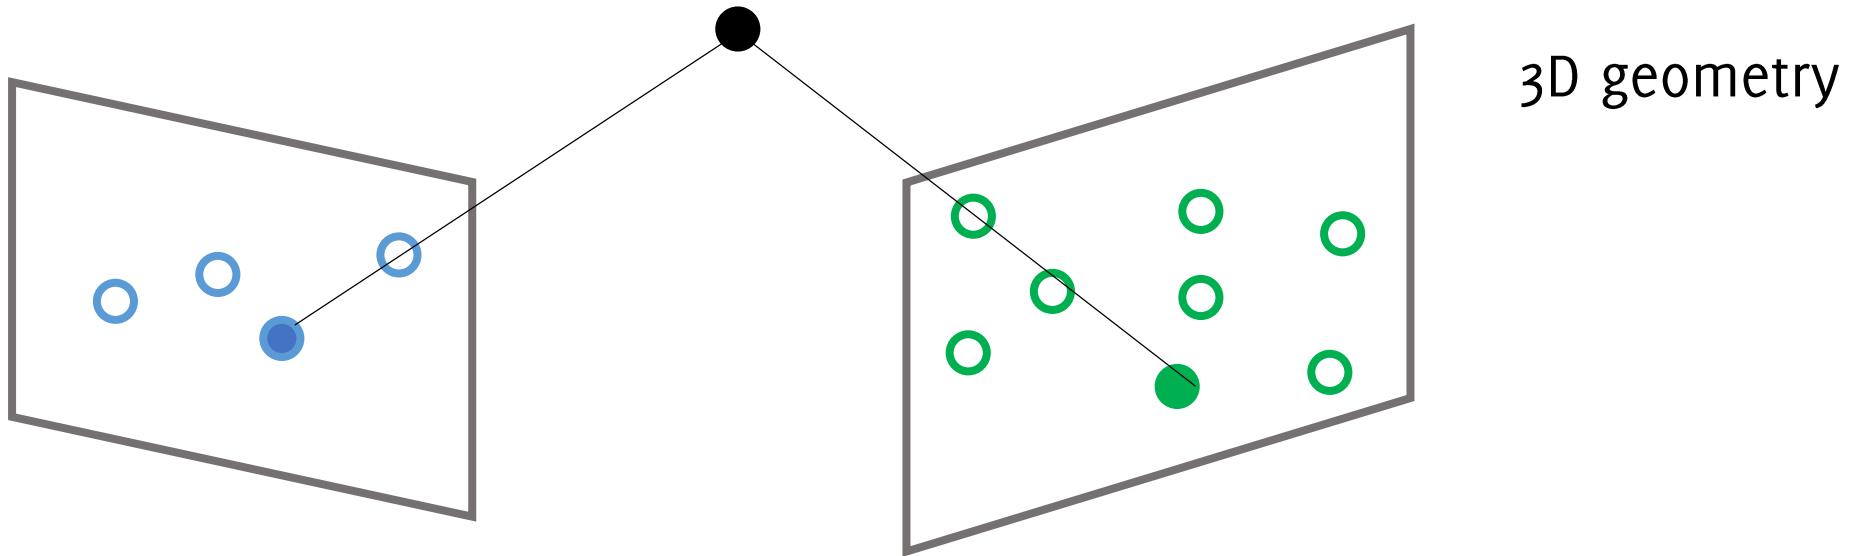
Until now we have searched for corresponding points leveraging on image appearance. But we can also exploit 3D geometric constraints



Do these two points look similar?

# Geometric verification

Until now we have searched for corresponding points leveraging on image appearance. But we can also exploit 3D geometric constraints

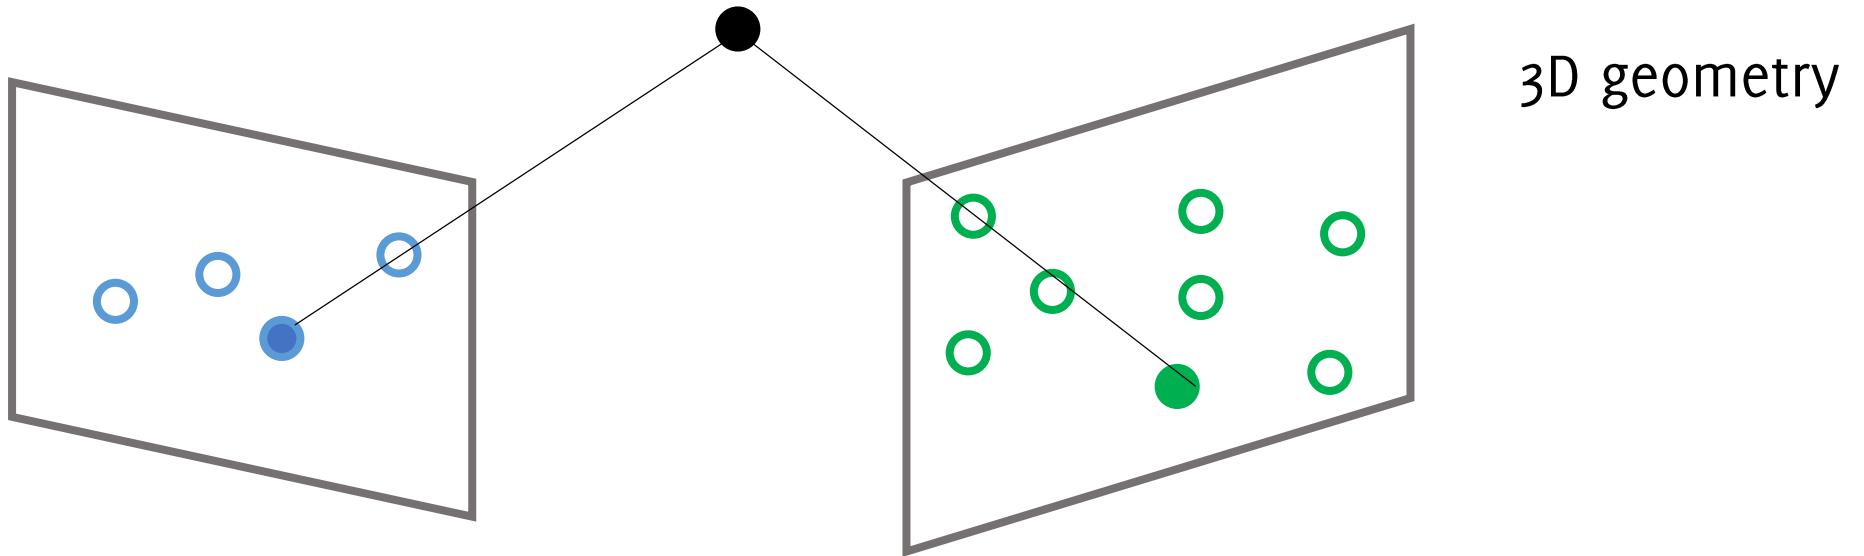


Do these two points look similar?

Can these two points be the image of the same 3D point in the scene?

# Geometric verification

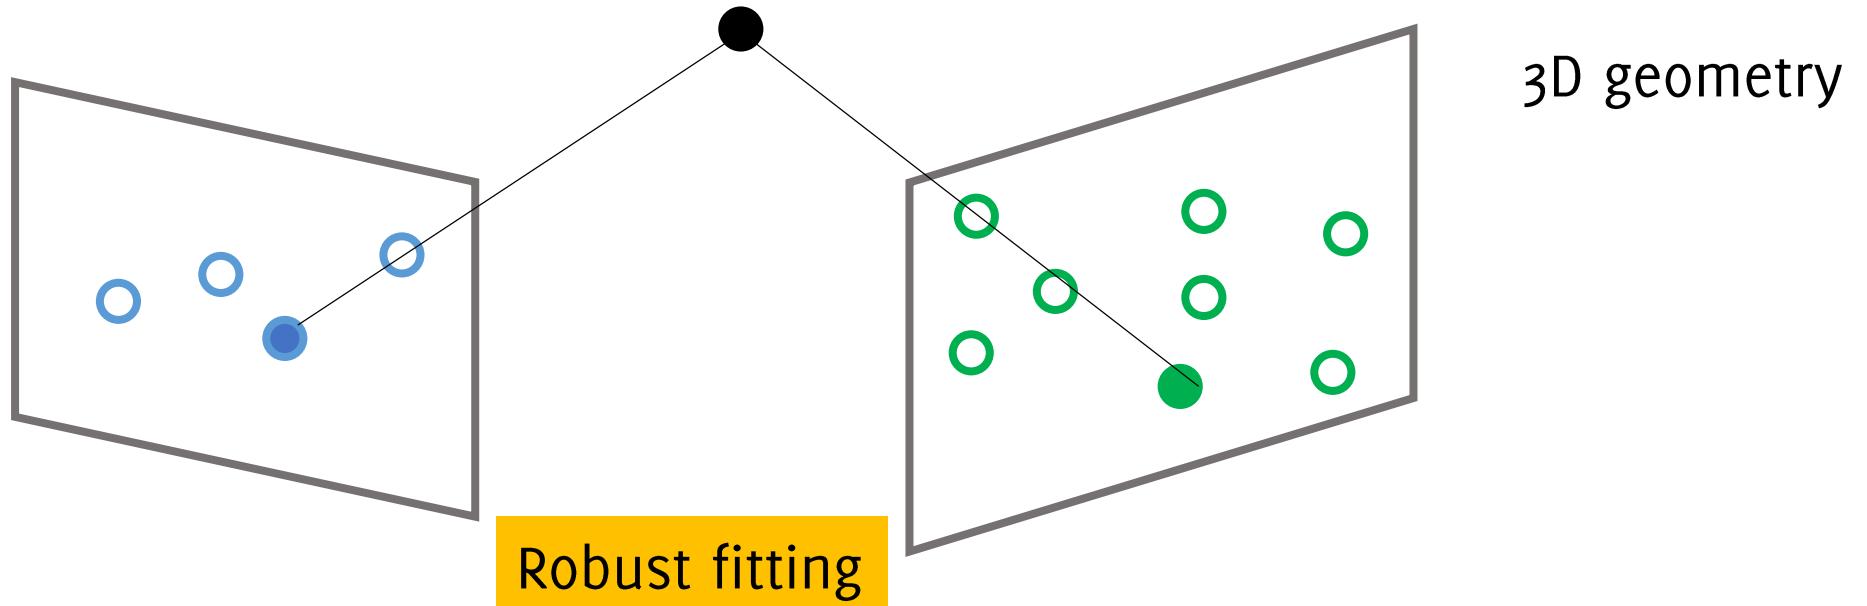
Until now we have searched for corresponding points leveraging on image appearance. But we can also exploit 3D geometric constraints



Can these two points be the image of the same 3D point in the scene?  
**Pairs of corresponding points must satisfy epipolar geometry**

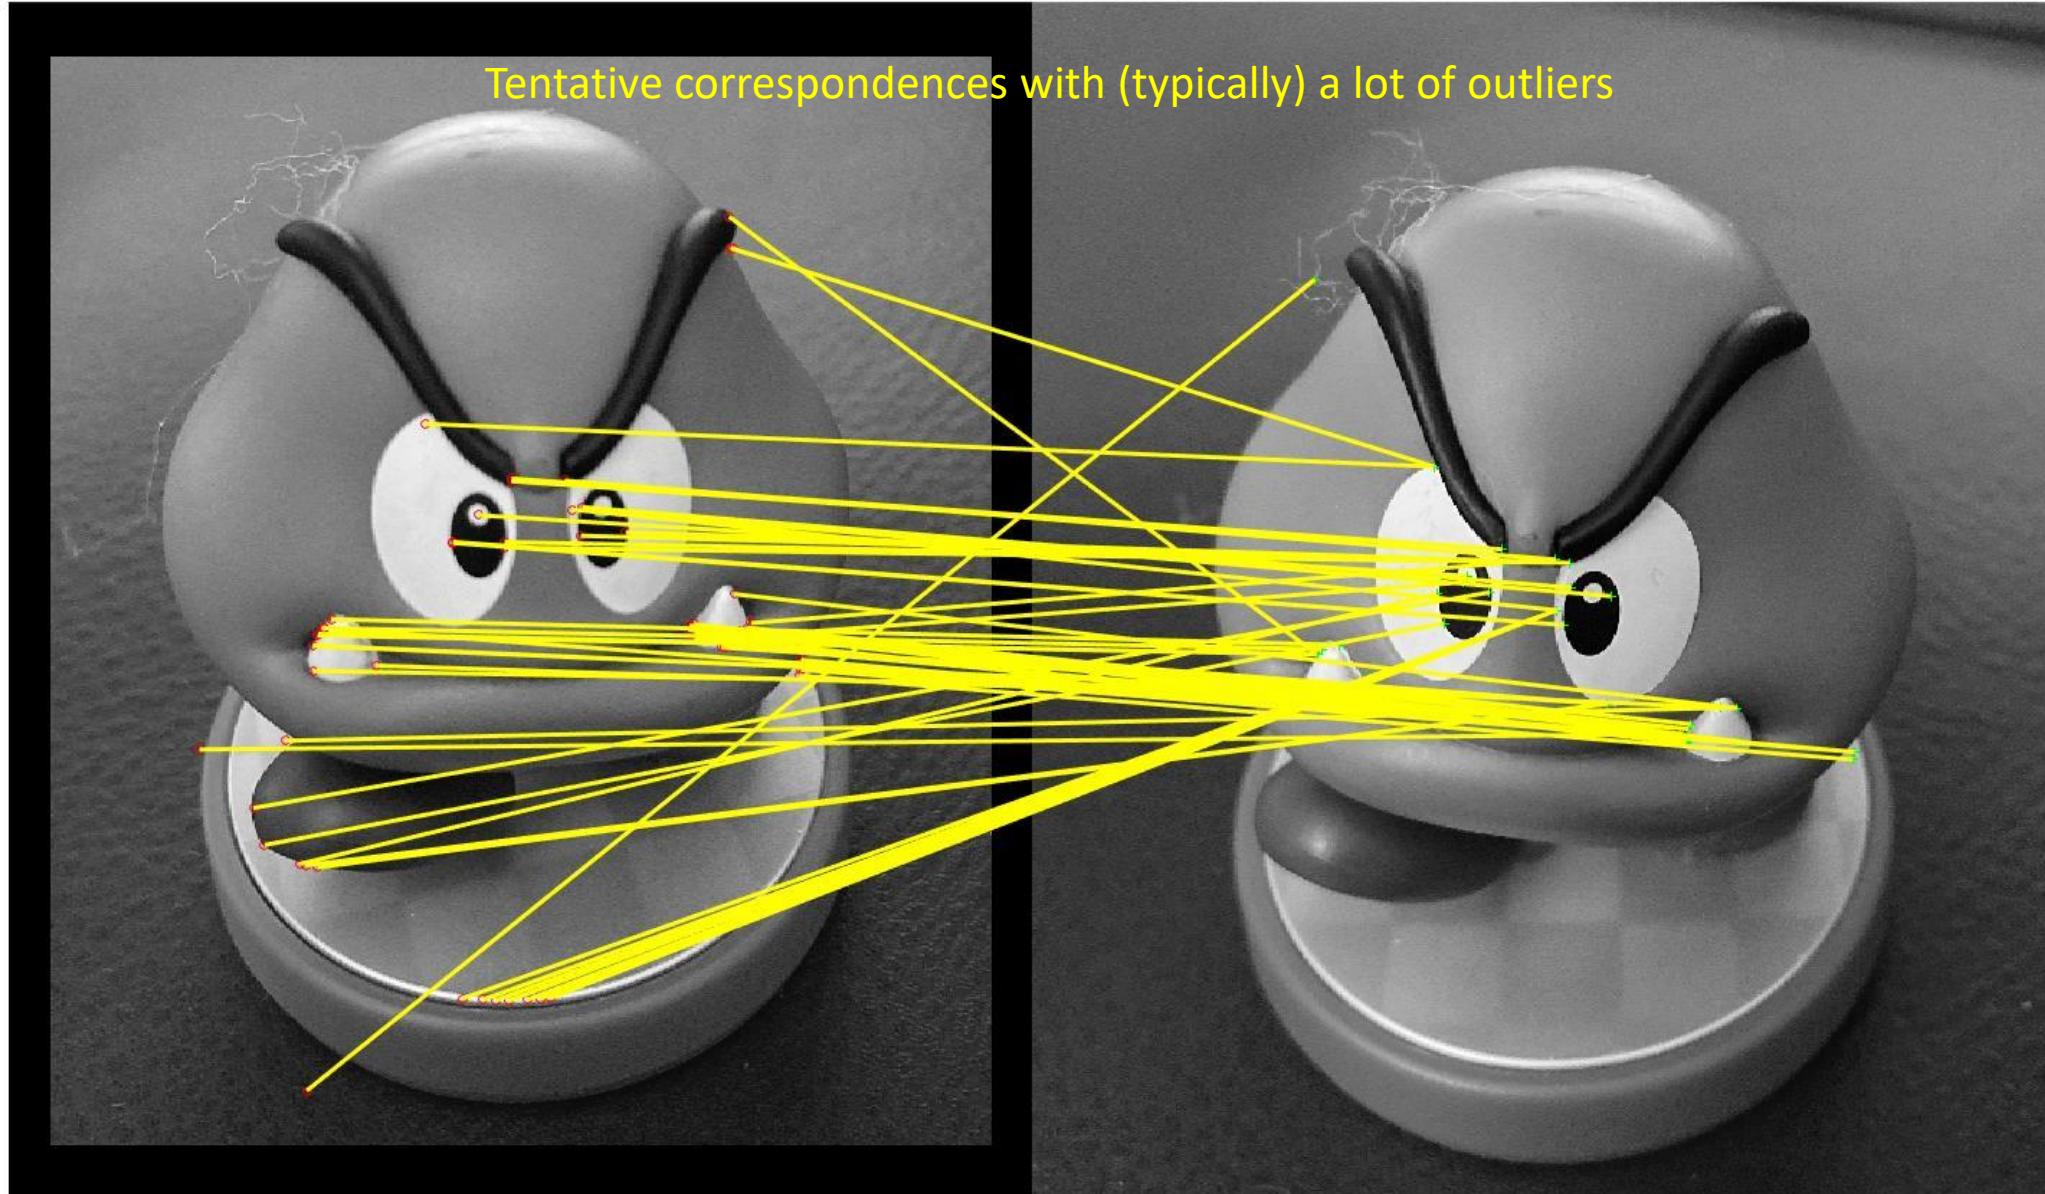
# Geometric verification

Until now we have searched for corresponding points leveraging on image appearance. But we can also exploit 3D geometric constraints



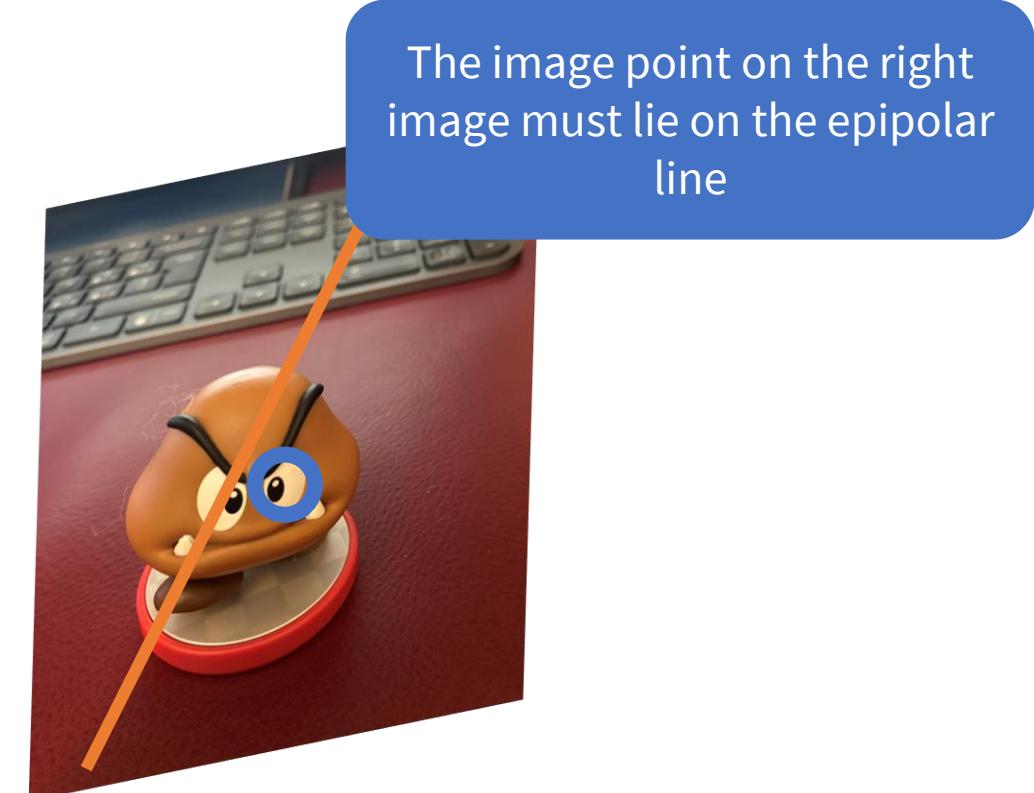
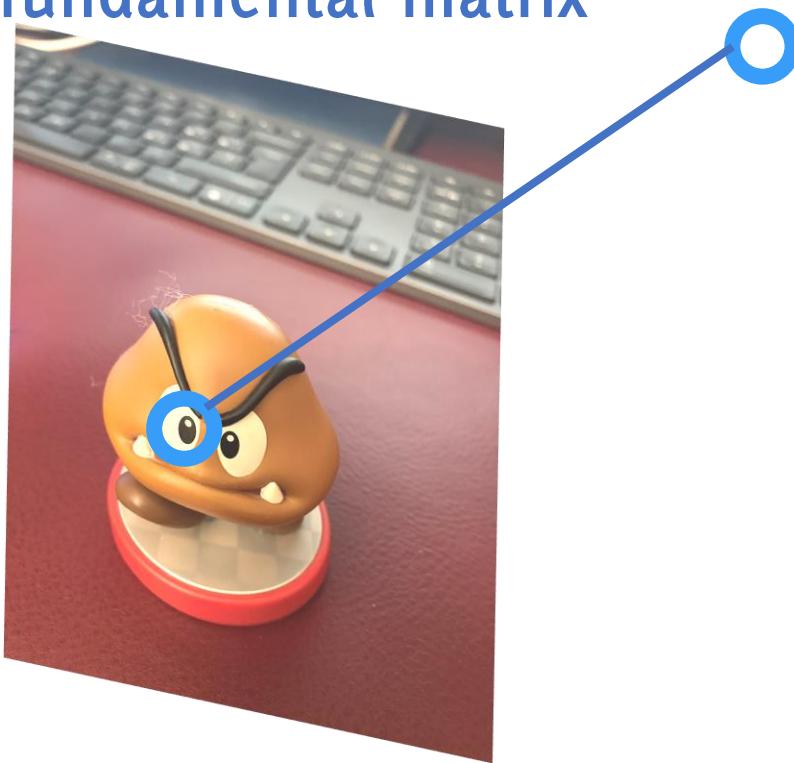
Robustly fit (e.g. via RANSAC) a fundamental matrix on matches and keep only the inliers

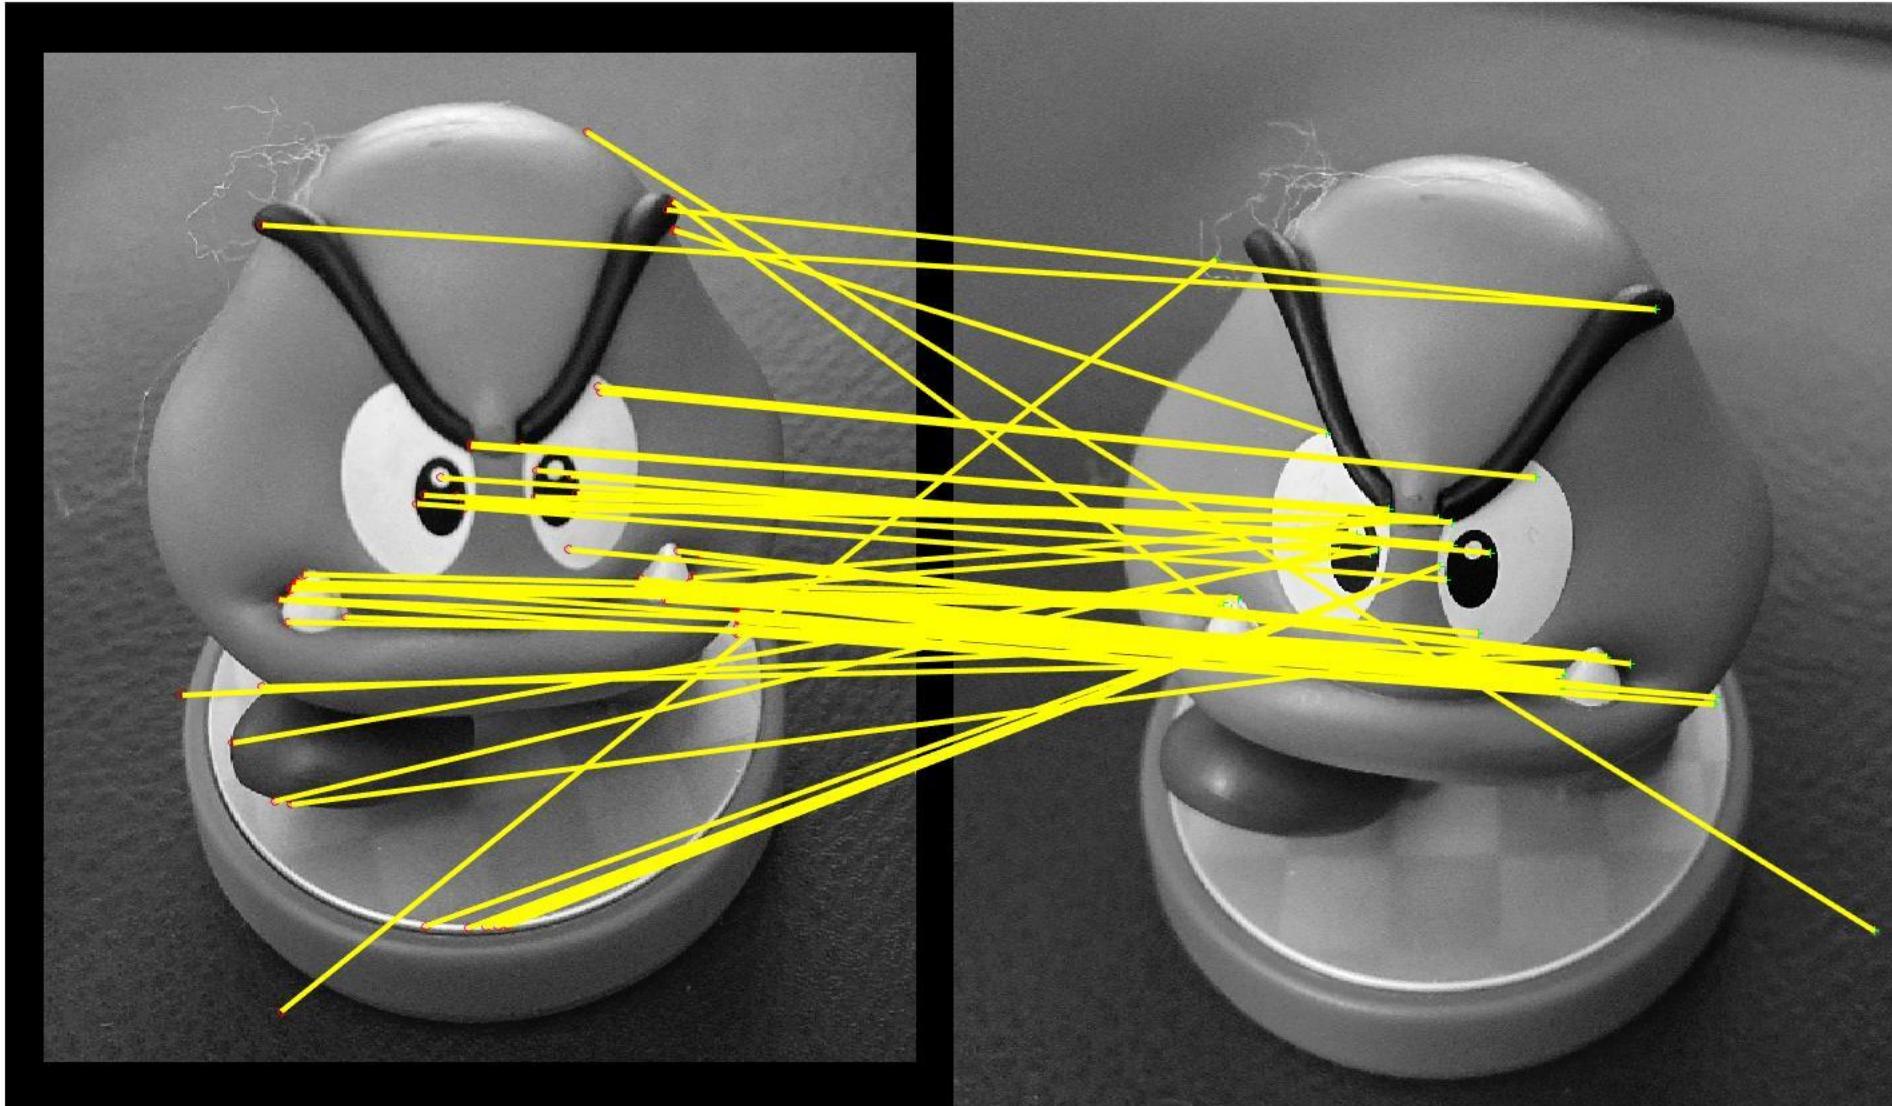
If the scene is planar or the camera movement is pure rotational an homography can be used to validate matches.



# Fit the epipolar matrix

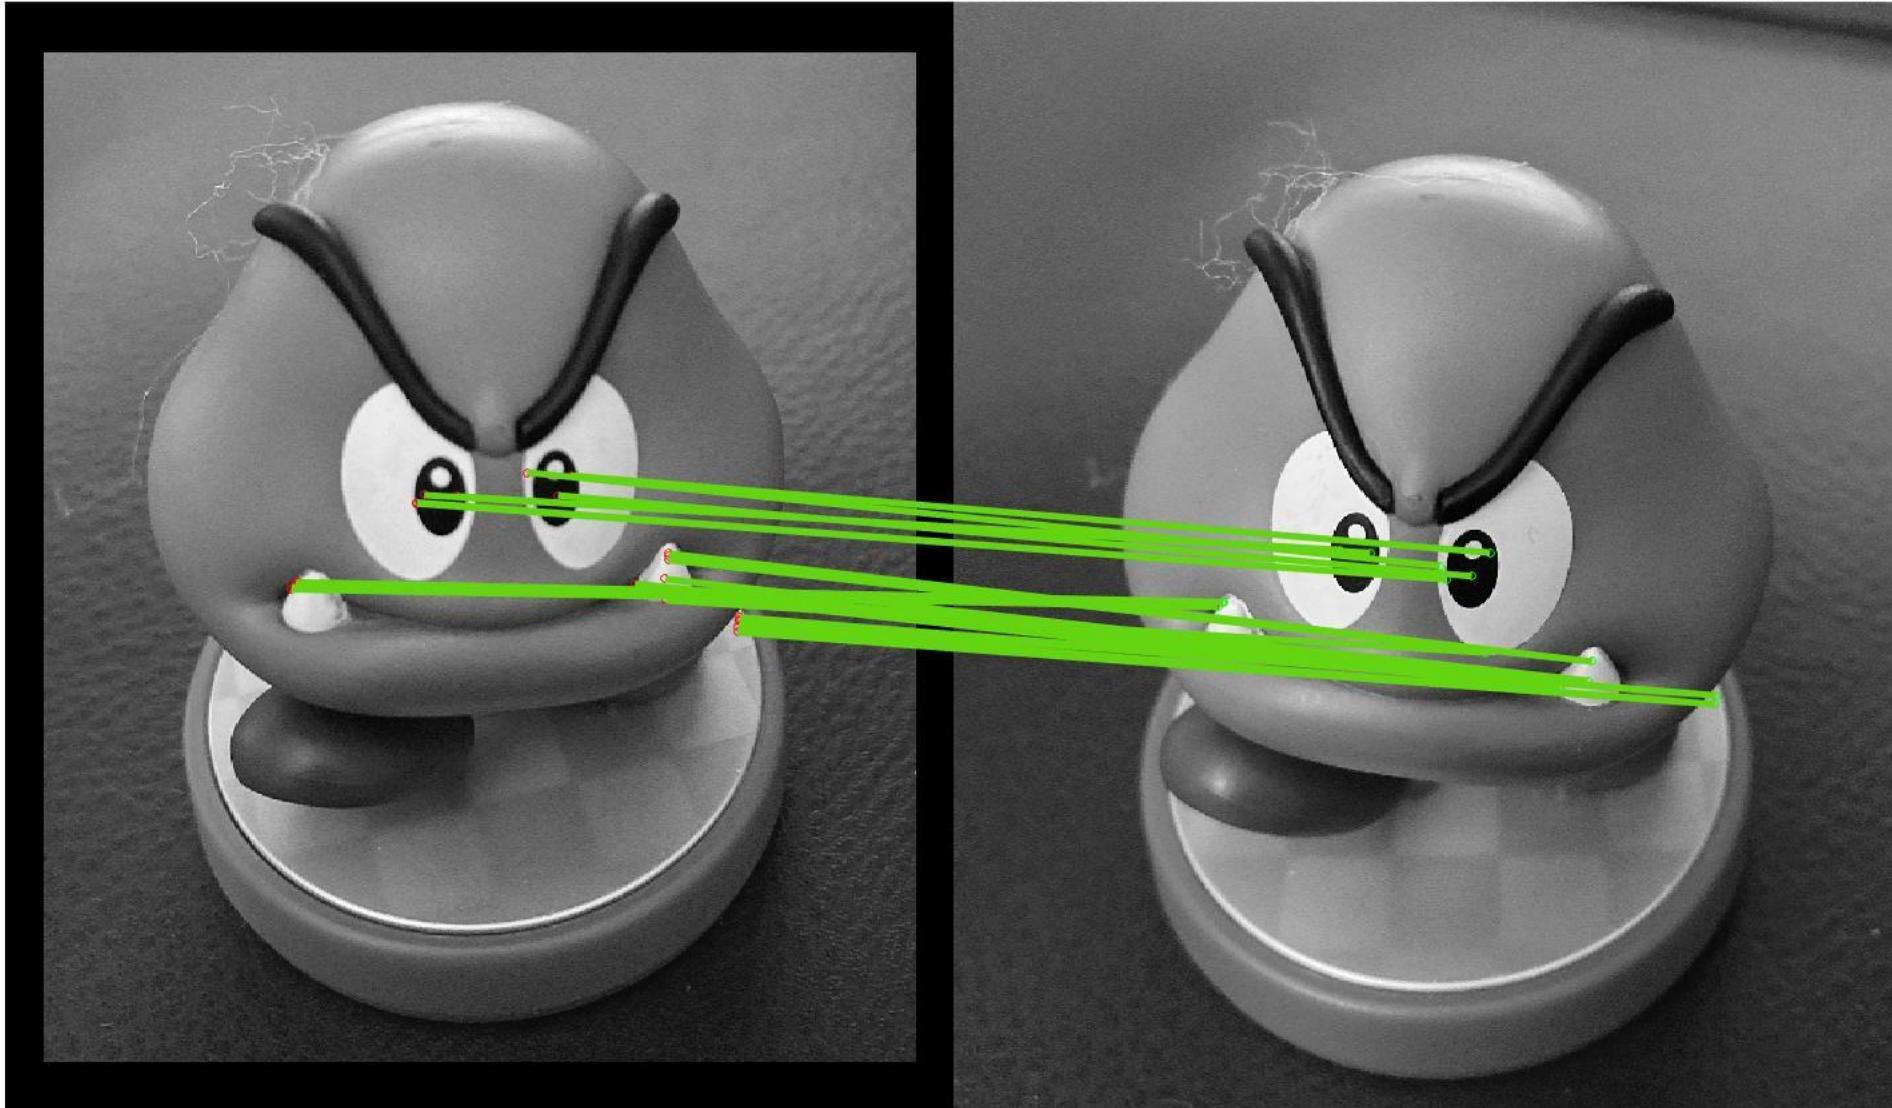
You can fit a parametric model to disambiguate between inliers and outliers: **fundamental matrix**





$X = \{(x_{\ell i}, x_{r i}) \mid i = 1, \dots, n\}$  data are set of correspondences

$F$  the model is a fundamental matrix



$X = \{(x_{\ell i}, x_{r i}) \mid i = 1, \dots, n\}$  data are set of correspondences

$F$  the model is a fundamental matrix

# In case of planar objects

The distance ratio test have discarded many false matches

Still, many matches belong to different objects.

We need to cluster features belonging to the same object

To this purpose we need a model:

- For instance: “*planar objects seen from two different views are related by an homography*”

Homographies are  $H \in \mathbb{R}^{3 \times 3}$  linear transformation from

$$H: \mathbb{R}^2 \rightarrow \mathbb{R}^2$$

which can be computed by 4 matches.

# Adding Homography Constraint



The whole processing can be iterated to match more template in the image.

# Adding Homography Constraint

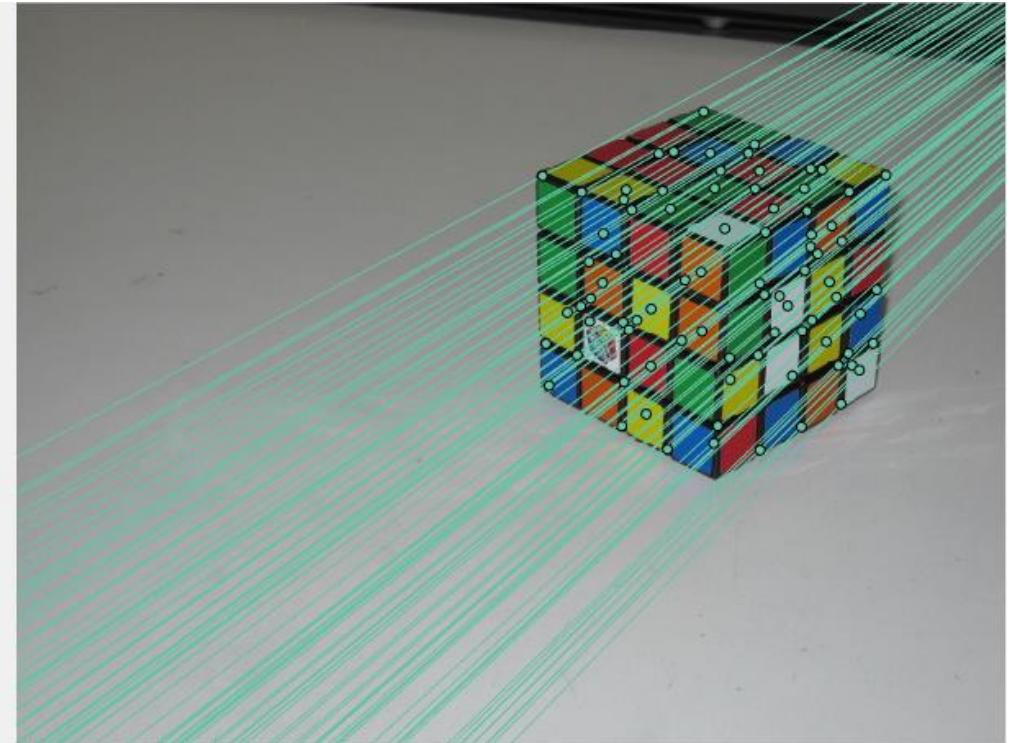
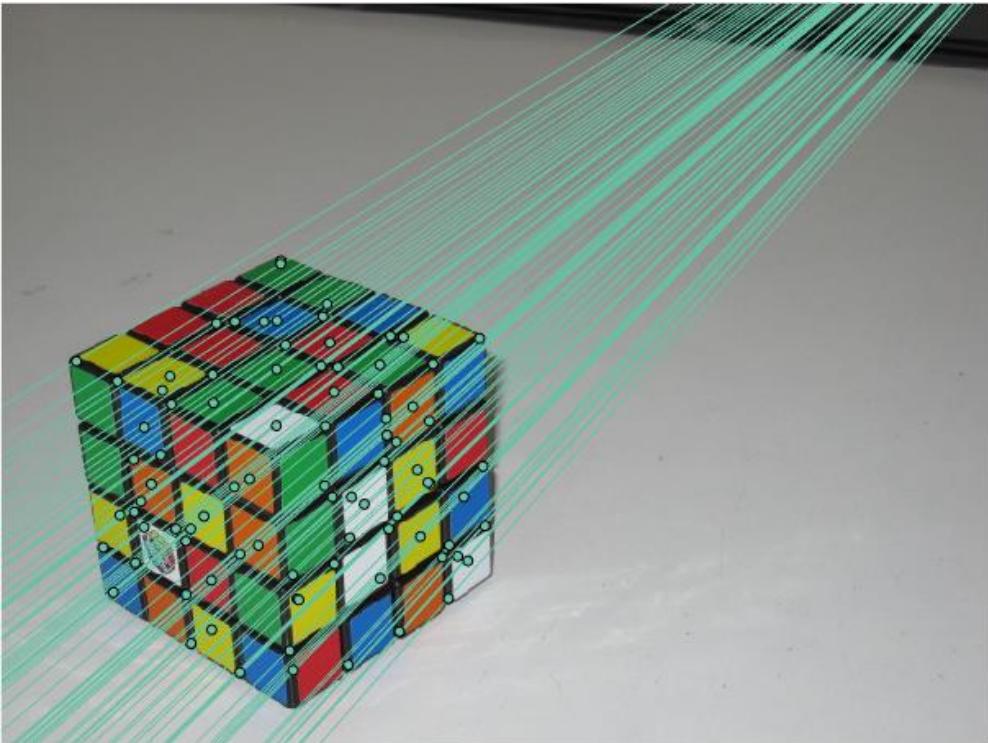


The whole processing can be iterated to match more template in the image.  
Each estimated homography can be used to rectify the each detected region and make it pixel-wise comparable with the template



# Robust Fitting

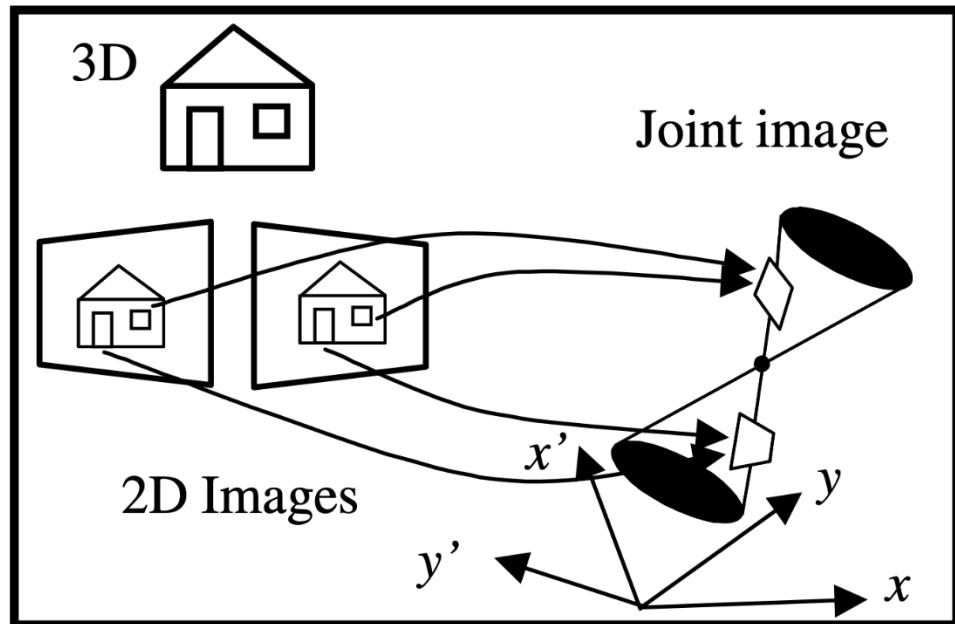
# Example: Estimating Fundamental Matrix



$$\theta = F; \\ X = \{x_1, x_2\} \quad x_2^t F x_1 = 0$$

# Example: Estimating Fundamental Matrix

It can be thought as a **generalization of a conic fitting problem** in the joint image space of pairs of correspondences.



Each match is a pair of points and defines a vector in  $\mathbb{R}^4$ .  
For each correct match, the pair of corresponding points satisfies the epipolar constraint (bilinear relation) and lies on a 4D cone that can be fitted from 8 correspondences.

$$\theta = F;$$

$$X = \{x_1, x_2\} \quad x_2^t F x_1 = 0$$

Image from Anandanand Avidan. Integrating Local Affine into Global Projective Images in the Joint Image Space. ECCV 2020

# Fitting

**Input:**  $X$  Data

**Output:**  $\theta$  parameters of the model that describes matches

Our aim is to fit  $\theta$  to the data  $X$  to recover a “geometric pattern” described by the model (structure)

In the previous example we want to recognize those correspondences that are compatible with a rigid ego-motion (of the scene/of the cameras).

This framework can be found in several computer vision applications.

# Example: Estimating Homographic Transformations



# Example: Line Fitting for Vanishing Point Estimation



$$\theta = (a, b, c), X = \{(x, y)\}$$
$$ax + by + c = 0$$

# Example: Conic Fitting



$$\theta = (r, c_x, c_y); X = \{(x, y)\}$$
$$(x - c_x)^2 + (y - c_y)^2 = r^2$$

# Ordinary Least Square

# Least Square Regression

Given a set of  $N$  points (or matches..)

$$X = \{(x_1, y_1), \dots, (x_N, y_N)\}$$

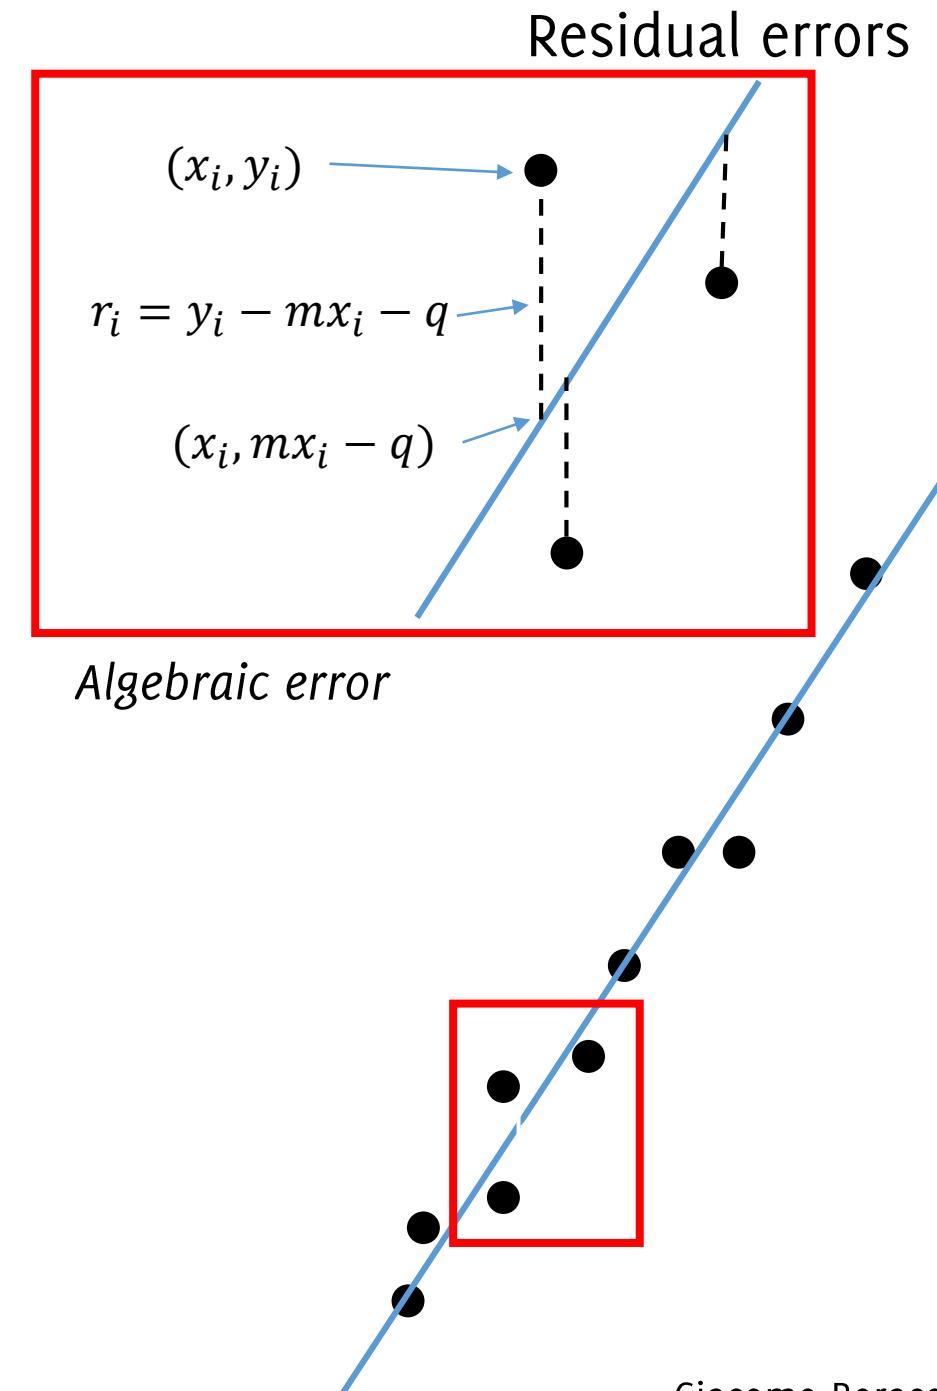
Given a parametric model

$$y = mx + q$$

Estimate the parameters  $\theta = \{m, q\}$  yielding  
the **best fit**

The best fit is the one **minimizing some**  
**residual error over  $X$**

$$E = \sum_{i=1}^N (r_i)^2 = \sum_{i=1}^N (y_i - mx_i - q)^2$$



This error does not make sense  
when the line is vertical

# Minimizing the point-line distance

Given a set of  $N$  points (or matches..)

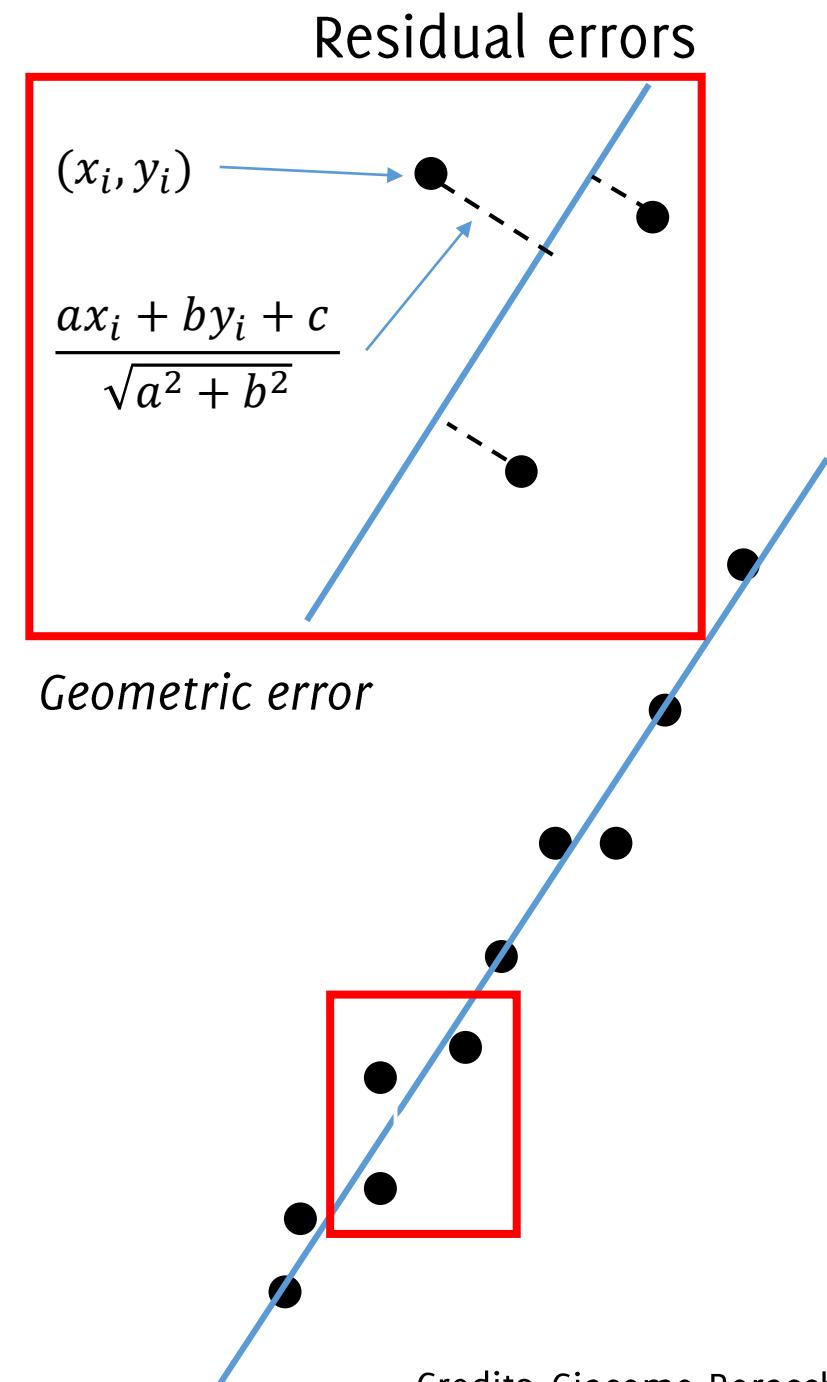
$$X = \{(x_1, y_1), \dots, (x_N, y_N)\}$$

Given a parametric model

$$ax + by + c = 0$$

Estimate the line parameters  $\theta = \{a, b, c\}$   
minimizing the residual error

$$E = \sum_{i=1}^N (ax_i + by_i + c)^2$$

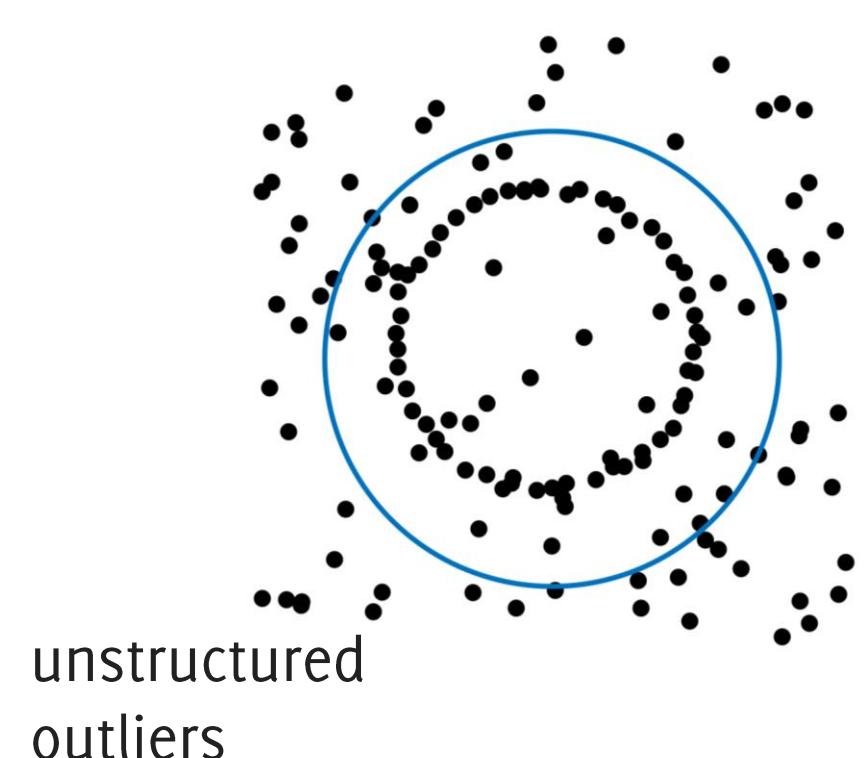
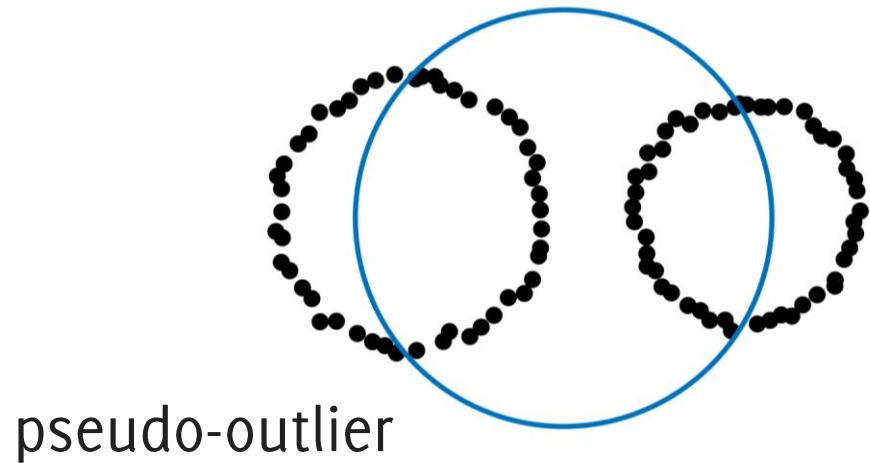
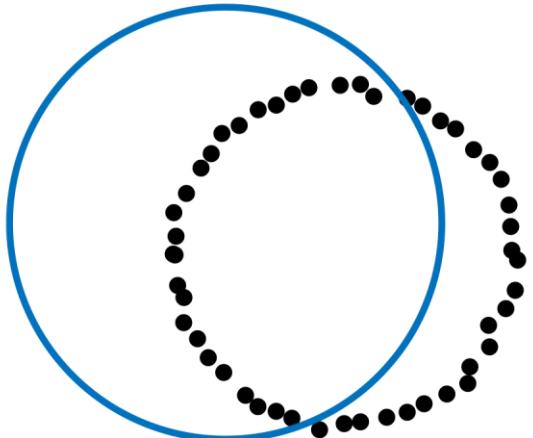


# Least squares breaks down

**Break down point:** the proportion of incorrect observations that can be handled before giving an arbitrarily large uncorrect result.

✗ Least squares has 0% breakdown point  
(the outlier might be arbitrarily large)

   
single outlier

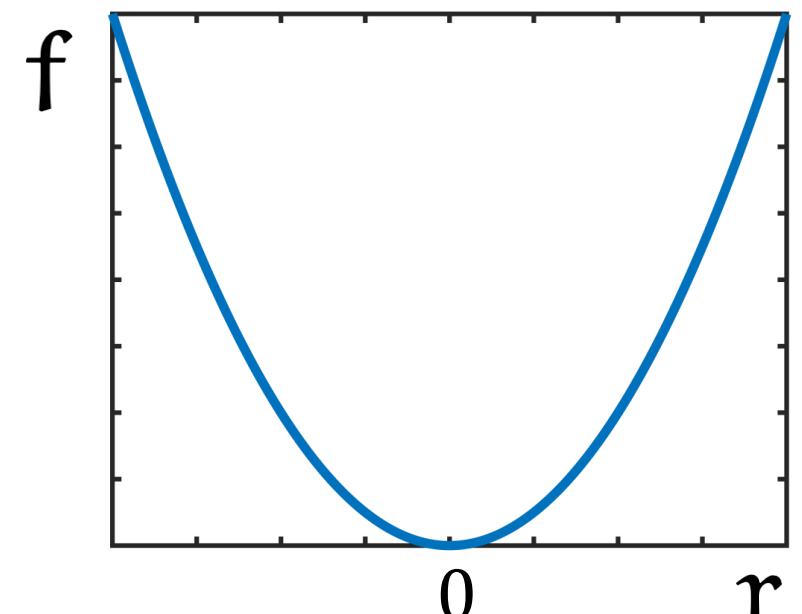


# Robustness to Outliers

# The loss function: least square

The loss so far is the sum of a function of all the residuals

$$E = \sum_{i=1}^N f(r_i), \text{ where } f(r_i) = r_i^2$$

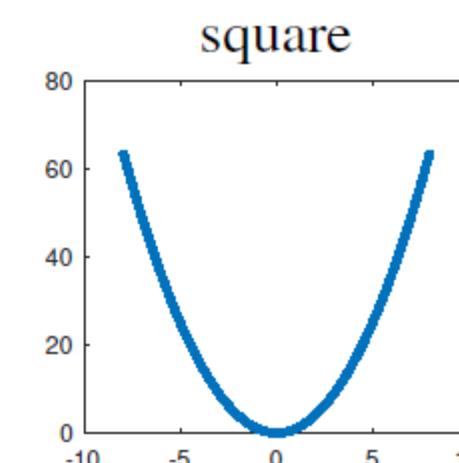
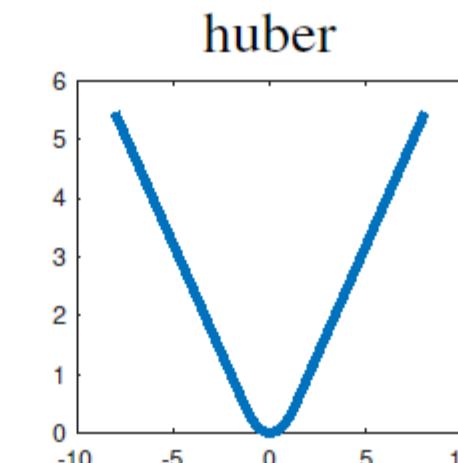
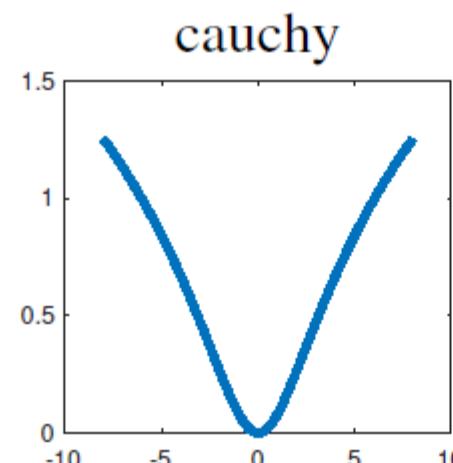
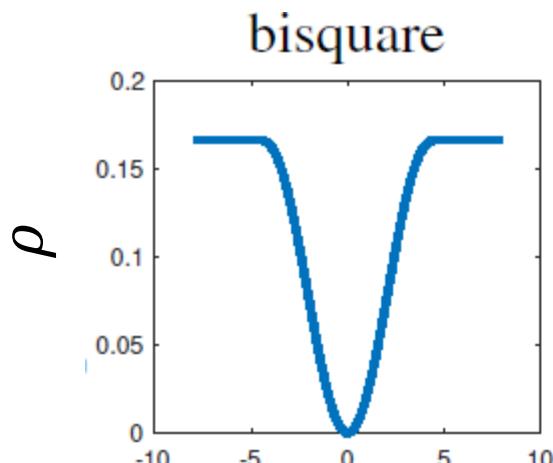


# M-Estimator

Replaces the squared loss with a different loss function which penalizes less large residual values (deemed to correspond to outliers)

$$\hat{\theta} = \operatorname{argmin}_{\theta} \sum_{i=1}^N \rho(r_i(\theta))$$

Where  $\rho$  a symmetric, positive-definite function having a unique minimum in zero



RanSaC

# Robust single model fitting: consensus maximisation

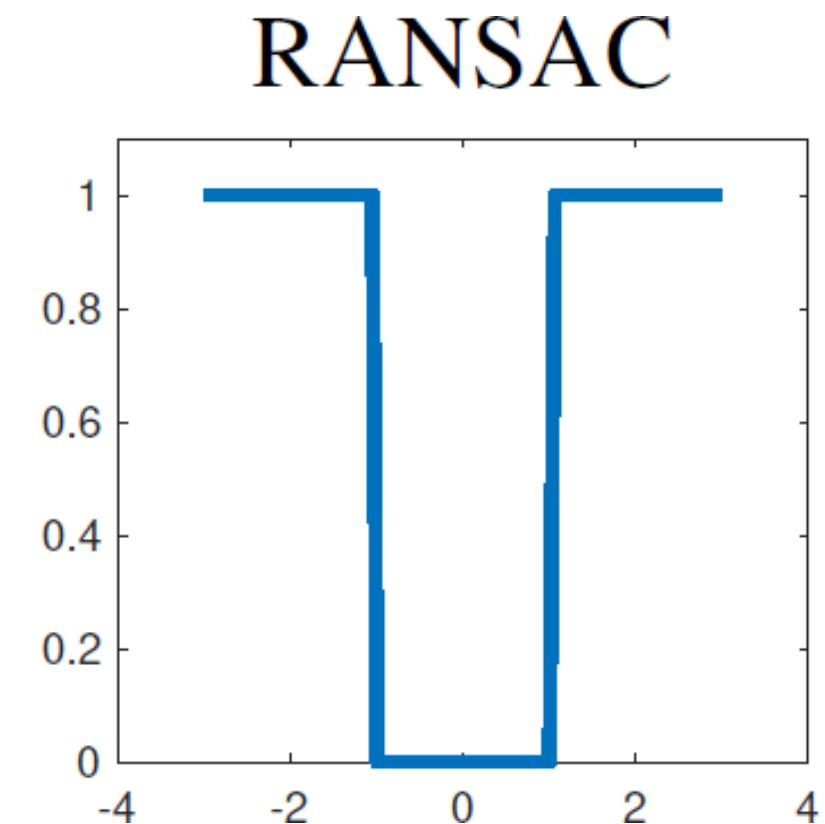
Instead of minimizing the residuals, **maximize the consensus**. Define:

- an inlier threshold  $\epsilon > 0$
- a consensus function  $\tilde{f}$  which is

$$\tilde{f}(r_i) = \begin{cases} 1, & r_i \leq \epsilon \\ 0, & r_i > \epsilon \end{cases}$$

Identify  $\hat{\theta}$  that maximizes the consensus

$$\hat{\theta} = \arg \min_{\theta} \sum_{i=1}^N 1 - \tilde{f}(r_i)$$



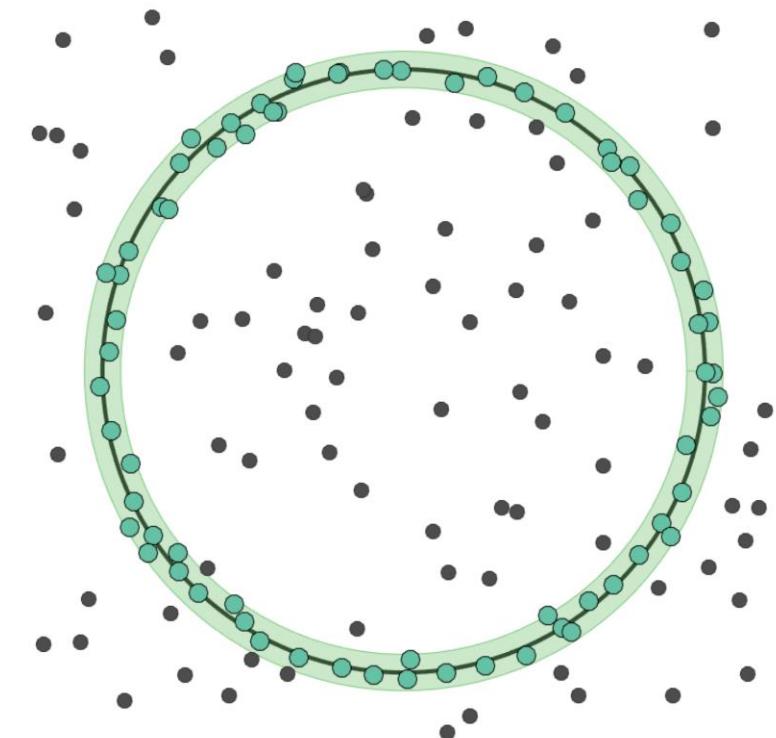
# The Consensus Set

Consensus set

$$CS(\theta, \epsilon) = \{x_i \mid r_i \leq \epsilon\}$$

Being  $r_i = r(x_i, \theta)$ , the residual of the model  $\theta$  at a point  $x_i$

The larger the consensus set  $CS(\theta, \epsilon)$ , the better the model  $\theta$



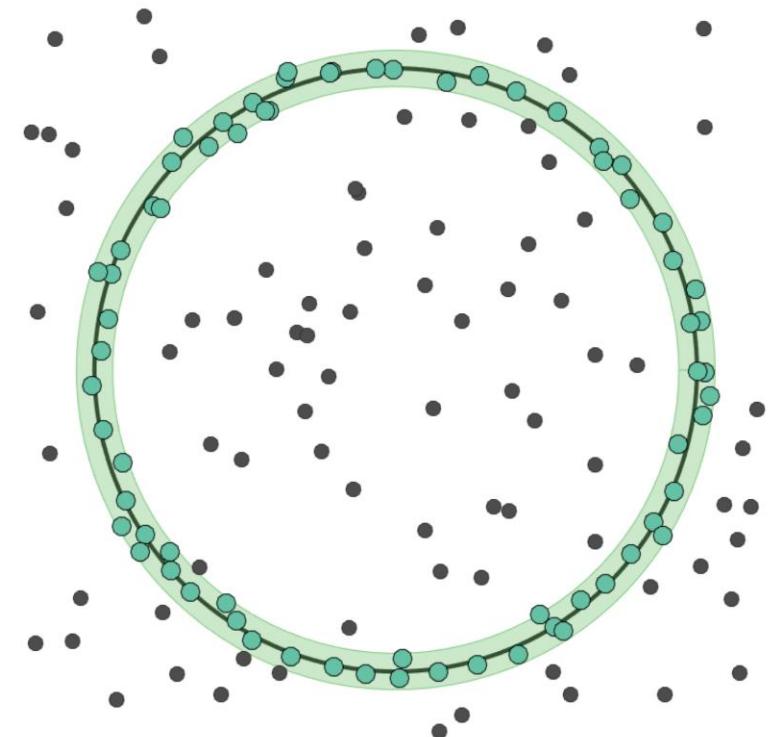
# The Consensus Set

Consensus set

$$\text{CS}(\theta, \epsilon) = \{x_i \mid r_i \leq \epsilon\}$$

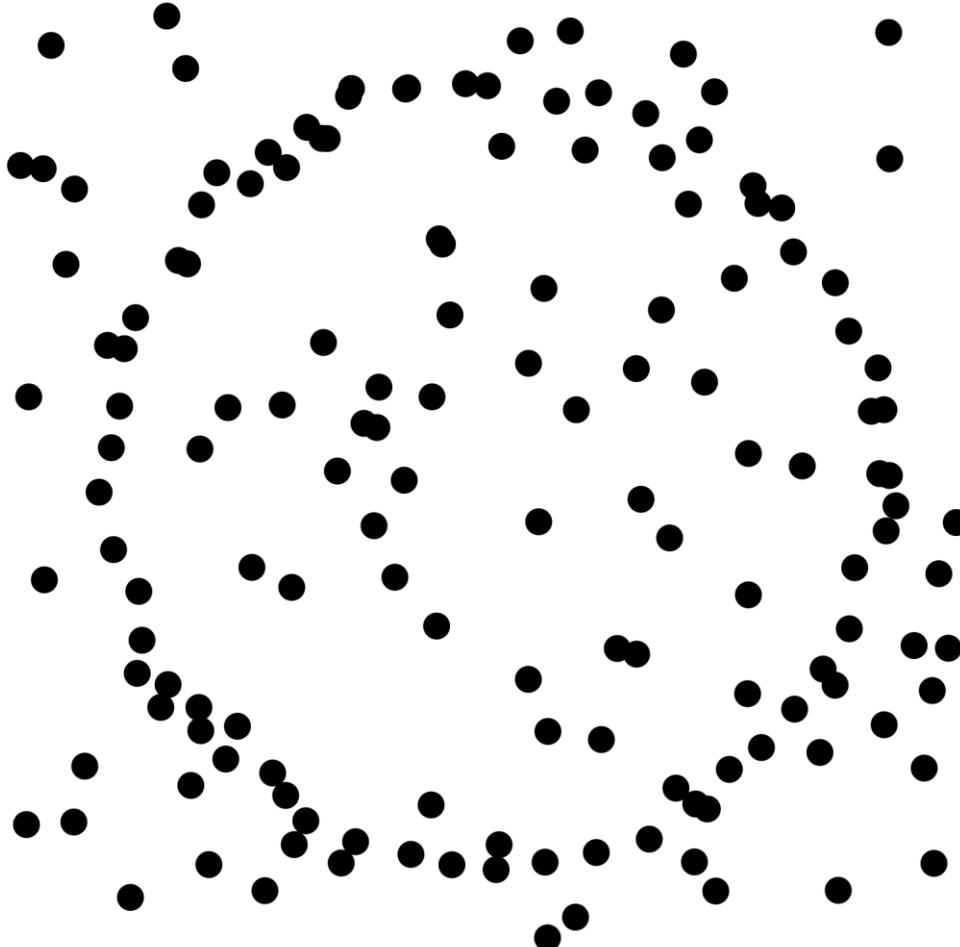
Being  $r_i = r(x_i, \theta)$ , the residual of the model  $\theta$  at a point  $x_i$

The larger the consensus set  $\text{CS}(\theta, \epsilon)$ , the better the model  $\theta$



We have been fitting lines so far, but  
**everything** holds for any parametric model  
(e.g. circle, conics, homographies,  
fundamental matrices)

# Randomized Sample Consensus [Fischler and Bolles 1981]



**Input:**  $X$  data,  $\epsilon$  inlier threshold,  $k_{\max}$  max iteration

**Output:**  $\theta^*$  model estimate

$J^* = -\infty, k = 0;$

**repeat**

Select randomly a minimal sample set  $S \subset X$ ;

Estimate parameters  $\theta$  on  $S$ ;

Evaluate  $J(\theta) = \sum_{x \in X} \hat{f}_\epsilon(r(x, \theta))$ ;

**if**  $J(\theta) > J^*$  **then**

$\theta^* = \theta$ ;

$J^* = J(\theta)$ ;

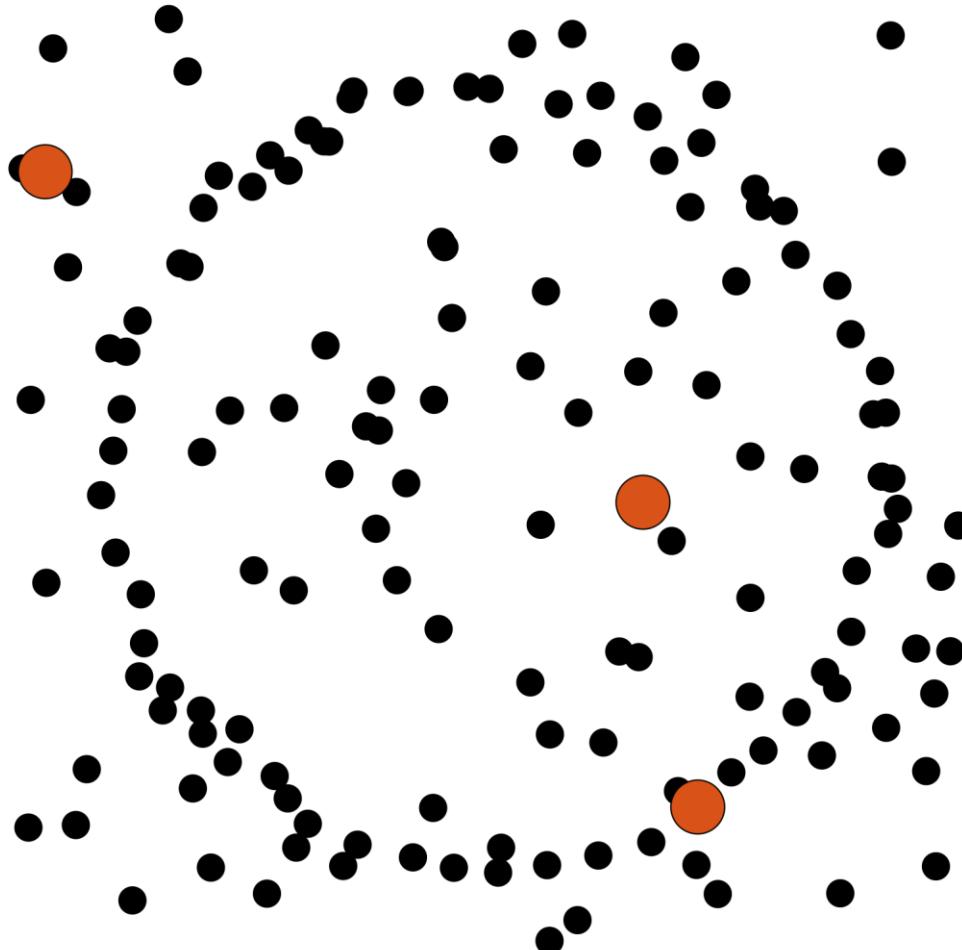
**end**

$k = k + 1$ ;

**until**  $k > k_{\max}$ ;

Optimize  $\theta^*$  on its inliers.

# Randomized Sample Consensus [Fischler and Bolles 1981]



**Input:**  $X$  data,  $\epsilon$  inlier threshold,  $k_{\max}$  max iteration

**Output:**  $\theta^*$  model estimate

$J^* = -\infty, k = 0;$

**repeat**

Select randomly a minimal sample set  $S \subset X$ ;

Estimate parameters  $\theta$  on  $S$ ;

Evaluate  $J(\theta) = \sum_{x \in S} \hat{f}_\epsilon(r(x, \theta))$ ;

**if**  $J(\theta) > J^*$  **then**

$\theta^* = \theta$ ;

$J^* = J(\theta)$ ;

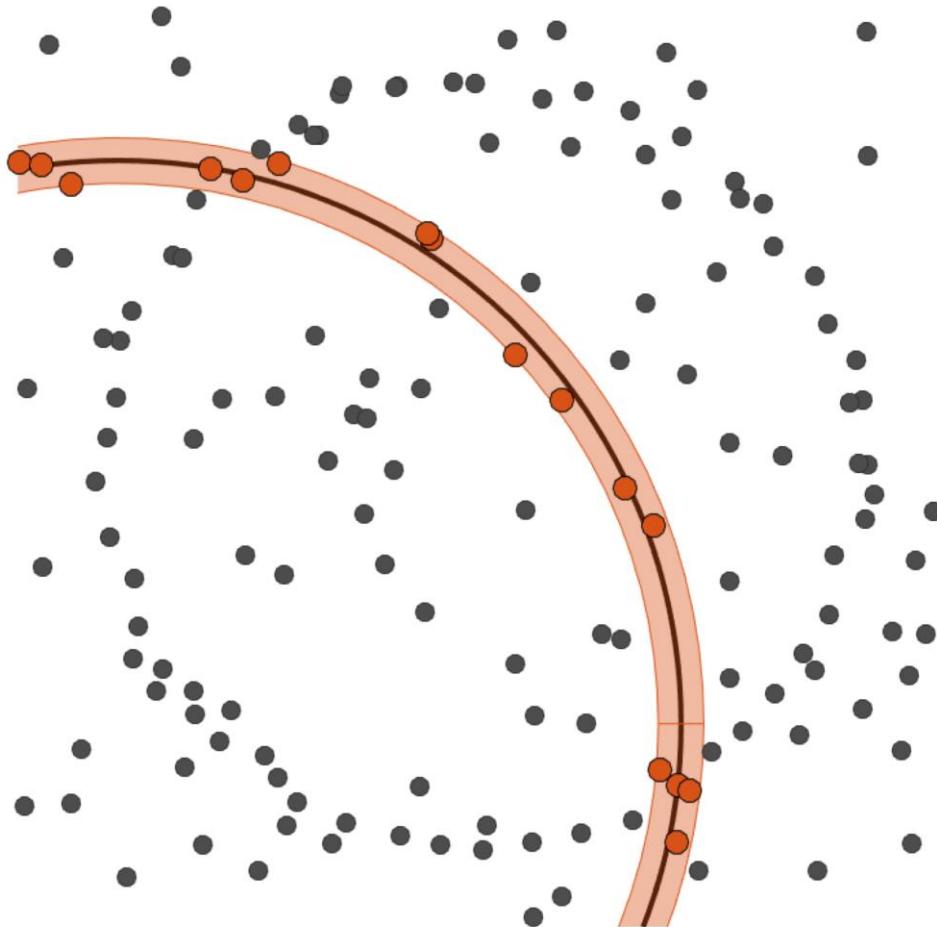
**end**

$k = k + 1$ ;

**until**  $k > k_{\max}$ ;

Optimize  $\theta^*$  on its inliers.

# Randomized Sample Consensus [Fischler and Bolles 1981]



**Input:**  $X$  data,  $\epsilon$  inlier threshold,  $k_{\max}$  max iteration

**Output:**  $\theta^*$  model estimate

$J^* = -\infty, k = 0;$

**repeat**

Select randomly a minimal sample set  $S \subset X$ ;

Estimate parameters  $\theta$  on  $S$ ;

Evaluate  $J(\theta) = \sum_{x \in X} \hat{f}_\epsilon(r(x, \theta))$ ;

**if**  $J(\theta) > J^*$  **then**

$\theta^* = \theta$ ;

$J^* = J(\theta)$ ;

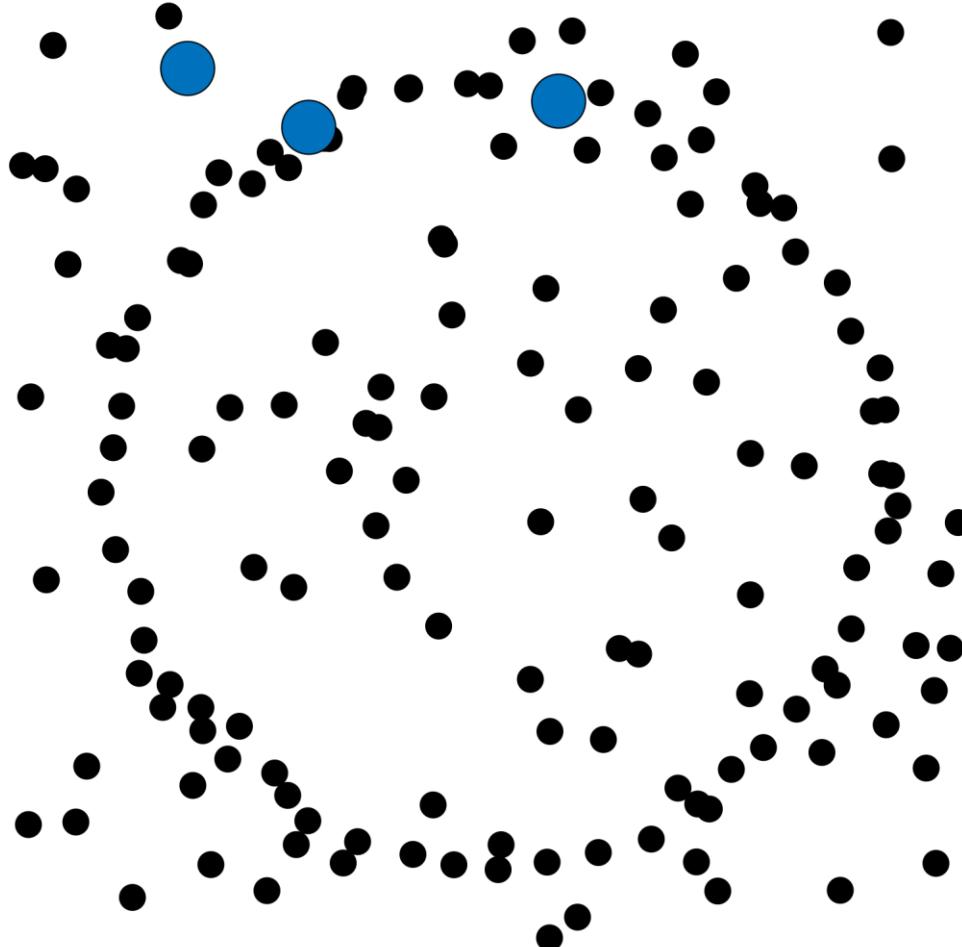
**end**

$k = k + 1$ ;

**until**  $k > k_{\max}$ ;

Optimize  $\theta^*$  on its inliers.

# Randomized Sample Consensus [Fischler and Bolles 1981]



**Input:**  $X$  data,  $\epsilon$  inlier threshold,  $k_{\max}$  max iteration

**Output:**  $\theta^*$  model estimate

$J^* = -\infty, k = 0;$

**repeat**

Select randomly a minimal sample set  $S \subset X$ ;

Estimate parameters  $\theta$  on  $S$ ;

Evaluate  $J(\theta) = \sum_{x \in S} \hat{f}_\epsilon(r(x, \theta))$ ;

**if**  $J(\theta) > J^*$  **then**

$\theta^* = \theta$ ;

$J^* = J(\theta)$ ;

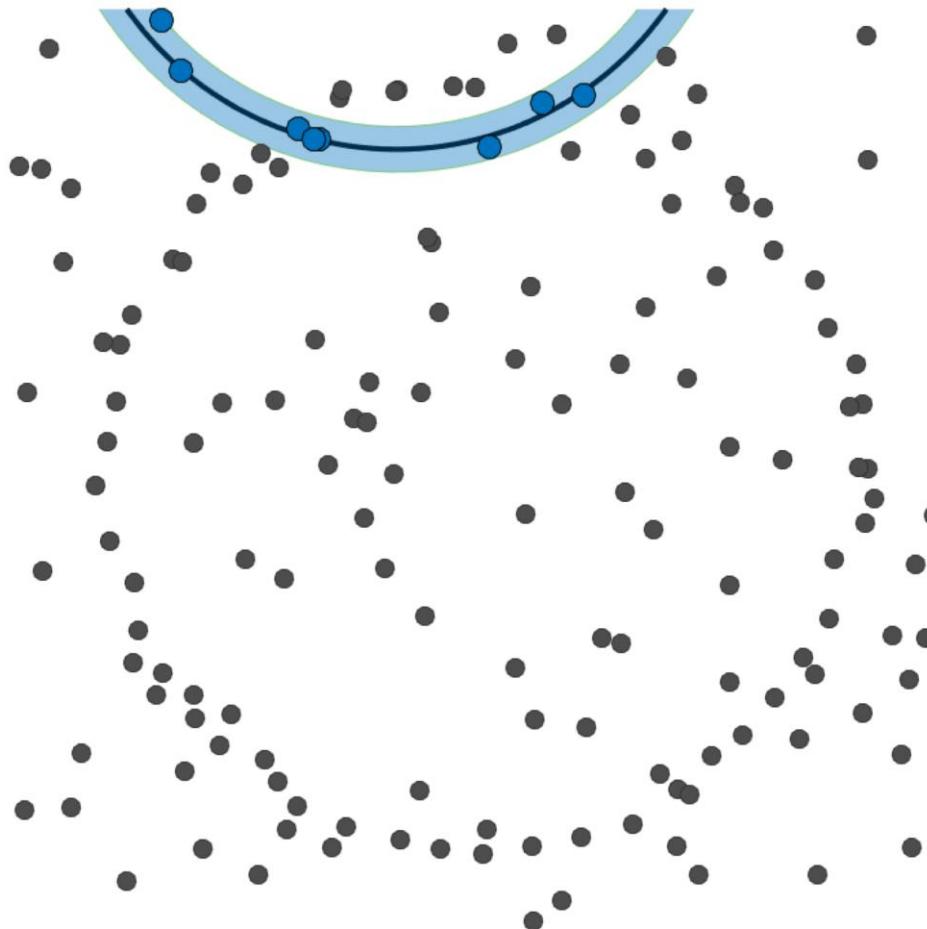
**end**

$k = k + 1$ ;

**until**  $k > k_{\max}$ ;

Optimize  $\theta^*$  on its inliers.

# Randomized Sample Consensus [Fischler and Bolles 1981]



**Input:**  $X$  data,  $\epsilon$  inlier threshold,  $k_{\max}$  max iteration

**Output:**  $\theta^*$  model estimate

$J^* = -\infty, k = 0;$

**repeat**

Select randomly a minimal sample set  $S \subset X$ ;

Estimate parameters  $\theta$  on  $S$ ;

Evaluate  $J(\theta) = \sum_{x \in X} \hat{f}_\epsilon(r(x, \theta))$ ;

**if**  $J(\theta) > J^*$  **then**

$\theta^* = \theta$ ;

$J^* = J(\theta)$ ;

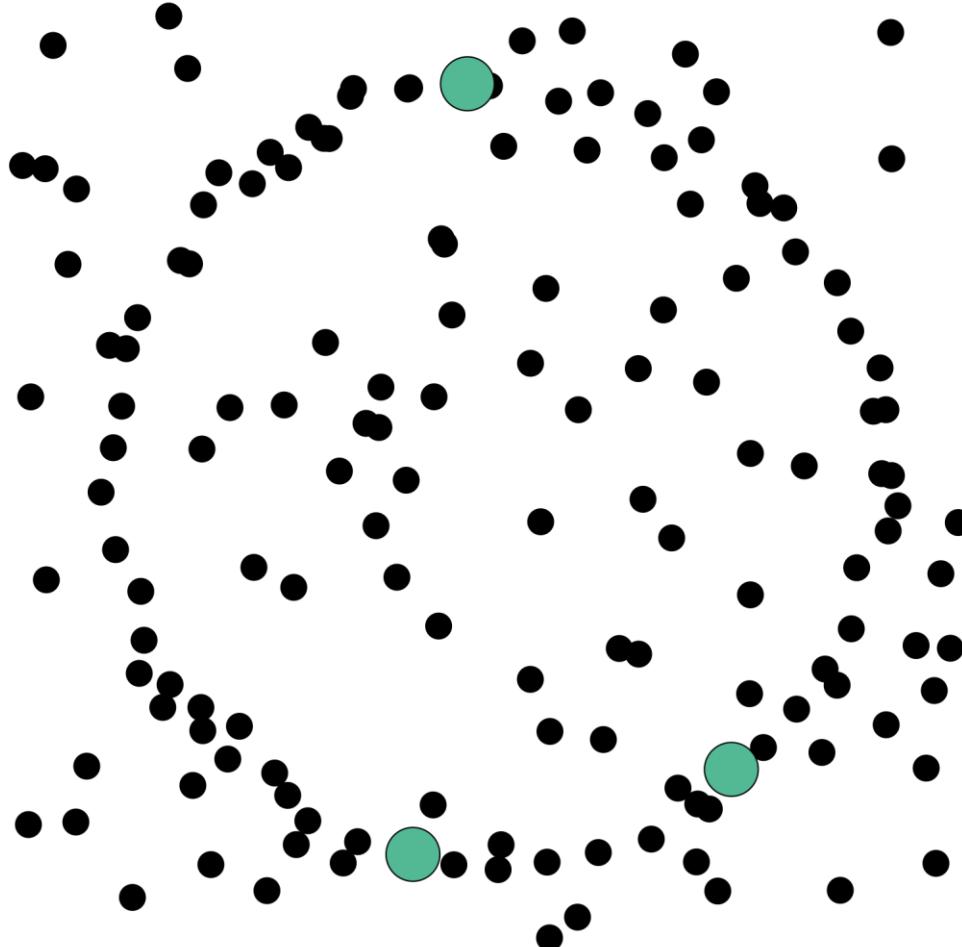
**end**

$k = k + 1$ ;

**until**  $k > k_{\max}$ ;

Optimize  $\theta^*$  on its inliers.

# Randomized Sample Consensus [Fischler and Bolles 1981]



**Input:**  $X$  data,  $\epsilon$  inlier threshold,  $k_{\max}$  max iteration

**Output:**  $\theta^*$  model estimate

$J^* = -\infty, k = 0;$

**repeat**

Select randomly a minimal sample set  $S \subset X$ ;

Estimate parameters  $\theta$  on  $S$ ;

Evaluate  $J(\theta) = \sum_{x \in X} \hat{f}_\epsilon(r(x, \theta))$ ;

**if**  $J(\theta) > J^*$  **then**

$\theta^* = \theta$ ;

$J^* = J(\theta)$ ;

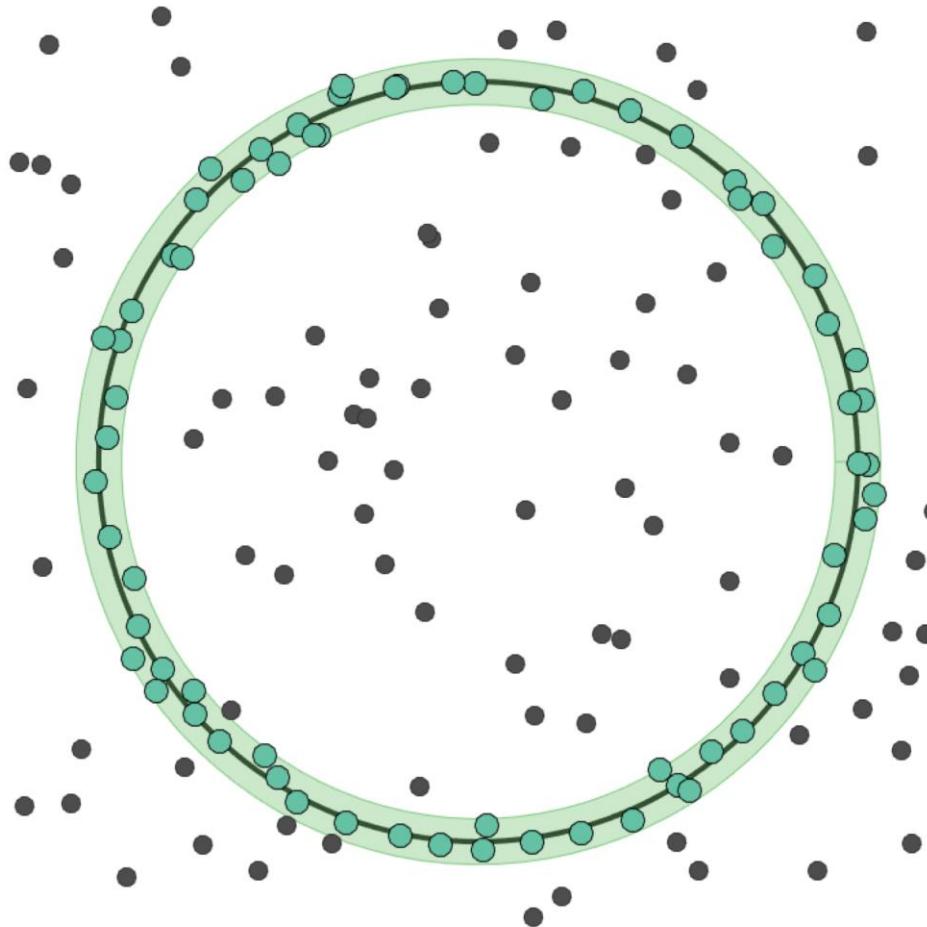
**end**

$k = k + 1$ ;

**until**  $k > k_{\max}$ ;

Optimize  $\theta^*$  on its inliers.

# Randomized Sample Consensus [Fischler and Bolles 1981]



**Input:**  $X$  data,  $\epsilon$  inlier threshold,  $k_{\max}$  max iteration

**Output:**  $\theta^*$  model estimate

$J^* = -\infty, k = 0;$

**repeat**

Select randomly a minimal sample set  $S \subset X$ ;

Estimate parameters  $\theta$  on  $S$ ;

Evaluate  $J(\theta) = \sum_{x \in X} \hat{f}_\epsilon(r(x, \theta))$ ;

**if**  $J(\theta) > J^*$  **then**

$\theta^* = \theta$ ;

$J^* = J(\theta)$ ;

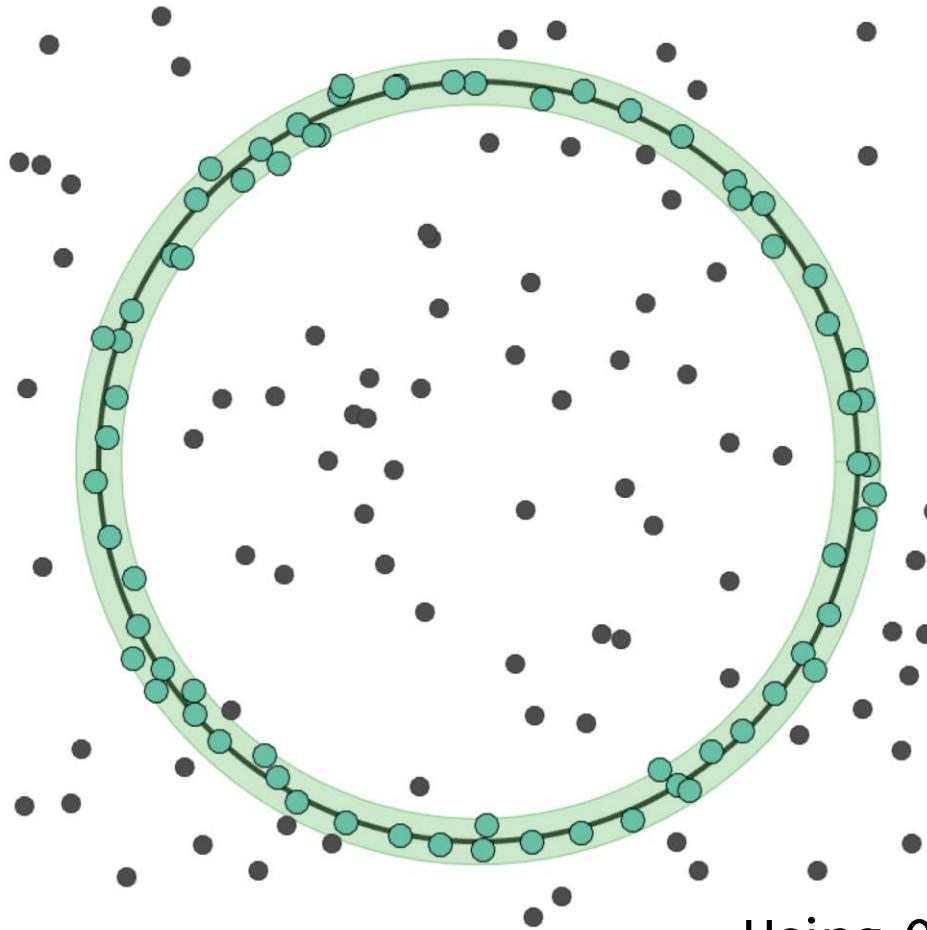
**end**

$k = k + 1$ ;

**until**  $k > k_{\max}$ ;

Optimize  $\theta^*$  on its inliers.

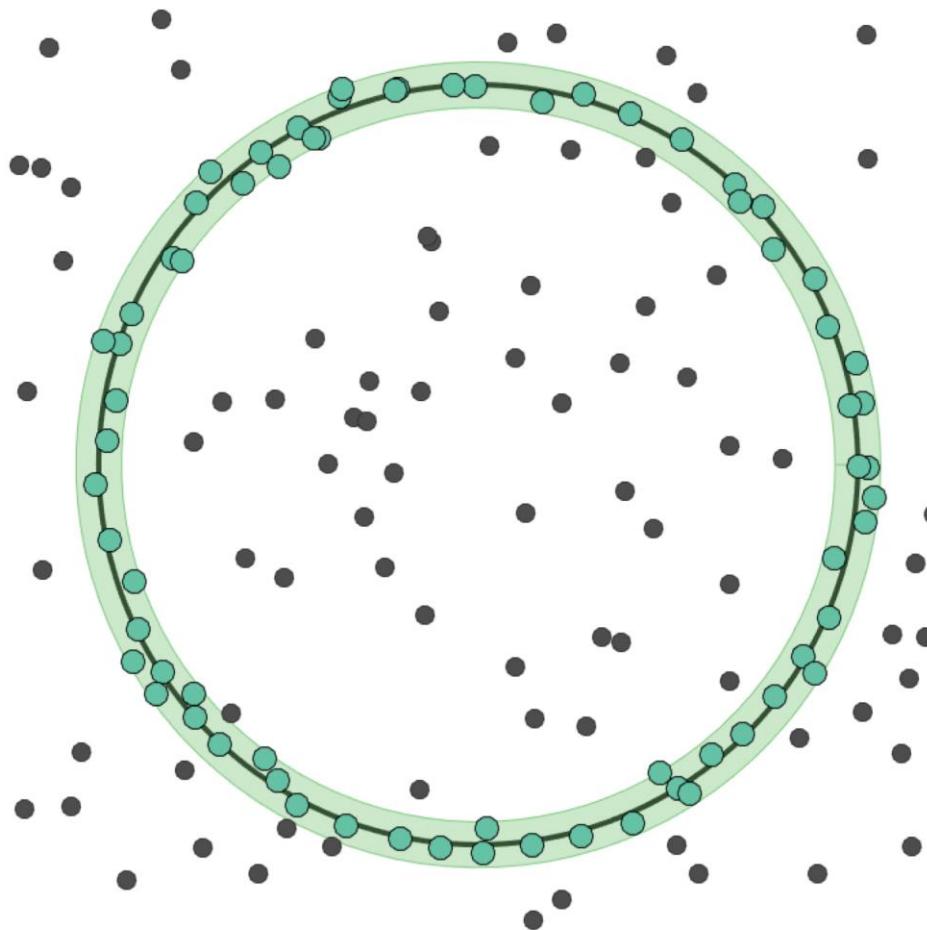
# Randomized Sample Consensus [Fischler and Bolles 1981]



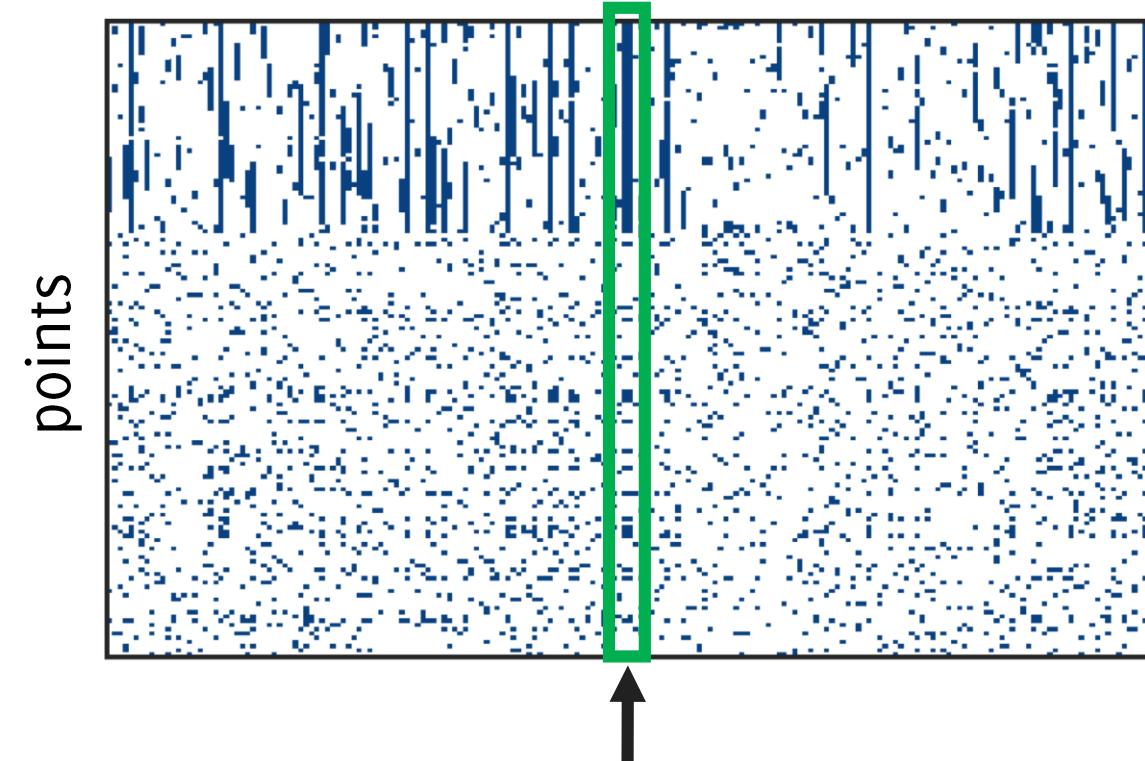
Using OLS,  
Increase stability of  
the results

**Input:**  $X$  data,  $\epsilon$  inlier threshold,  $k_{\max}$  max iteration  
**Output:**  $\theta^*$  model estimate  
 $J^* = -\infty, k = 0;$   
**repeat**  
    Select randomly a minimal sample set  $S \subset X$ ;  
    Estimate parameters  $\theta$  on  $S$ ;  
    Evaluate  $J(\theta) = \sum_{x \in X} \hat{f}_\epsilon(r(x, \theta))$ ;  
    **if**  $J(\theta) > J^*$  **then**  
         $\theta^* = \theta$ ;  
         $J^* = J(\theta)$ ;  
    **end**  
     $k = k + 1$ ;  
**until**  $k > k_{\max}$ ;  
Optimize  $\theta^*$  on its inliers.

# Randomized Sample Consensus

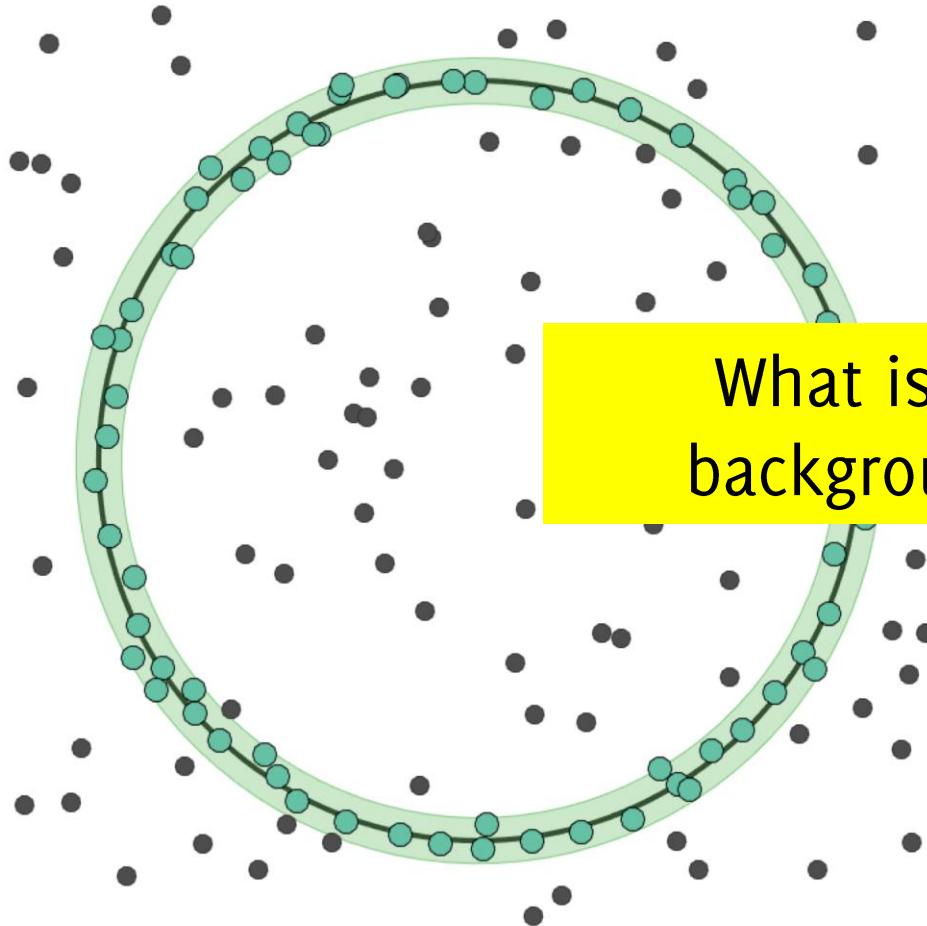


Data driven search of model space  
 $H = \{\theta_1, \theta_2, \dots, \theta_m\} \approx \Theta$



pick the column with the maximum sum

# Randomized Sample Consensus

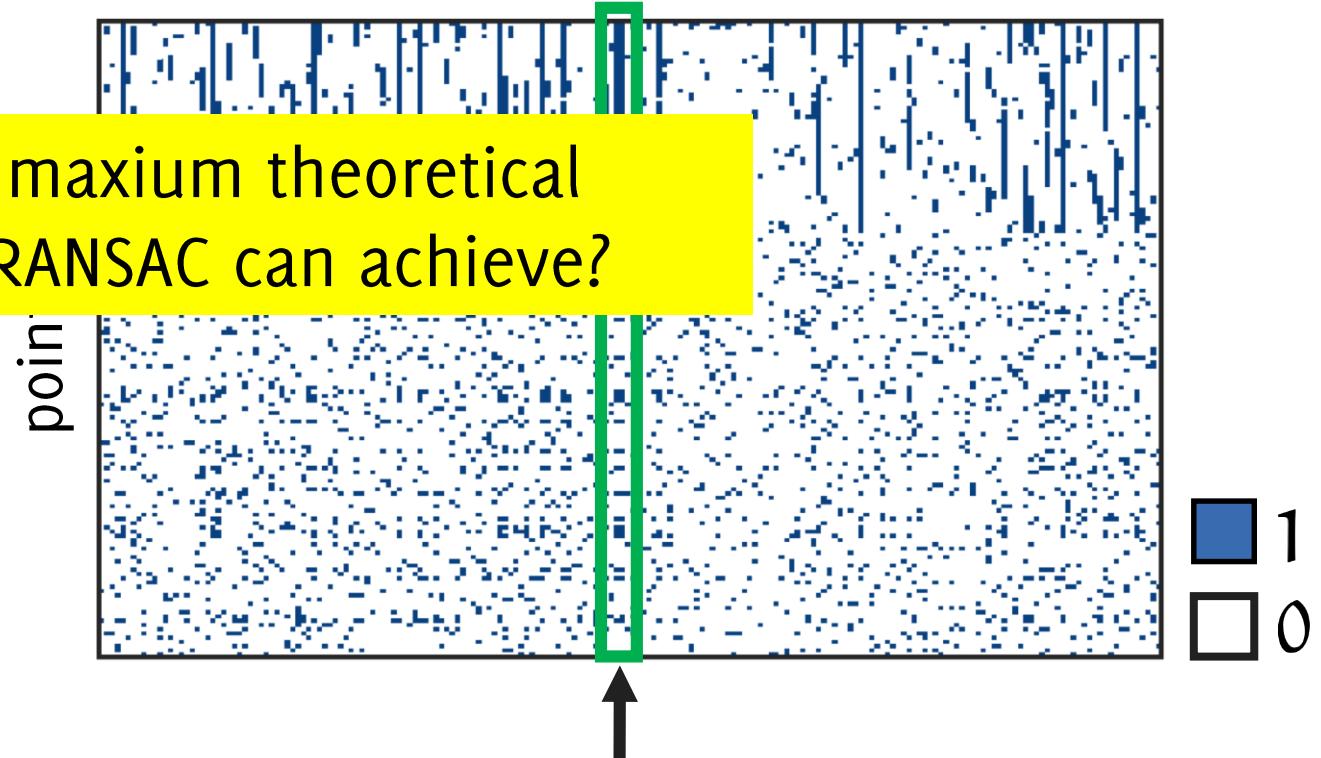


What is the maximum theoretical background RANSAC can achieve?

Data driven search of model space

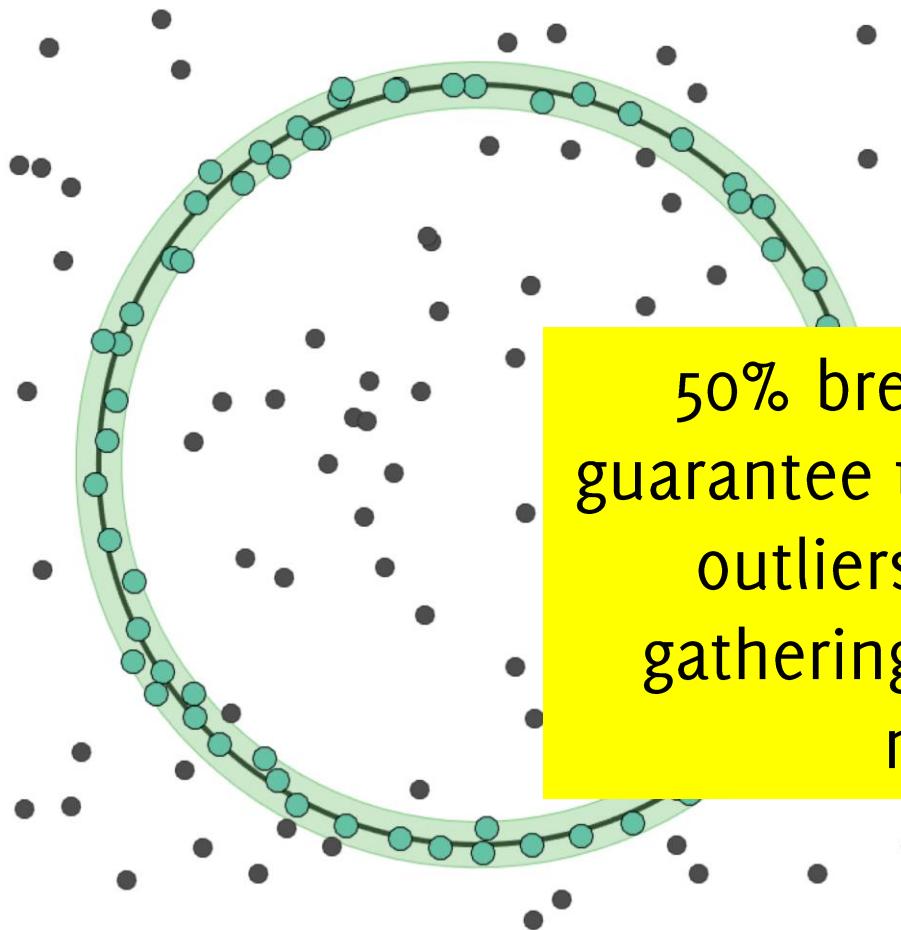
$$H = \{\theta_1, \theta_2, \dots, \theta_m\} \approx \Theta$$

tentative models



pick the column with the maximum sum

# Randomized Sample Consensus

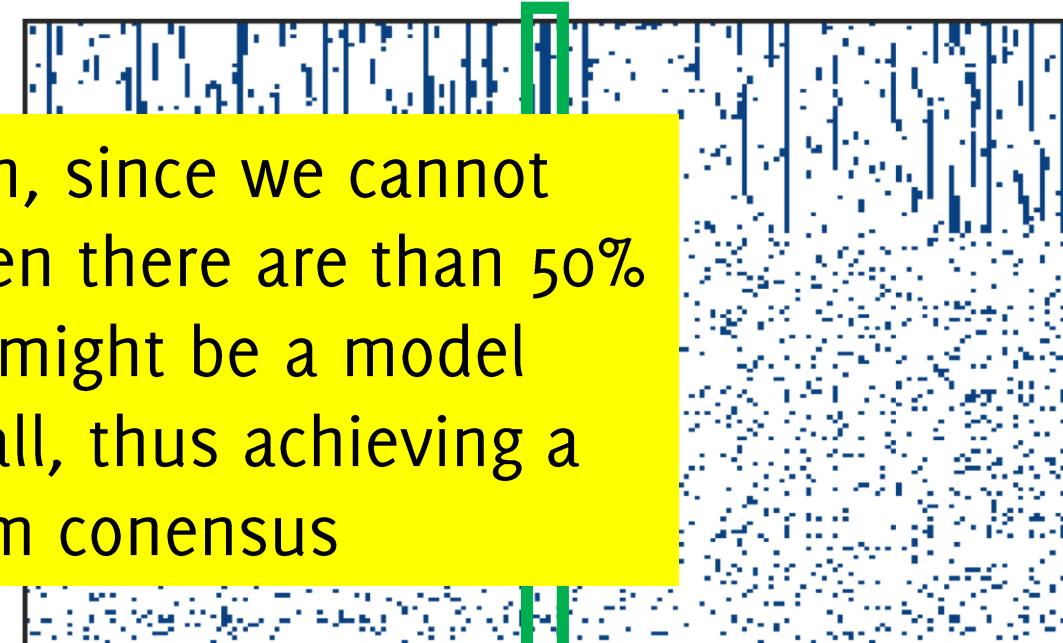


50% break-down, since we cannot guarantee that when there are than 50% outliers, there might be a model gathering them all, thus achieving a maximum consensus

Data driven search of model space

$$H = \{\theta_1, \theta_2, \dots, \theta_m\} \approx \Theta$$

tentative models



1

0

pick the column with the maximum sum

# Ransac: practical issues

**The size of the minimum sample  $S$ :** this is the bare minimum number of points to fit the parametric model at hand

**The inlier threshold  $\epsilon$ :** this can be estimated from the noise in the data

**The number of max iterations  $n$ :** the criteria for selecting the number of samples  $n$ : “Choose  $n$  so that, with probability  $p$ , at least one random sample is without outliers (e.g.  $p = 0.99$ ) “

# The maximum number of iterations $n$

Let  $e$  the probability of a sample to be an outlier and  $1 - e$  the probability of an inlier (can be estimated / provided by a-priori information)

The probability that all  $s$  points are inliers:  $(1 - e)^s$

The probability that at least one point in  $S$  is an outlier:  $1 - (1 - e)^s$   
(this is the probability for a sample  $S$  to yield the right model)

The probability that all the  $n$  selected set contain outliers

$$(1 - (1 - e)^s)^n$$

The probability that at least one the  $n$  set is without outliers:

$$1 - (1 - (1 - e)^s)^n$$

Set  $n$  to have the above probability below a parameter  $p$

$$p = (1 - (1 - (1 - e)^s)^n) \rightarrow n = \log(1 - p) / \log(1 - (1 - e)^s)$$

# Ransac: practical issues

Choose  $n$  so that, with probability  $p$ , at least one random sample is free from outliers (e.g.  $p = 0.99$ ) (outlier ratio:  $e$ )

$$n = \log(1 - p) / \log(1 - (1 - e)^s)$$

s	proportion of outliers $e$							
	5%	10%	20%	25%	30%	40%	50%	
2	2	3	5	6	7	11	17	
3	3	4	7	9	11	19	35	
4	3	5	9	13	17	34	72	
5	4	6	12	17	26	57	146	
6	4	7	16	24	37	97	293	
7	4	8	20	33	54	163	588	
8	5	9	26	44	78	272	1177	

# Ransac: details

Repeat  $n$  times:

- Draw  $s$  points uniformly at random
- Fit line to these  $s$  points
- Find inliers to this line among the remaining points (i.e., points whose distance from the line is less than  $t$ )
- Update  $n$ 
  - $e = 1 - \frac{\text{number of inliers}}{\text{number of points}}$
  - $n = \log(1 - p) / \log(1 - (1 - e)^s)$

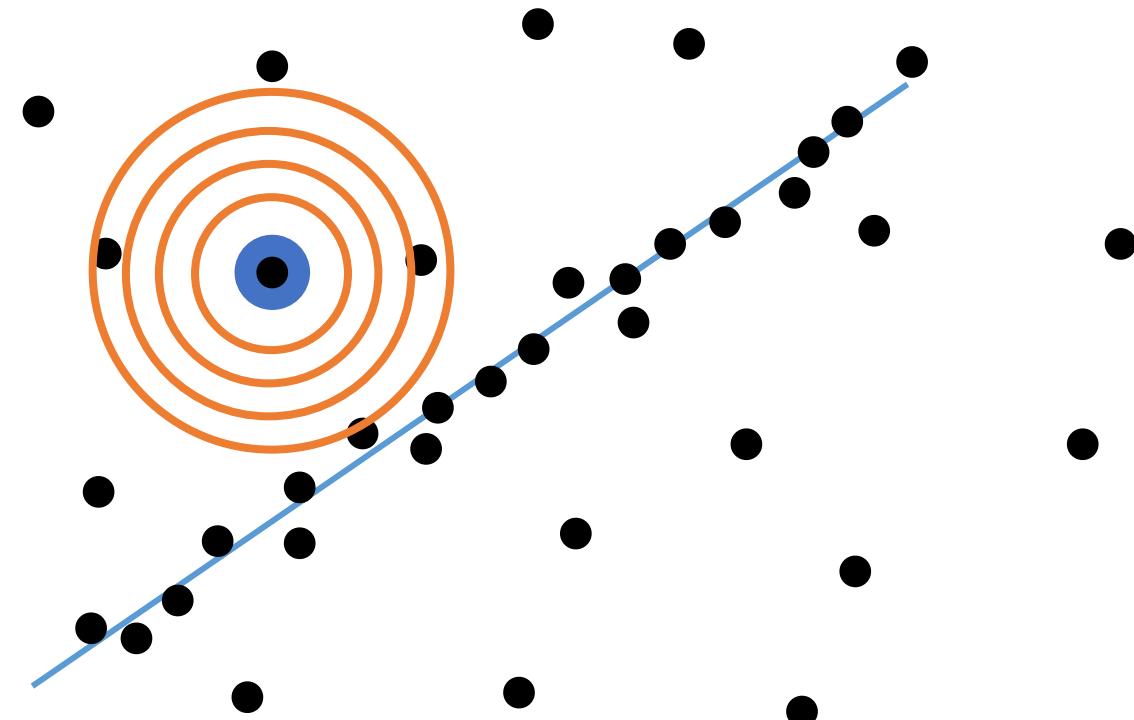
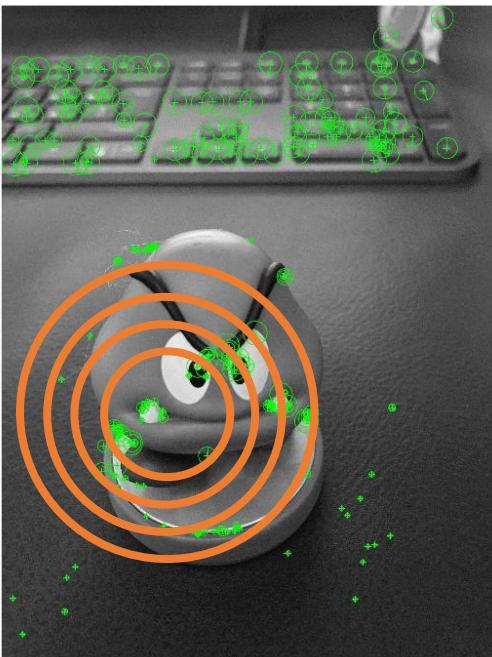
Choose the best model

Re-estimate the line with the inliers only through ordinary least square

# The sampling strategy

In some applications inliers tend to be localized:

- draw uniformly a first point
- draw the other samples from the neighborhood of the initial point



Barath, Noskova, Ivashechkin, Matas: MAGSAC++, a fast, reliable and accurate robust estimator. CVPR 2020

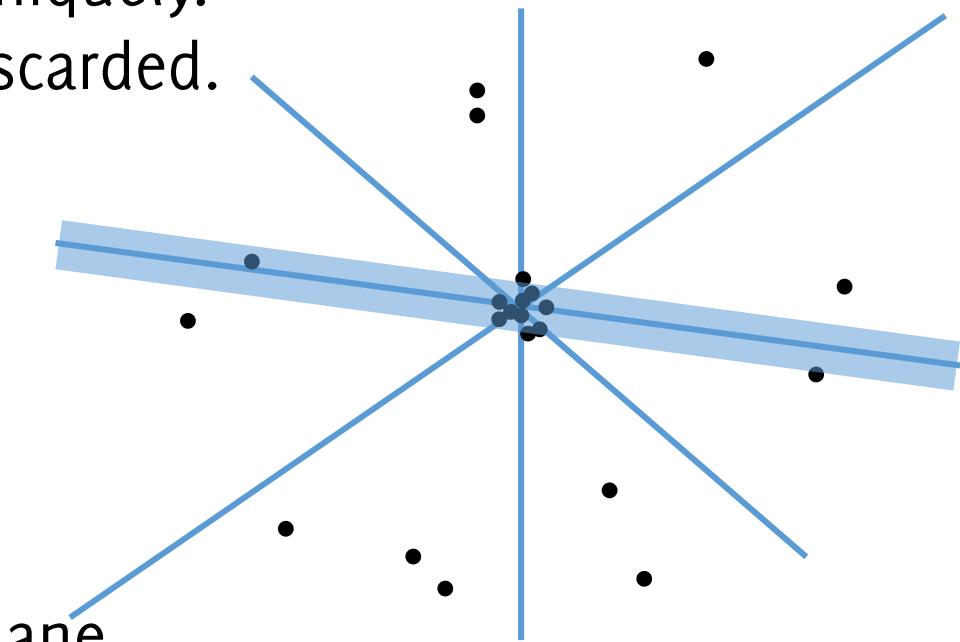


Daniel Barath and Jiri Matas, Graph-Cut RANSAC; CVPR 2018

Luca Magri

# Degeneracy check

A degenerate configuration is a configuration of points in the minimum sample set for which it is not possible to define the model uniquely. Typically, has higher random support and must be discarded.



## Examples:

For fundamental matrix, planar points or dominant plane.

For homography, collinear points.

When using SVD, when the rank of the coefficient matrix drops with respect the desired dimension.

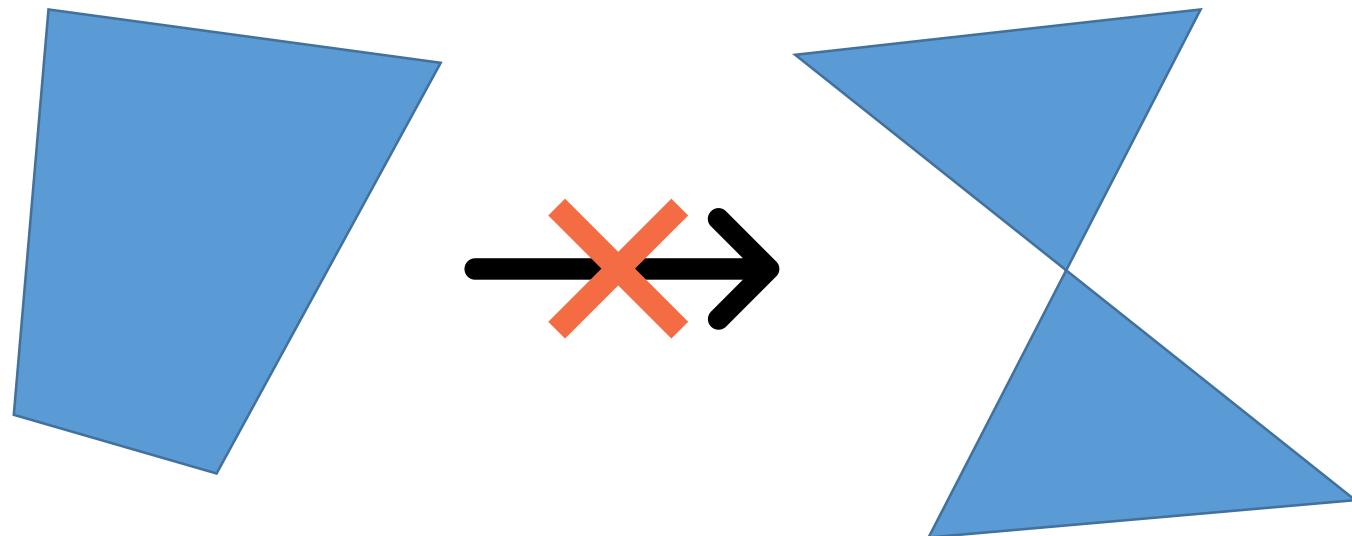


# Sample checks

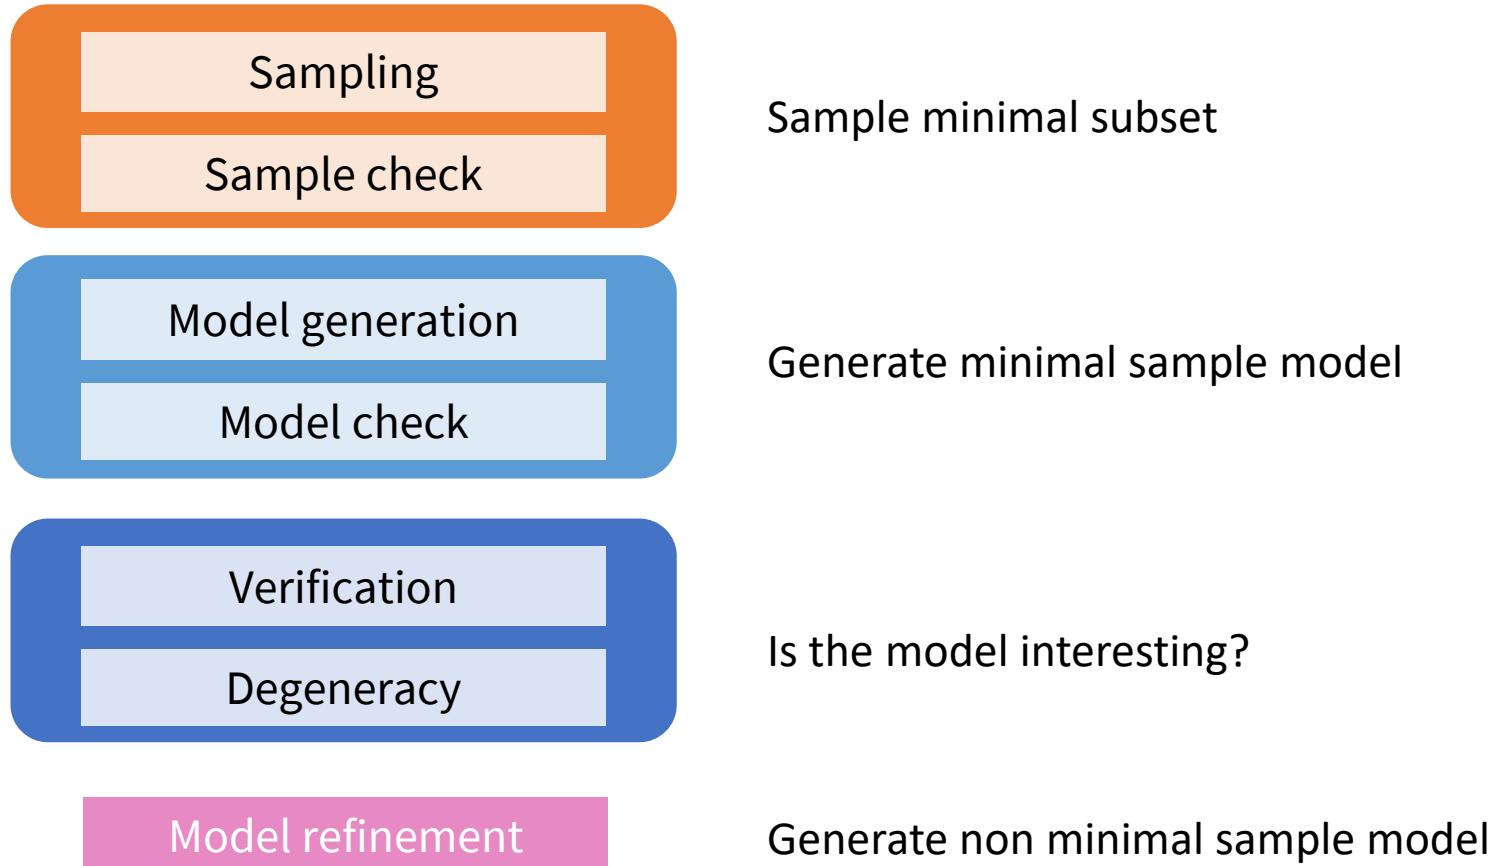
Before the verification step, minimal sample set can be checked:

Cameras can only observe points in front of them (remember chirality constraints used for the essential matrix ).

Homography transformation need to preserve convexity:



# RanSaC implementation with bells and whistles



Raguram, Chum, Pollefeys, Matas, Frham Usac a Universal Framework for Random  
Sample Consensus. TPAMI 2013

# Ransac

## Pros:

- very popular (>22900 citations in Google Scholar)
- many improvements have been proposed
- very versatile
- agnostic on outlier percentage
- mild assumption: know the scale noise  
to set the inlier threshold

## Cons:

- can take longer than expected

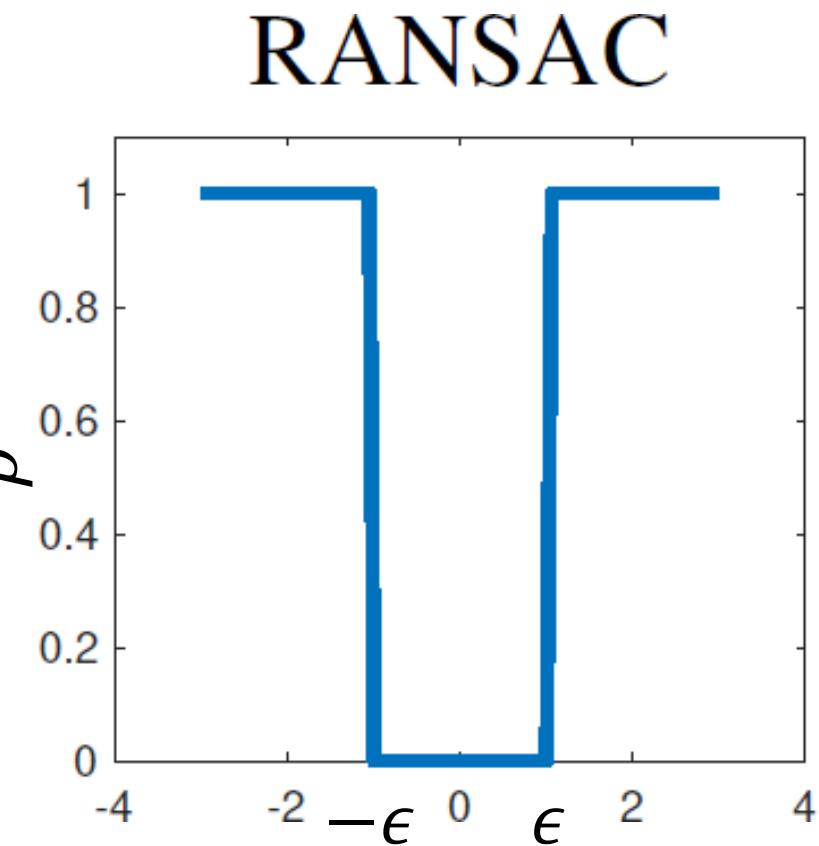
# Ransac as M-estimator

(Steward 1999) RanSaC can be seen as a particular M-estimator since the **loss it minimizes** is the number of points having residual above the inlier threshold  $\epsilon$

$$f(r_i) = \begin{cases} 1, & r_i > \epsilon \\ 0, & r_i \leq \epsilon \end{cases}$$

Of course selecting inlier threshold  $\epsilon$  is very critical

Ransac achieves a theoretical breakdown of 50% of outliers, but in practice, provided a good selection of  $\epsilon$ , this can be even higher

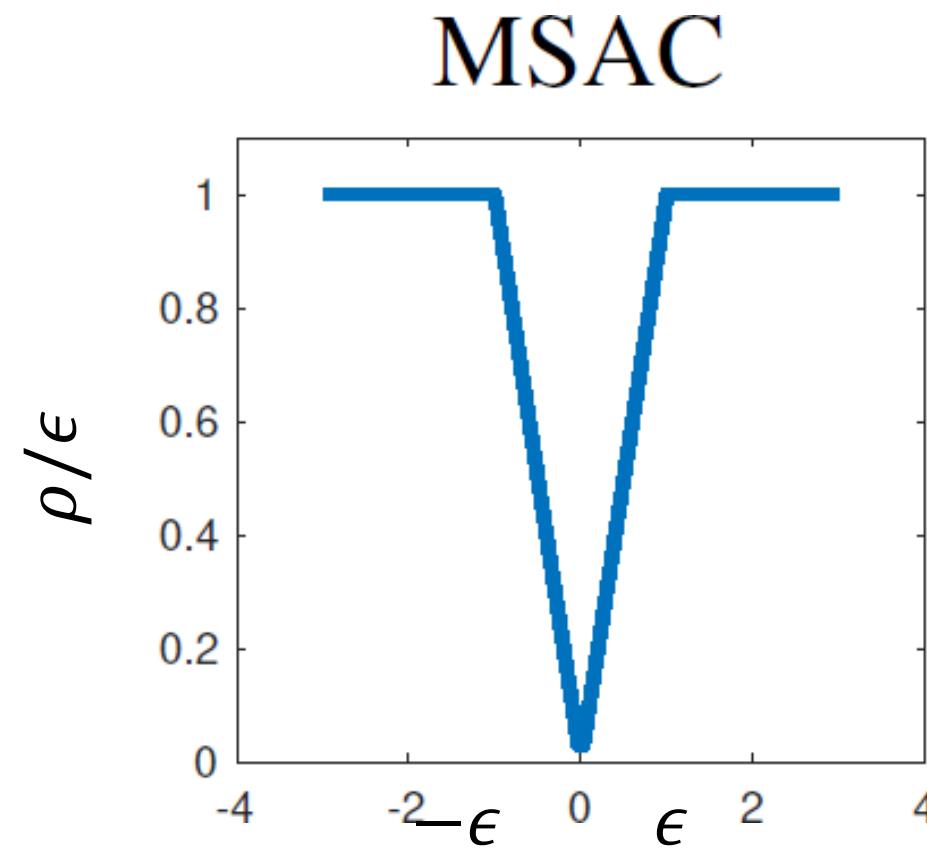


# MSAC

(Torr and Zisserman 2000) a different loss function to be minimized within the RanSaC framework

$$f(r_i) = \begin{cases} r_i, & r_i > \epsilon \\ \epsilon, & r_i \leq \epsilon \end{cases}$$

This turns to be more effective and should be preferred to RanSaC



# Ransac vs MSaC

## Ransac

**Input:**  $X$  data,  $\epsilon$  inlier threshold,  $k_{\max}$  max iteration

**Output:**  $\theta^*$  model estimate

$J^* = -\infty, k = 0;$

**repeat**

Select randomly a minimal sample set  $S \subset X$ ;

Estimate parameters  $\theta$  on  $S$ ;

$$\text{Evaluate } J(\theta) = \sum_{x \in X} \hat{f}_\epsilon(r(x, \theta));$$

**if**  $J(\theta) > J^*$  **then**

$\theta^* = \theta$ ;

$J^* = J(\theta)$ ;

**end**

$k = k + 1$ ;

**until**  $k > k_{\max}$ ;

Optimize  $\theta^*$  on its inliers.

## MSaC

**Input:**  $X$  data,  $\epsilon$  inlier threshold,  $k_{\max}$  max iteration

**Output:**  $\theta^*$  model estimate

$J^* = +\infty, k = 0;$

**repeat**

Select randomly a minimal sample set  $S \subset X$ ;

Estimate parameters  $\theta$  on  $S$ ;

Estimate inlier set  $I = \{x \in X : r(x, \theta)^2 < \epsilon^2\}$ ;

$$\text{Evaluate } J(\theta) = \sum_{x \in I} r(x, \theta) + (|X| - |I|)\epsilon;$$

**if**  $J(\theta) < J^*$  **then**

$\theta^* = \theta$ ;

$J^* = J(\theta)$ ;

**end**

$k = k + 1$ ;

**until**  $k > k_{\max}$ ;

Optimize  $\theta^*$  on its inliers.

# Video Analytics

**math & sport**



**POLITECNICO  
MILANO 1863**

DIPARTIMENTO DI ELETTRONICA  
INFORMAZIONE E BIOINGEGNERIA



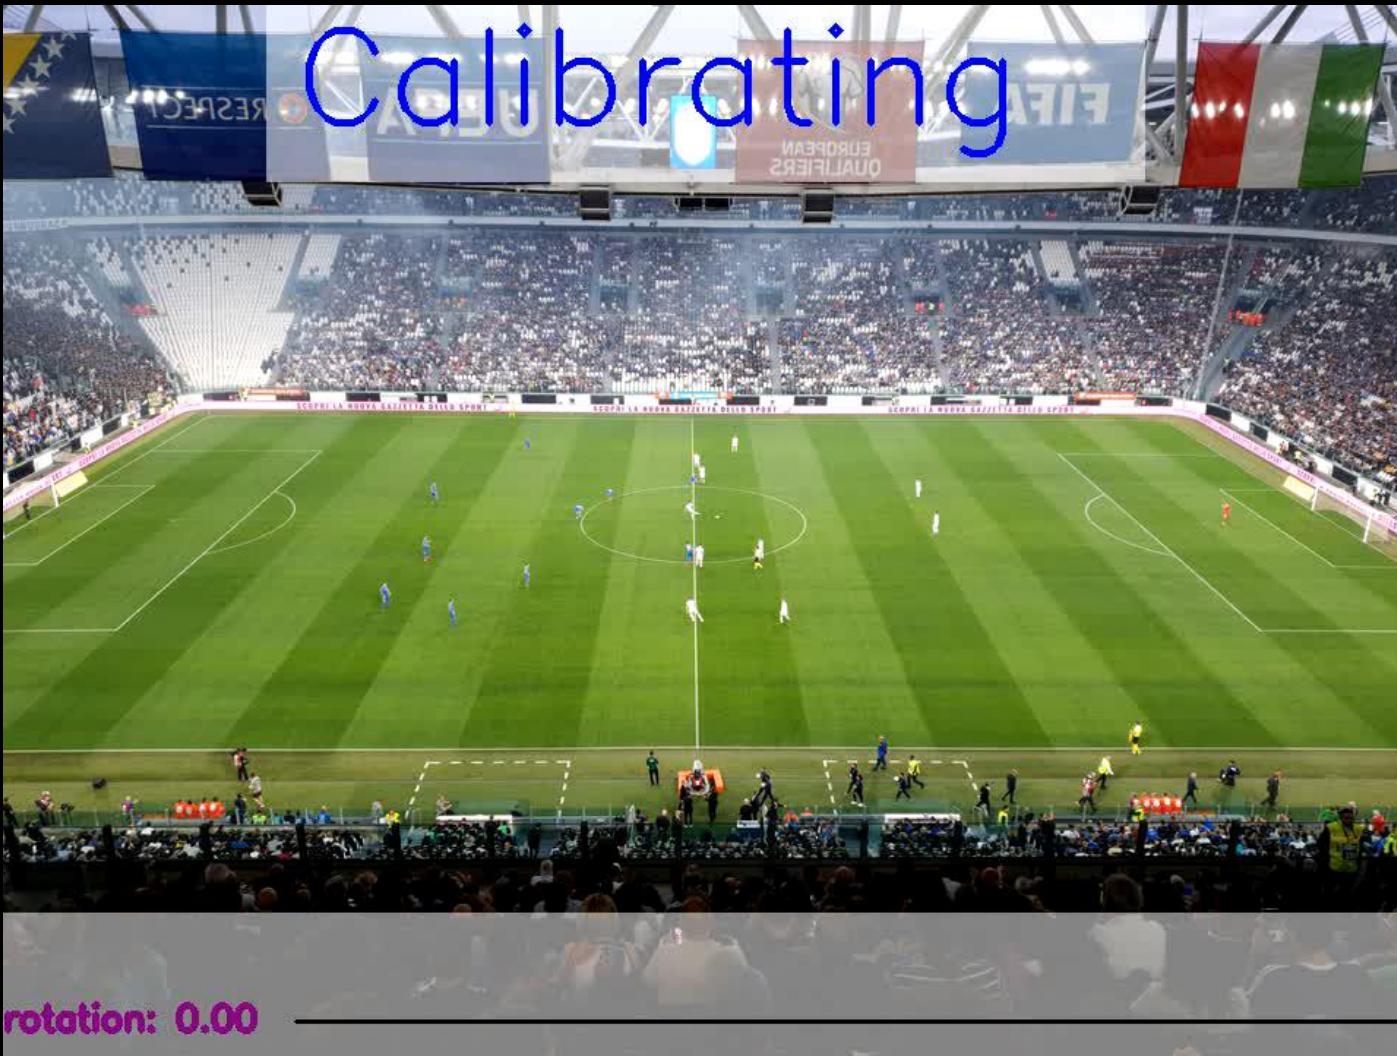
# math & sport



Powered by



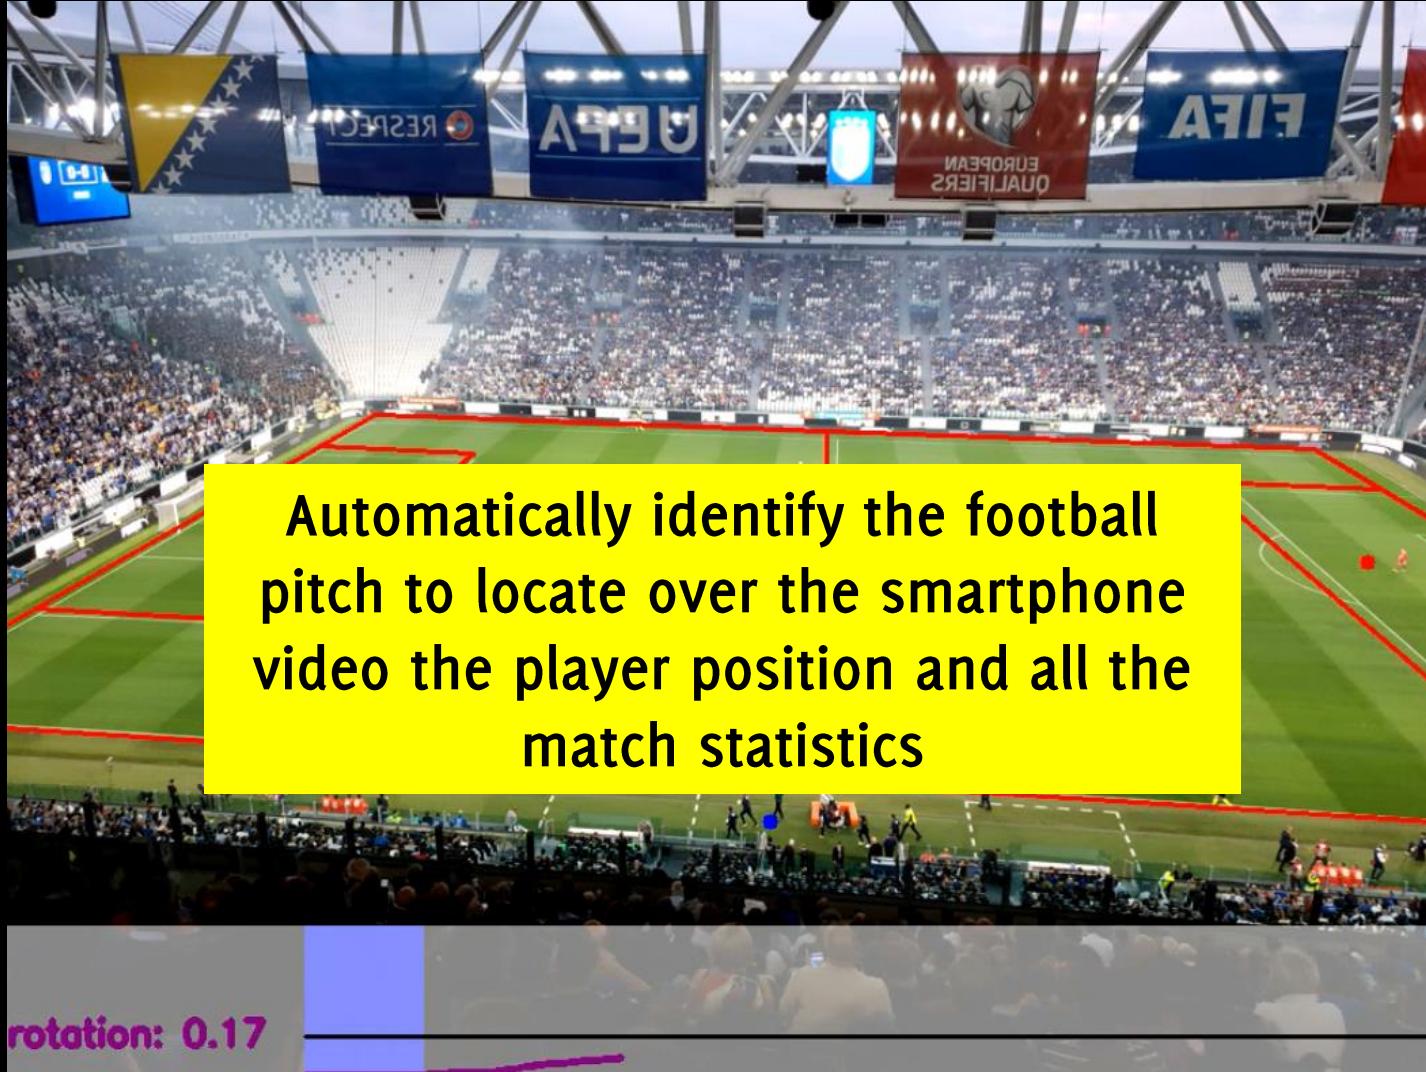
# Calibrating



rotation: 0.00

Powered by

# math & sport

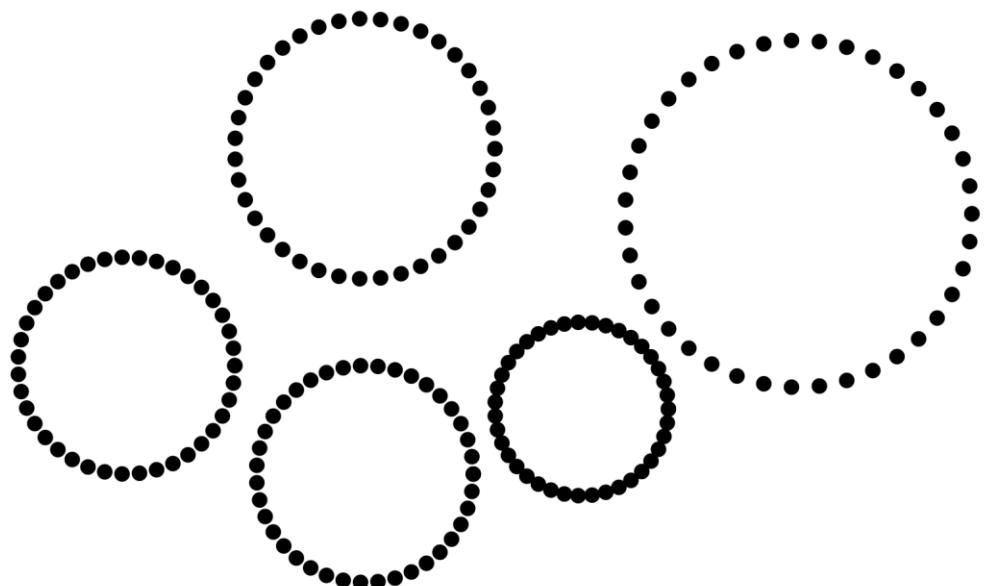


Powered by

# Multi-Model Fitting

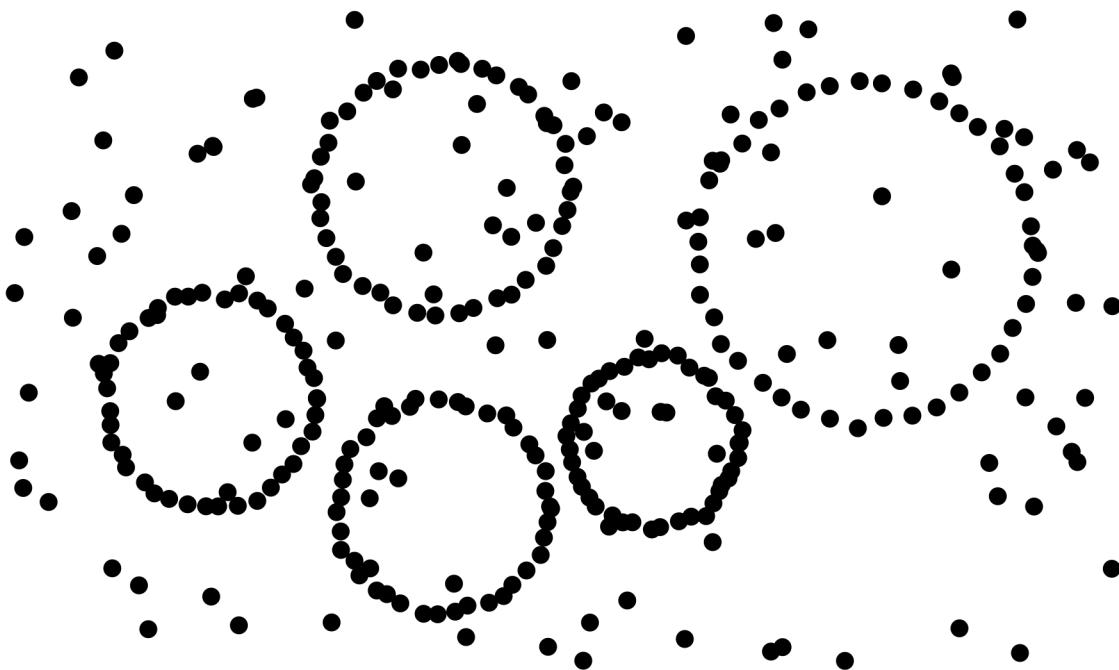
# The problem of multi-model fitting (or structure recovery)

Given a set of data  $X = \{x_1, \dots, x_N\} \subset \mathbb{R}^d$ , possibly corrupted by noise and outliers, and a family of geometric models  $\Theta$



# The problem of multi-model fitting (or structure recovery)

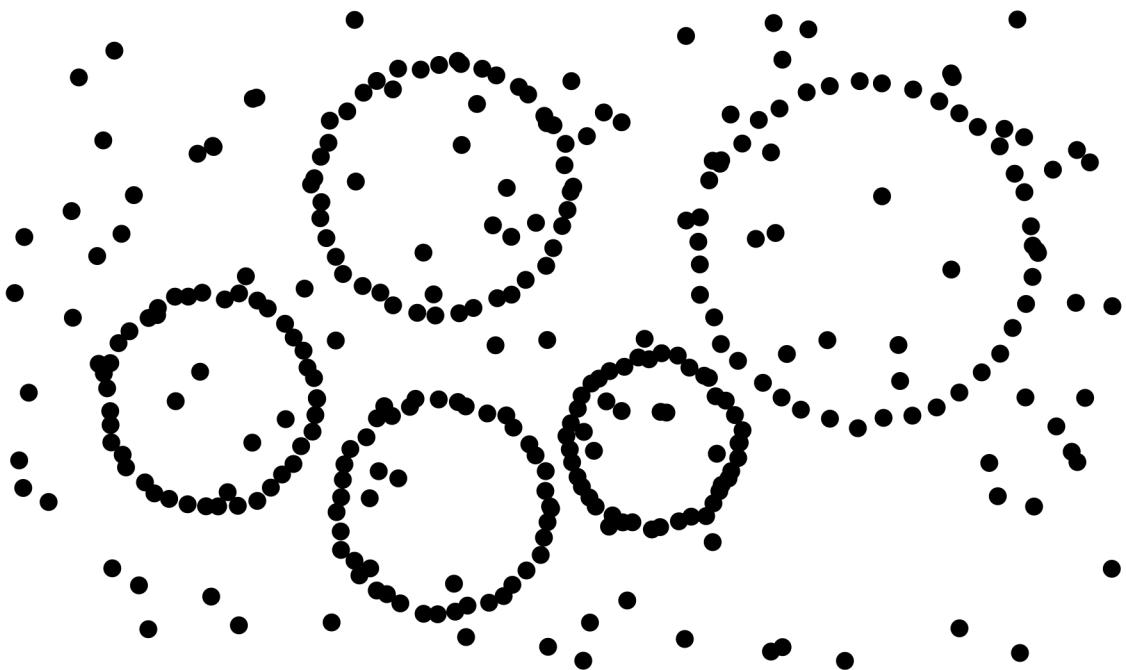
Given a set of data  $X = \{x_1, \dots, x_N\} \subset \mathbb{R}^d$ , possibly corrupted by noise and outliers, and a family of geometric models  $\Theta$



# The problem of multi-model fitting (or structure recovery)

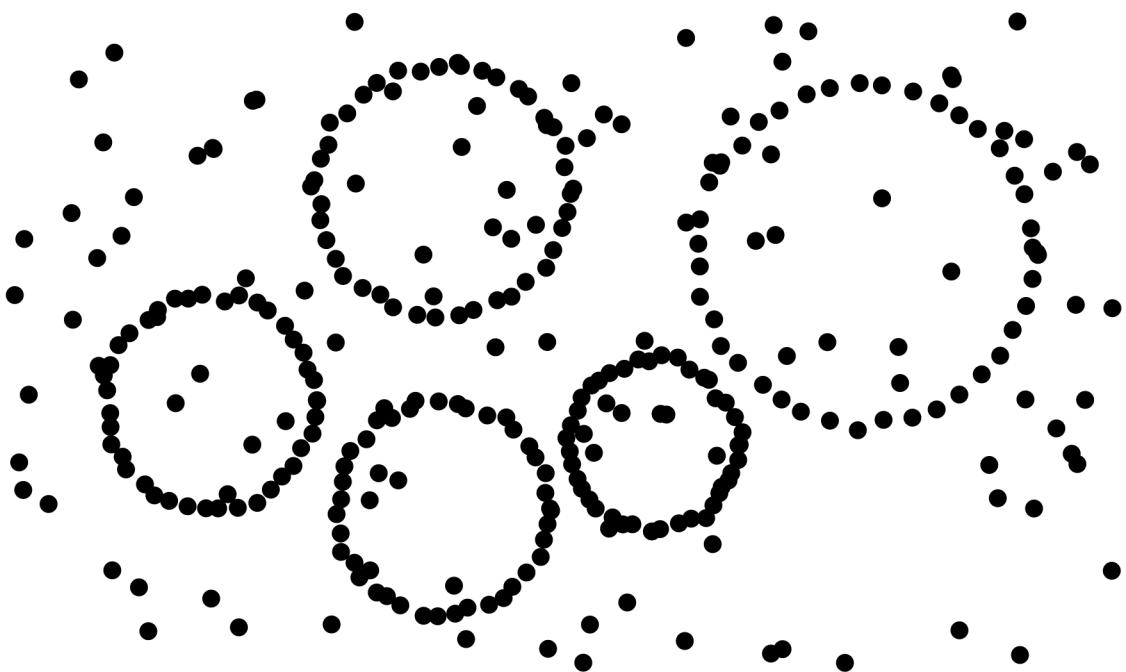
**Given** a set of data  $X = \{x_1, \dots, x_N\} \subset \mathbb{R}^d$ , possibly corrupted by noise and outliers, and a family of geometric models  $\Theta$

**Goal:** automatically estimate the models that best explain the data/discover the structures hidden in the data

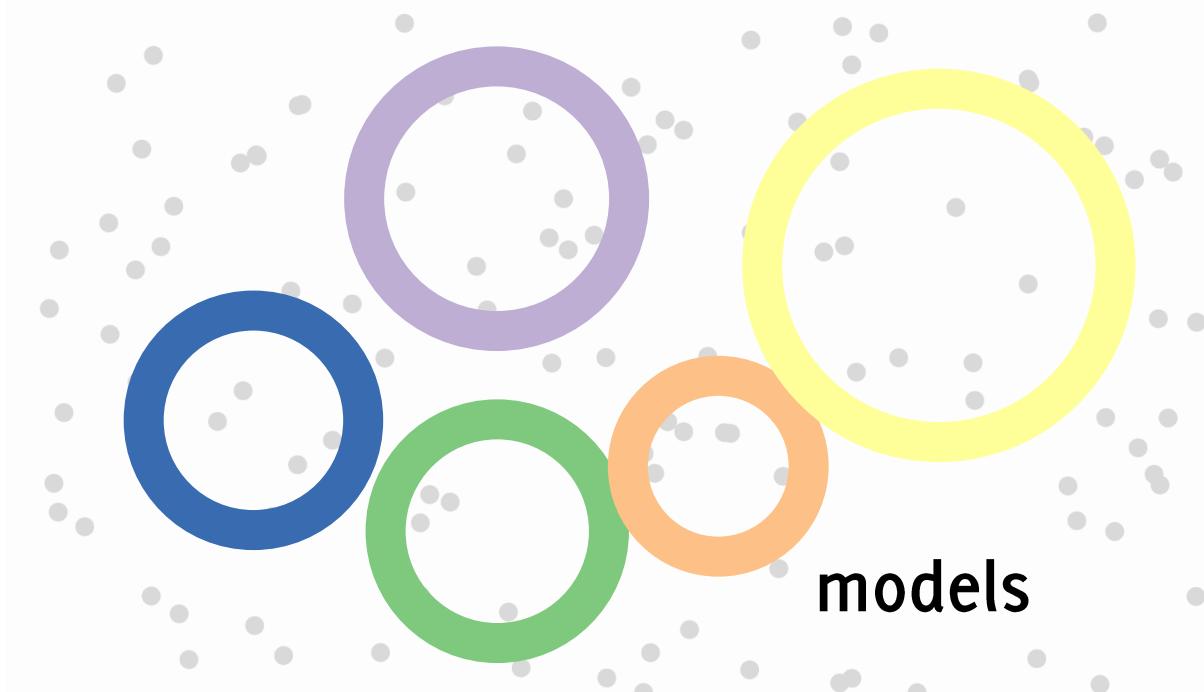


# The problem of multi-model fitting (or structure recovery)

**Given** a set of data  $X = \{x_1, \dots, x_N\} \subset \mathbb{R}^d$ , possibly corrupted by noise and outliers, and a family of geometric models  $\Theta$



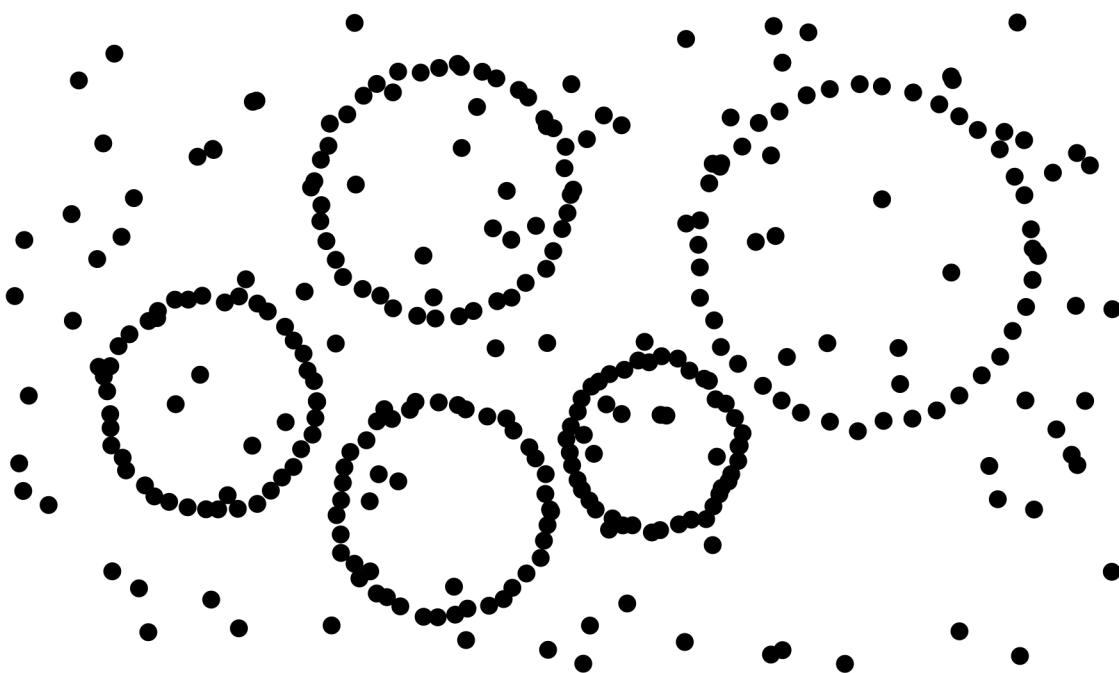
**Goal:** automatically estimate the models that best explain the data/discover the structures hidden in the data



“in the eye of the beholder”, mathematical descriptions of the data that an observer fits

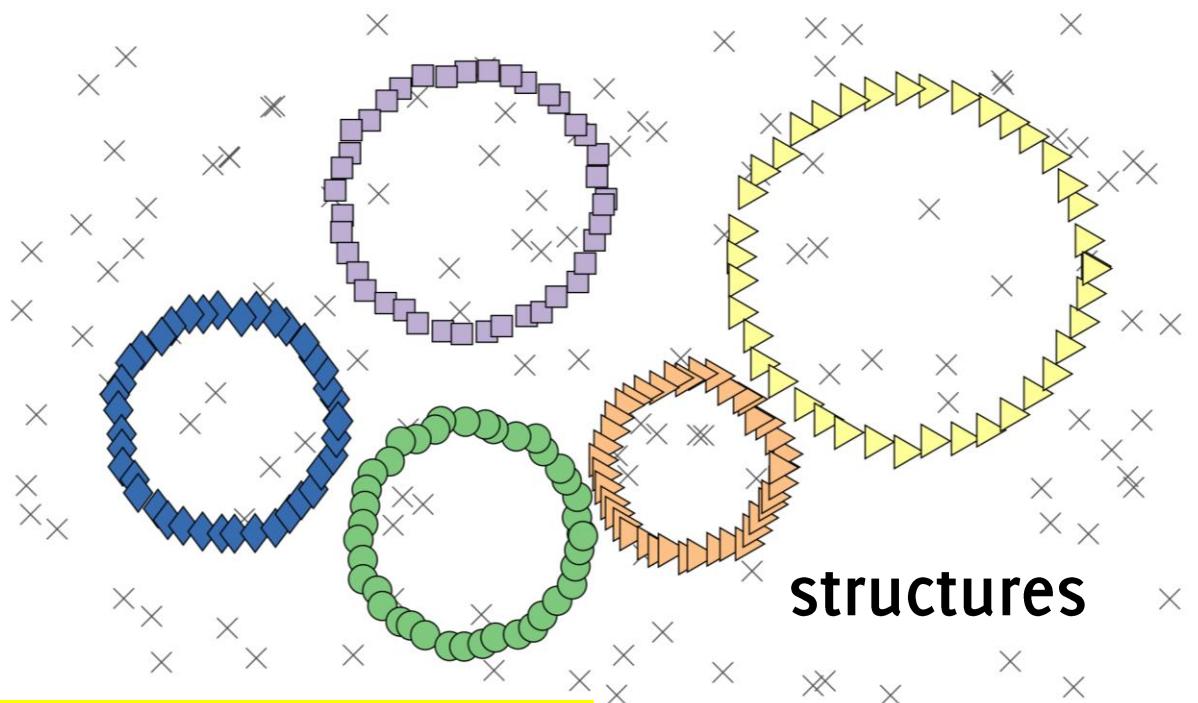
# The problem of multi-model fitting (or structure recovery)

Given a set of data  $X = \{x_1, \dots, x_N\} \subset \mathbb{R}^d$ , possibly corrupted by noise and outliers, and a family of geometric models  $\Theta$

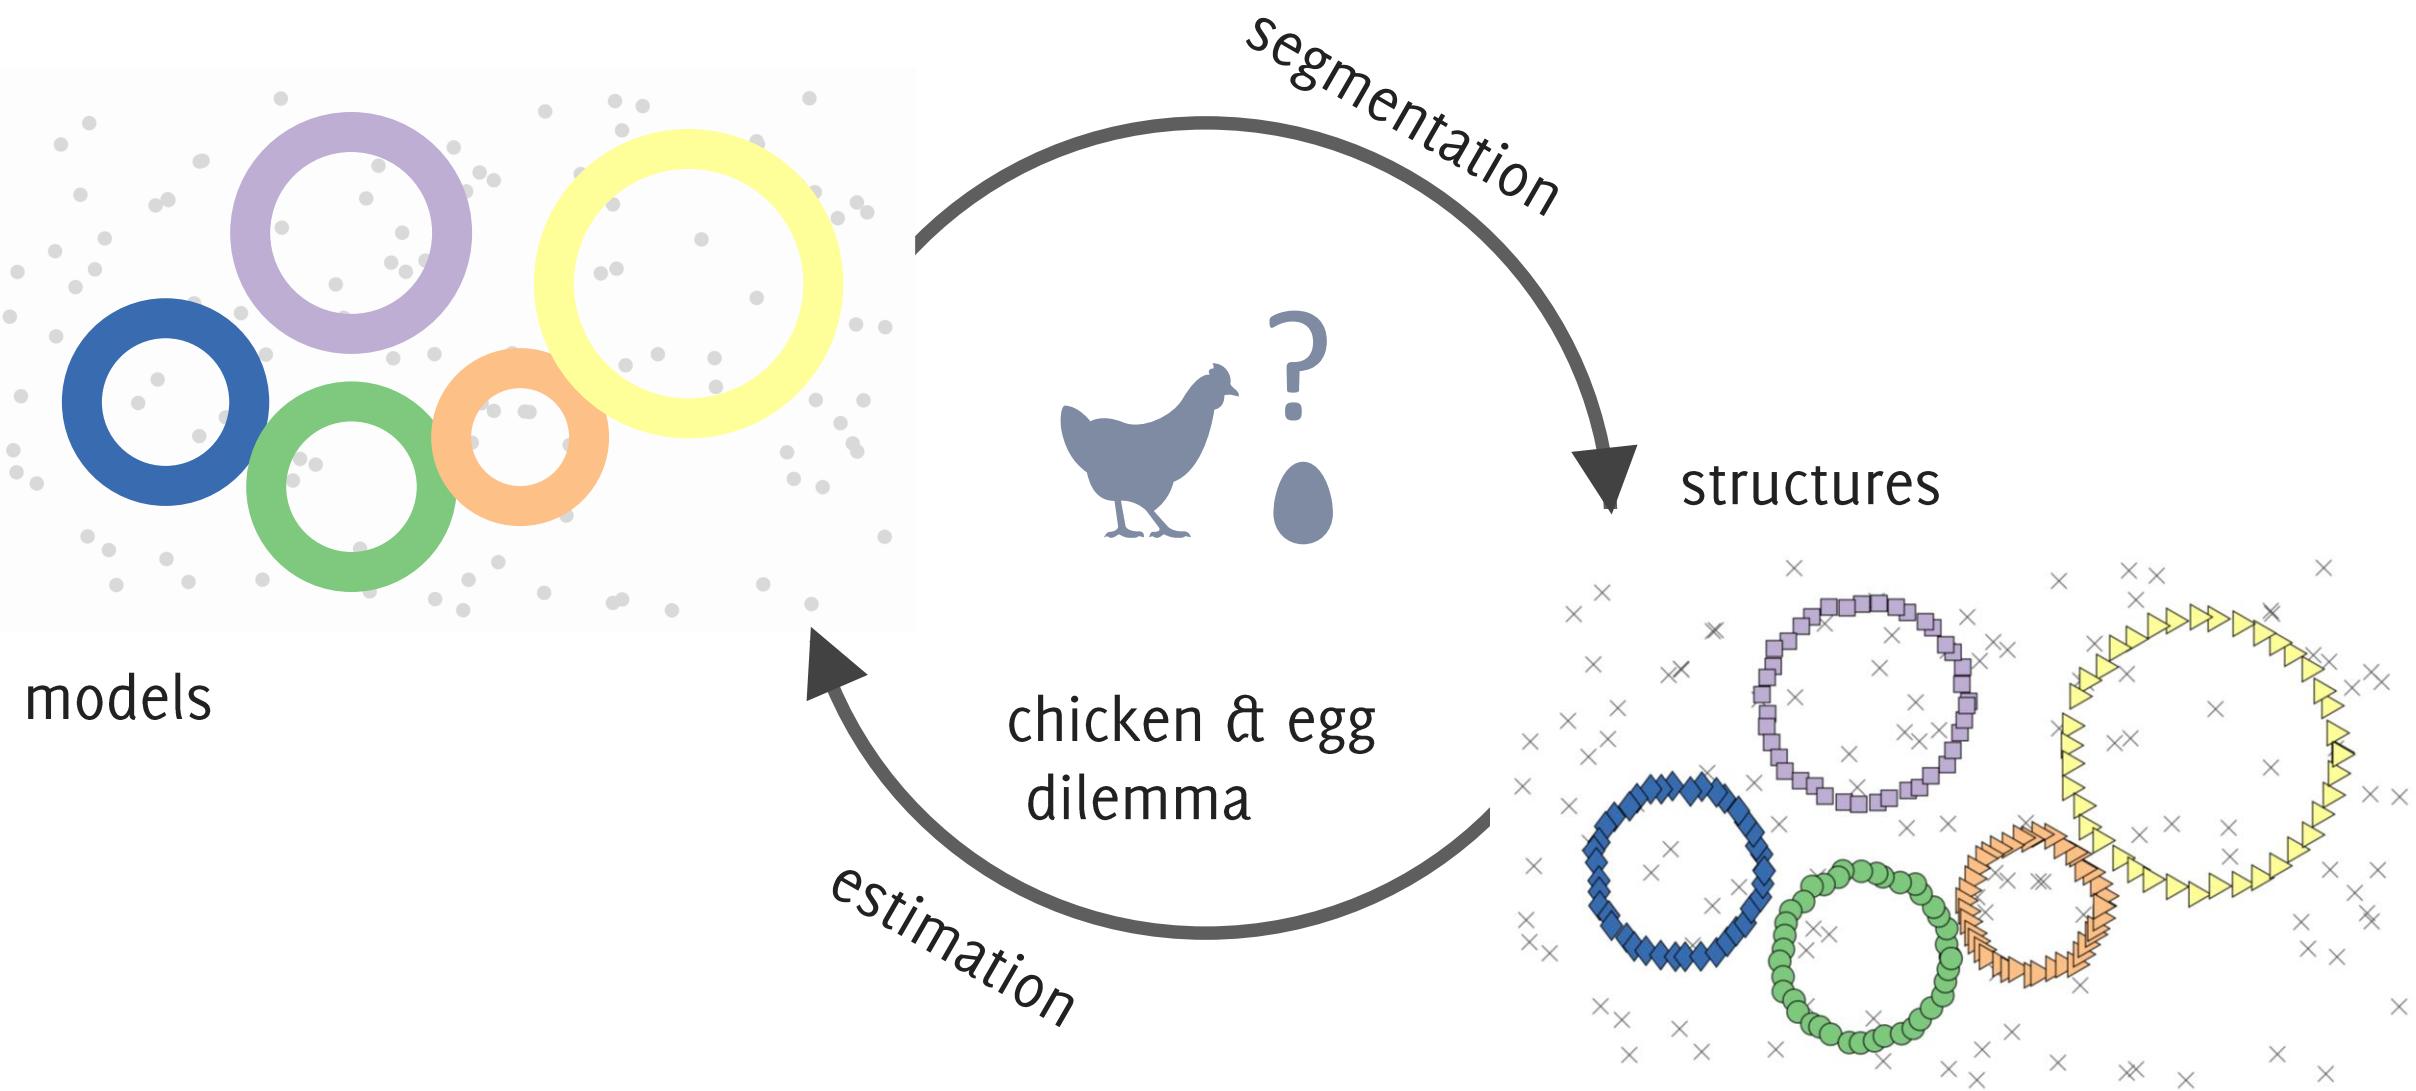


relations among the data, intrinsic to data

Goal: automatically estimate the models that best explain the data/discover the structures hidden in the data



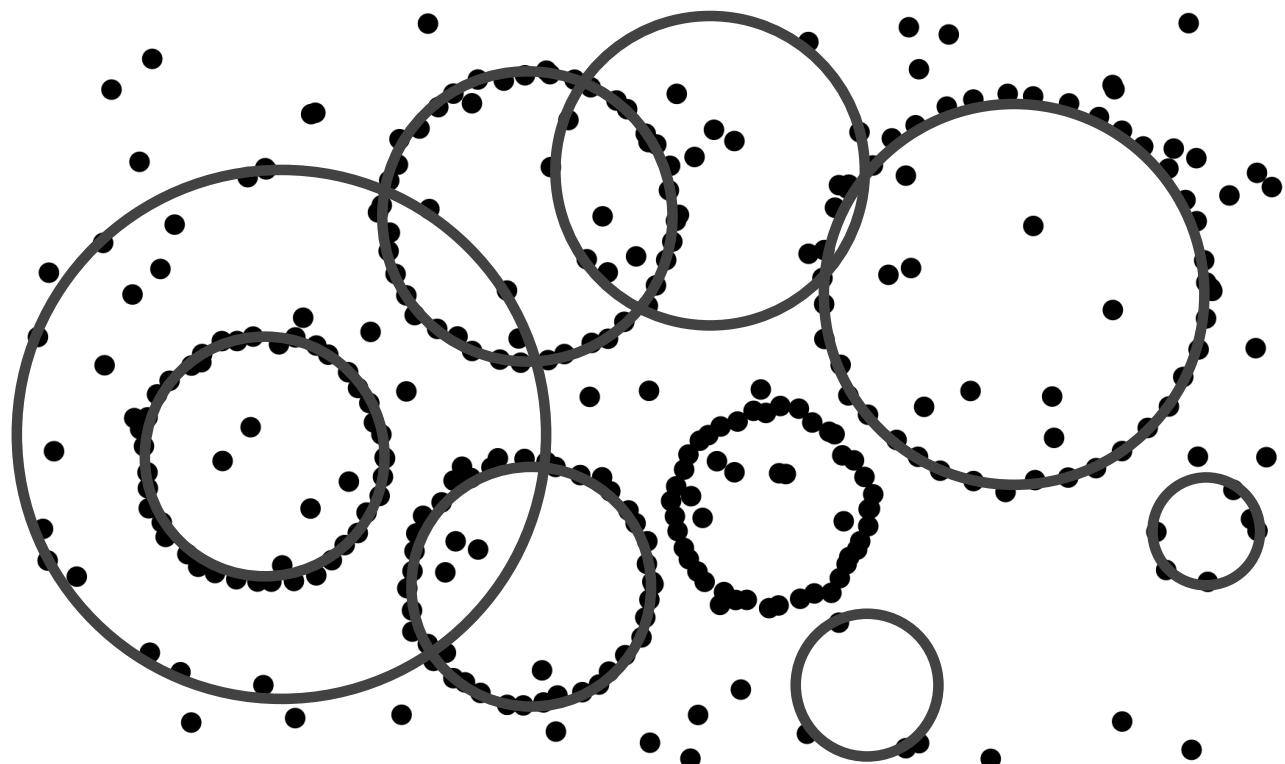
# The Challenges of multi-model fitting



# The Challenges of multi-model fitting

Number of models?

Inliers/outliers?

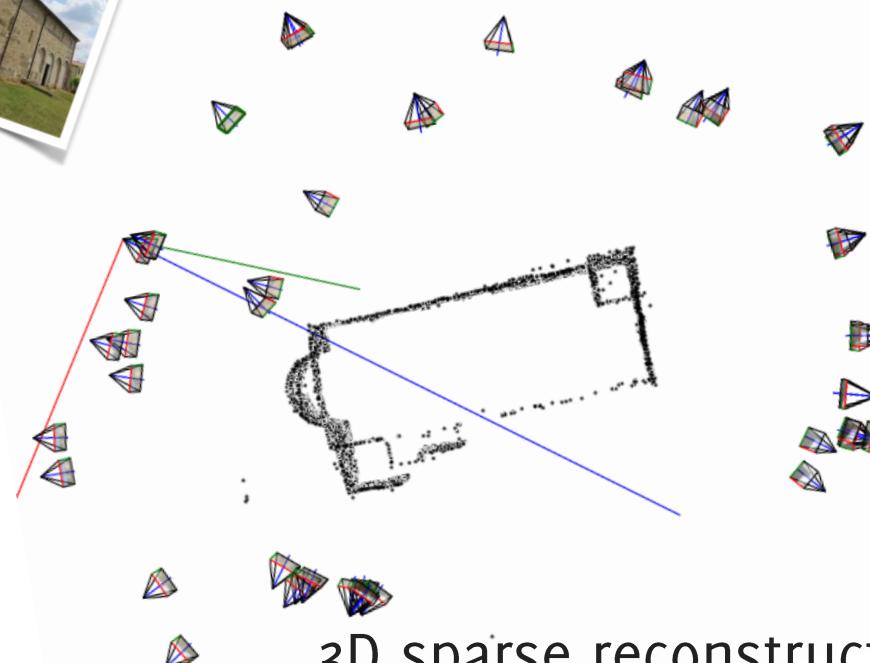


ill posed

# Multi model fitting applications: primitive fitting

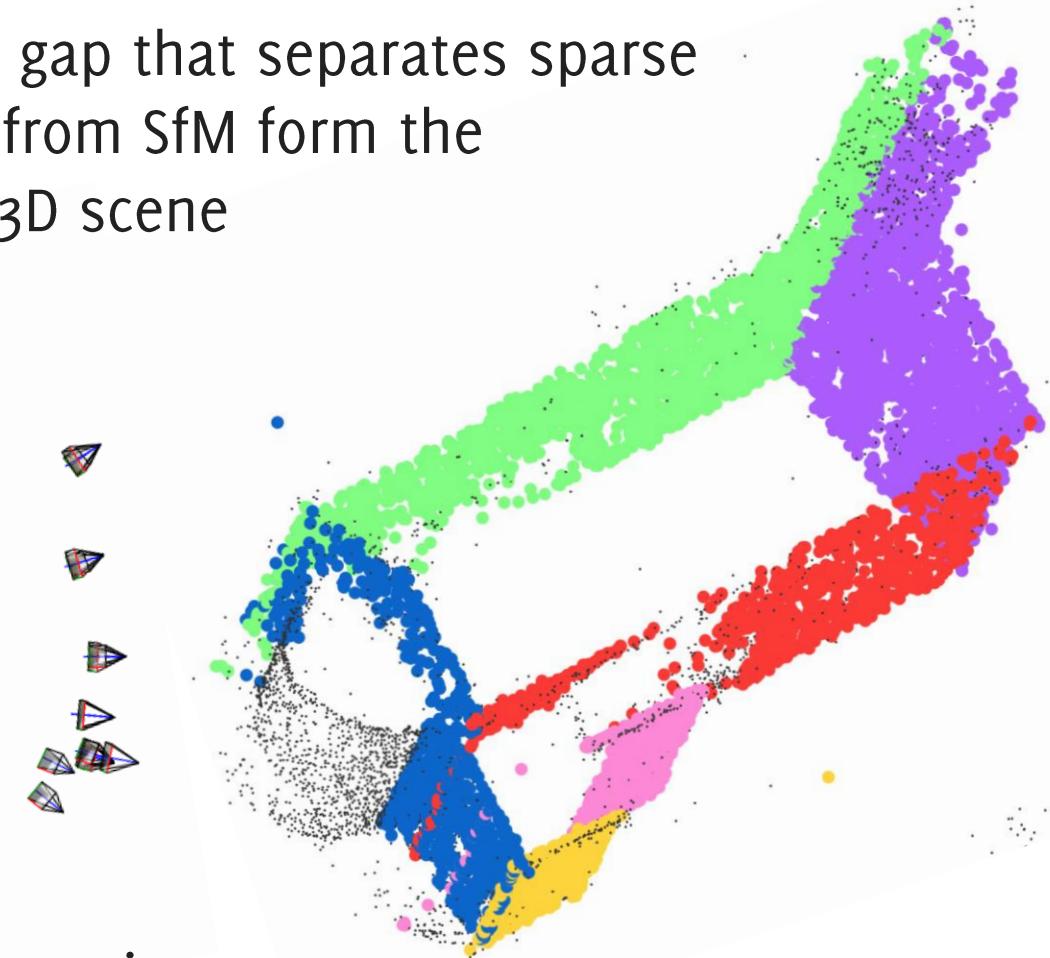


input images



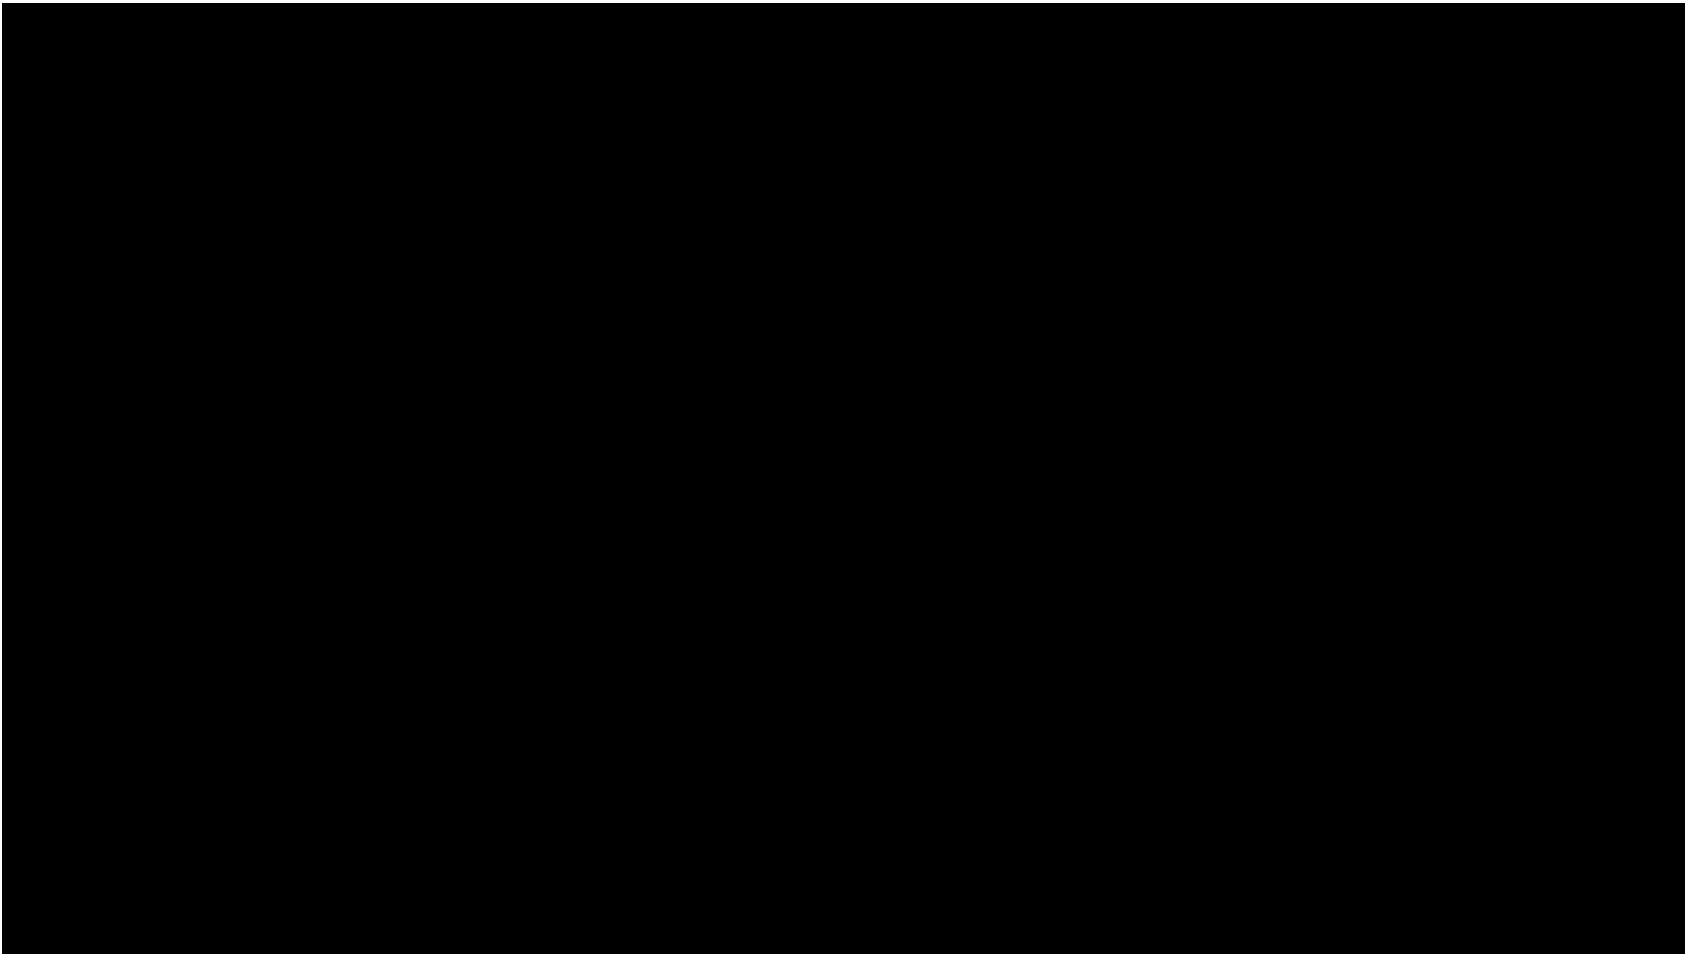
3D sparse reconstruction

Bridge the semantic gap that separates sparse point cloud coming from SfM form the understanding of a 3D scene



$$X \subset \mathbb{R}^3, \Theta = \text{planes}$$

# Multimodel fitting for 3D scattered data



L. Magri, and A. Fusiello. "Reconstruction of interior walls from point cloud data with min-hashed J-linkage." 2018 3DV

L. Magri, and A. Fusiello. "IMPROVING AUTOMATIC RECONSTRUCTION OF INTERIOR WALLS FROM POINT CLOUD DATA." International Archives of the Photogrammetry, Remote Sensing & Spatial Information Sciences (2019).

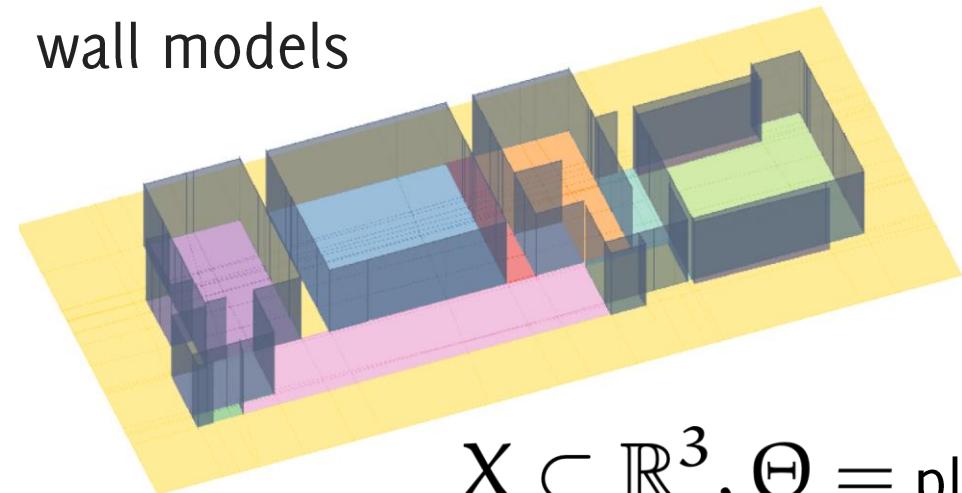
L. Magri, and Andrea Fusiello. "T-linkage: A continuous relaxation of j-linkage for multi-model fitting." CVPR 2014

# Multi model fitting applications: scan2bim

scanned point cloud

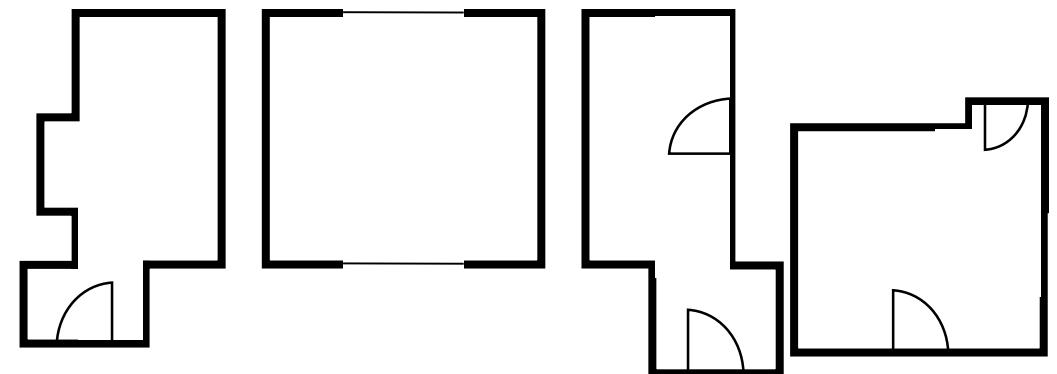


wall models



$$X \subset \mathbb{R}^3, \Theta = \text{planes}$$

floor-plan



Given a scanned point cloud of an interior environment, detect its primary facility surfaces – such as floors, walls, and ceilings.

$$X \subset \mathbb{R}^2, \Theta = \text{lines}$$

# Multi model fitting applications: two view geometry

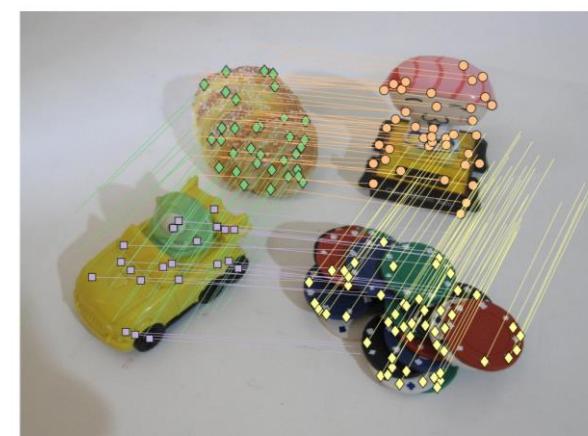
Geometric fit on corresponding matches across two images

plane detection



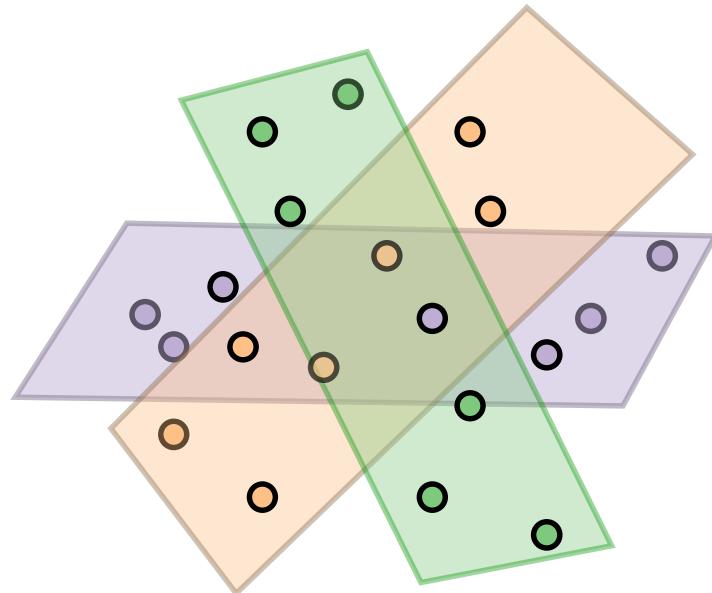
$X \subset \mathbb{R}^4$ ,  $\Theta = \text{homographies}$

epipolar geometry



$X \subset \mathbb{R}^4$ ,  $\Theta = \text{fundamental matrices}$

# Multi model fitting applications: subspace clustering



$$X \subset \mathbb{R}^d, \Theta = \text{subspaces}$$

3D Video segmentation



Face clustering

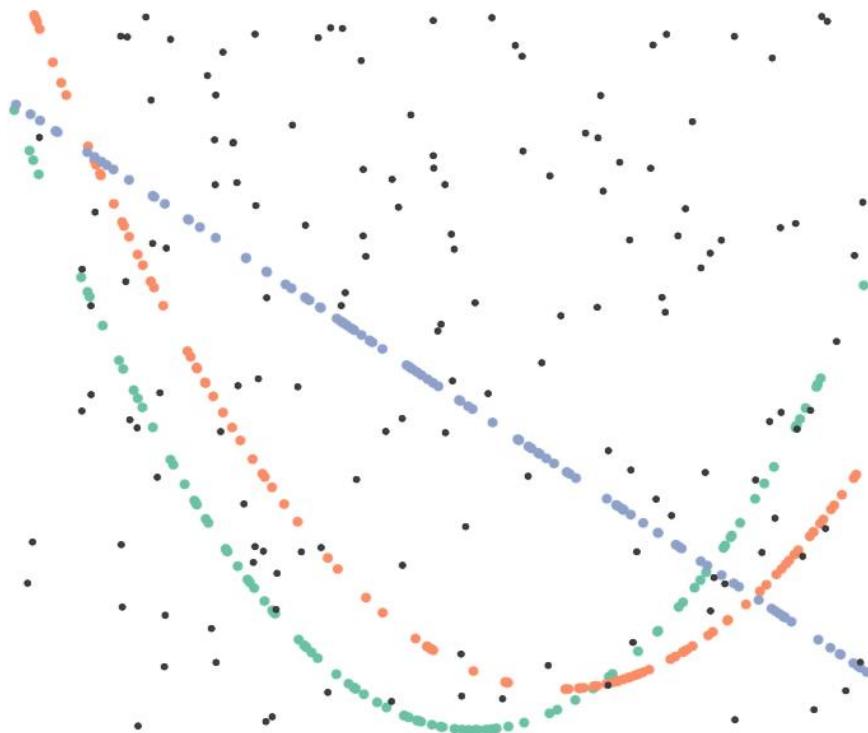


# Template Detection

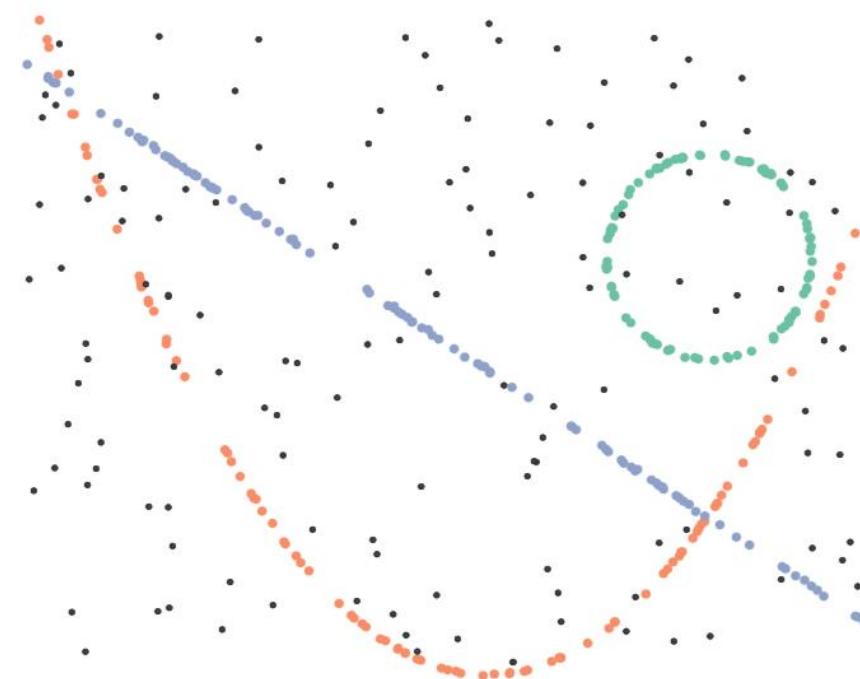


Our Collaboration with an Italian Company T&O

# Multimodel (and multi-class) fitting



(a)

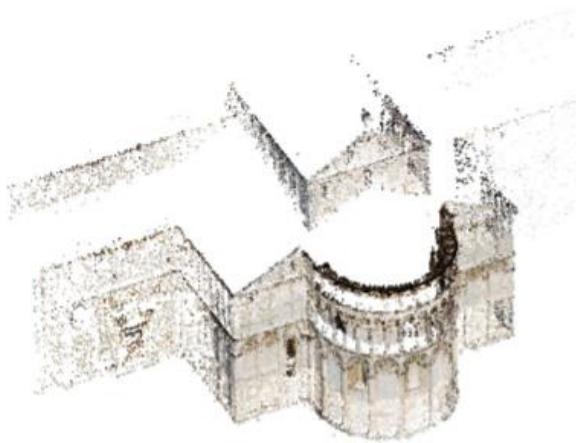
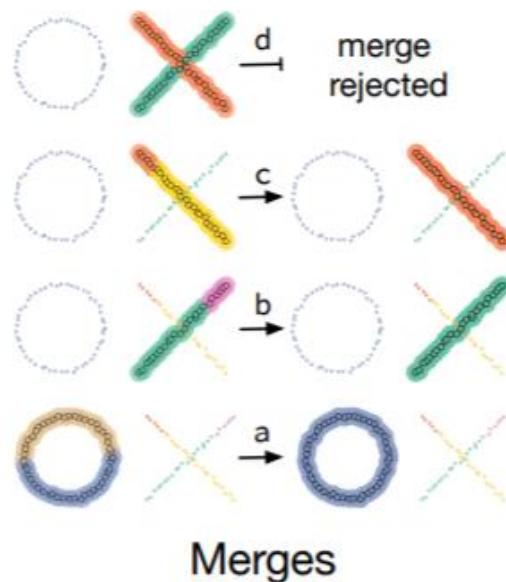
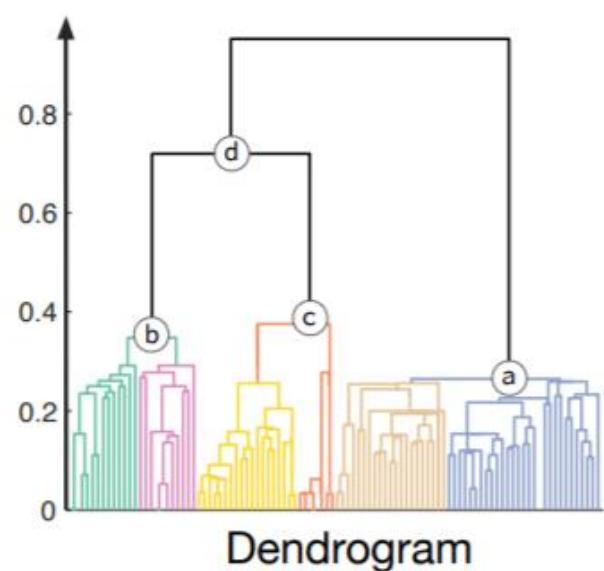


(b)

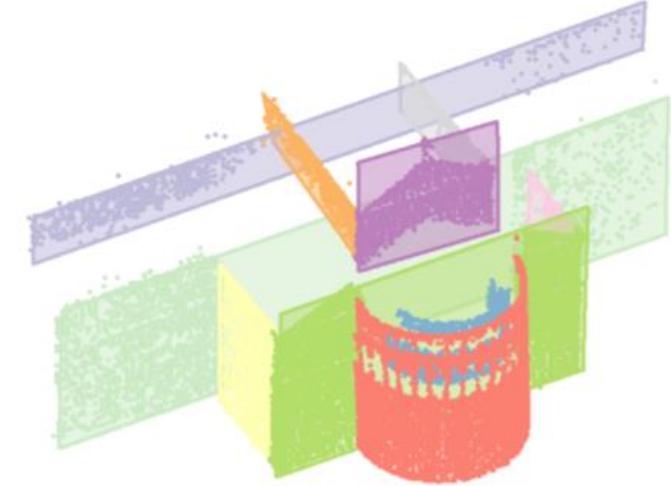


(c)

# MultiLink



(a) Input point cloud



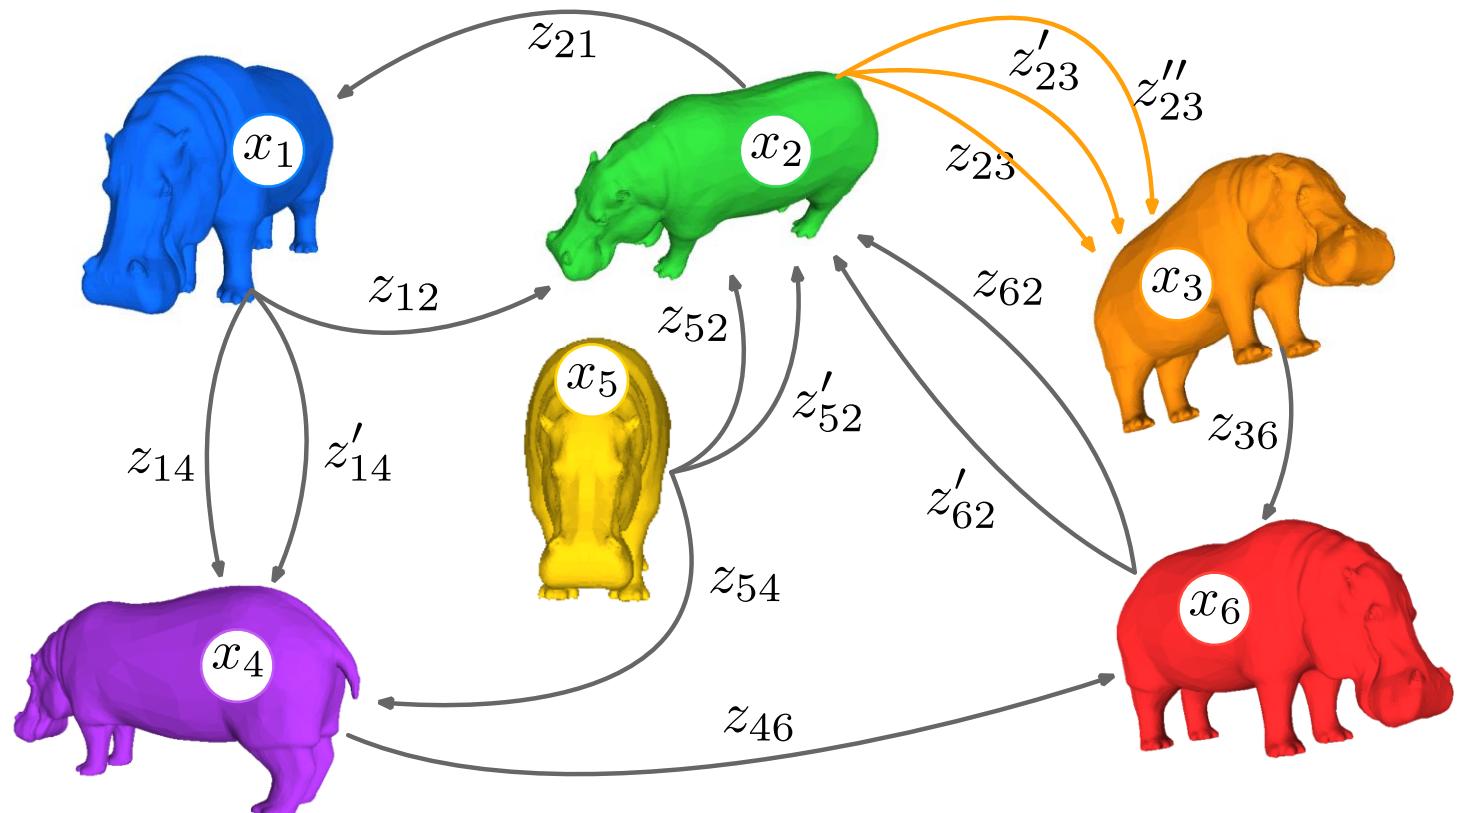
(b) Recovered structures

Figure 2: MultiLink combines single-linkage clustering and GRIC. Clusters are merged as long as the GRIC score improves when fitting suitable models on-the-fly. Colors indicate how cluster aggregation proceeds in the dendrogram.

# Synchronization

Robustly synchronize multiple transformations that can be represented in a matrix group:

- Perform global 3D registration avoiding drifting typical of incremental approaches.
- Graph-partitioning to synchronize very large graphs



# Multi-model Fitting Solutions

# Several solutions

## Optimization-based

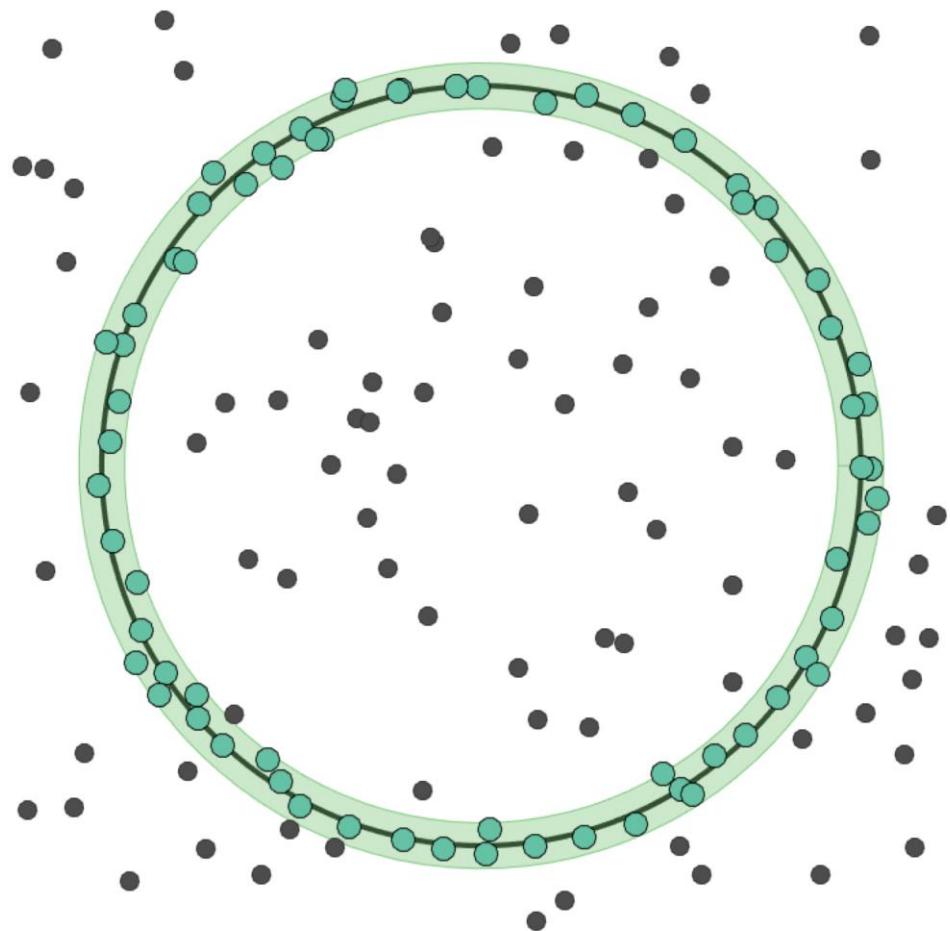
- [Xu et Al PatRecLet 90] Randomize Hough Transform
- [Zuliani ICIP 05] Multi Ransac
- [Isack and Boykov CVPR 10] Pearl
- [Magri and Fusiello CVPR 17] Set coverage
- [Barath and Matas ECCV 18] MultiX
- [Barath and Matas ICCV 19] ProgressiveX

## Clustering-based

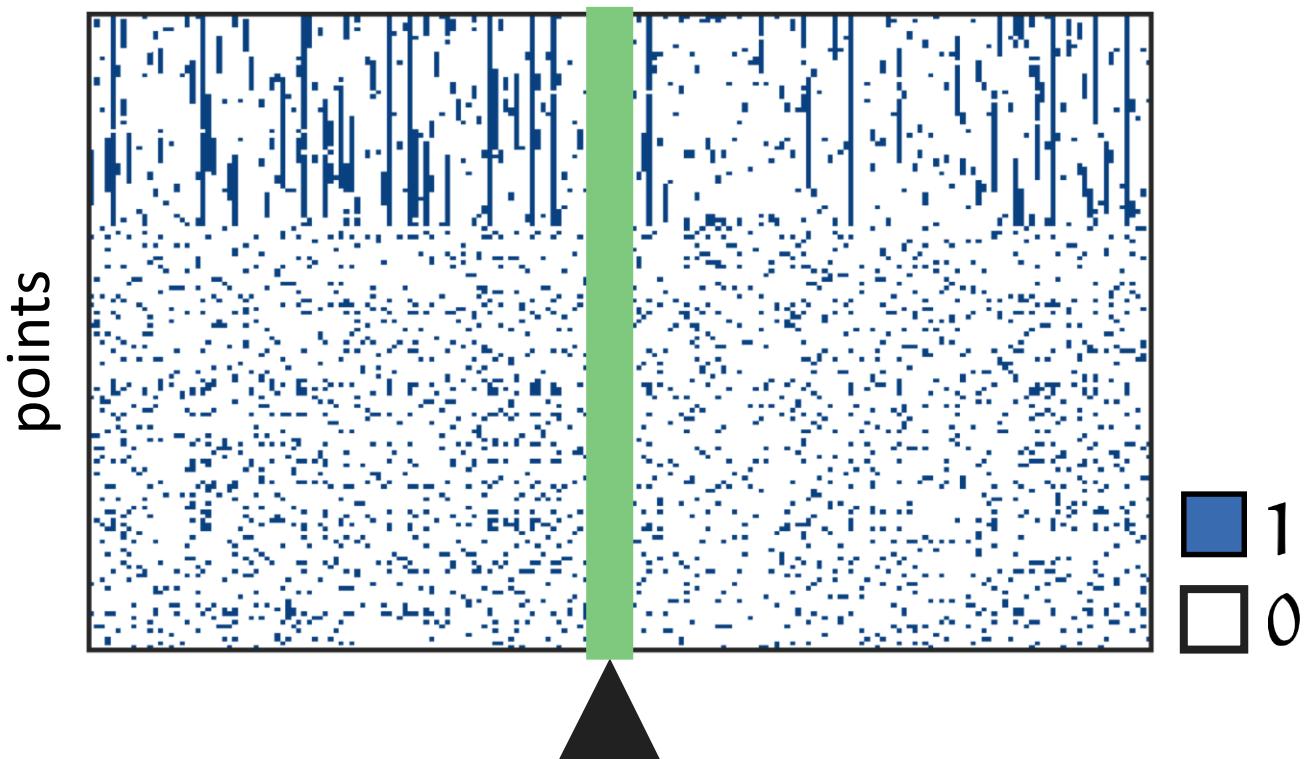
- [Zhang and Koseckà ECCV 06] Residual analysis
- [Toldo and Fusiello ECCV 08] J-linkage
- [Chin et al ICCV 09] Ordered residual kernel
- [Purkait et Al ECCV 14] Hypergraph clustering
- [Magri and Fusiello CVPR 14] T-linkage
- [Magri and Fusiello BMVC 15] Matrix factorization
- [Magri and Fusiello CVPR 19] Cascaded T-linkage
- [Magri et Al, CVPR 21] Multi-Link

Clustering provides a simple and powerful formulation based on few intelligible parameters

# Let's go back to RanSaC



Data driven search of model space  
 $H = \{\theta_1, \theta_2, \dots, \theta_m\} \approx \Theta$   
tentative models



# Sequential RanSaC

Start RanSaC on the dataset  $X$  searching for the best fit for a single instance of the model



Line fitting example

# Sequential RanSaC

Start RanSaC on the dataset  $X$  searching for the best fit for a single instance of the model

Once detected a model  $\theta_1$ , keep the model and remove all the inliers



Line fitting example

# Sequential RanSaC

Start RanSaC on the dataset  $X$  searching for the best fit for a single instance of the model

Once detected a model  $\theta_1$ , keep the model and remove all the inliers from  $X$



Line fitting example

# Sequential RanSaC

Start RanSaC on the dataset  $X$  searching for the best fit for a single instance of the model

Once detected a model  $\theta_1$ , keep the model and remove all the inliers from  $X$

Iterate through the remaining points



Line fitting example

# Sequential RanSaC

Start RanSaC on the dataset  $X$  searching for the best fit for a single instance of the model

Once detected a model  $\theta_1$ , keep the model and remove all the inliers from  $X$

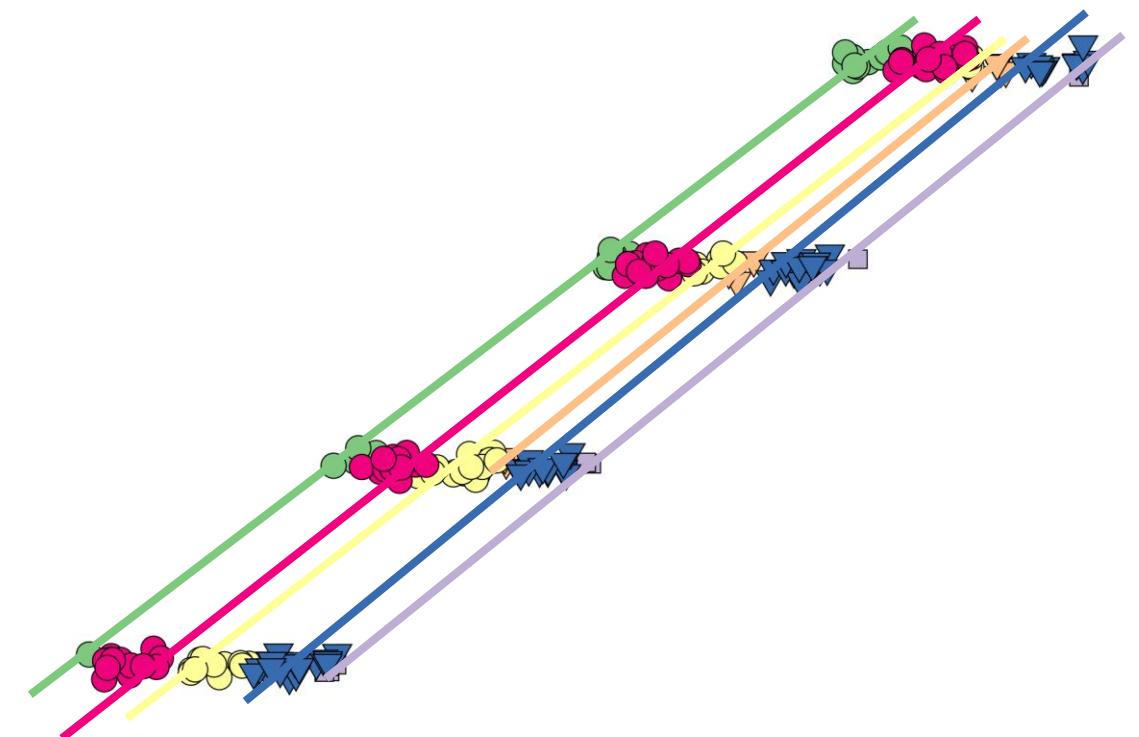
Iterate through the remaining points until there are no models with a sufficiently large consensus



**Line fitting example**

# Sequential RanSaC

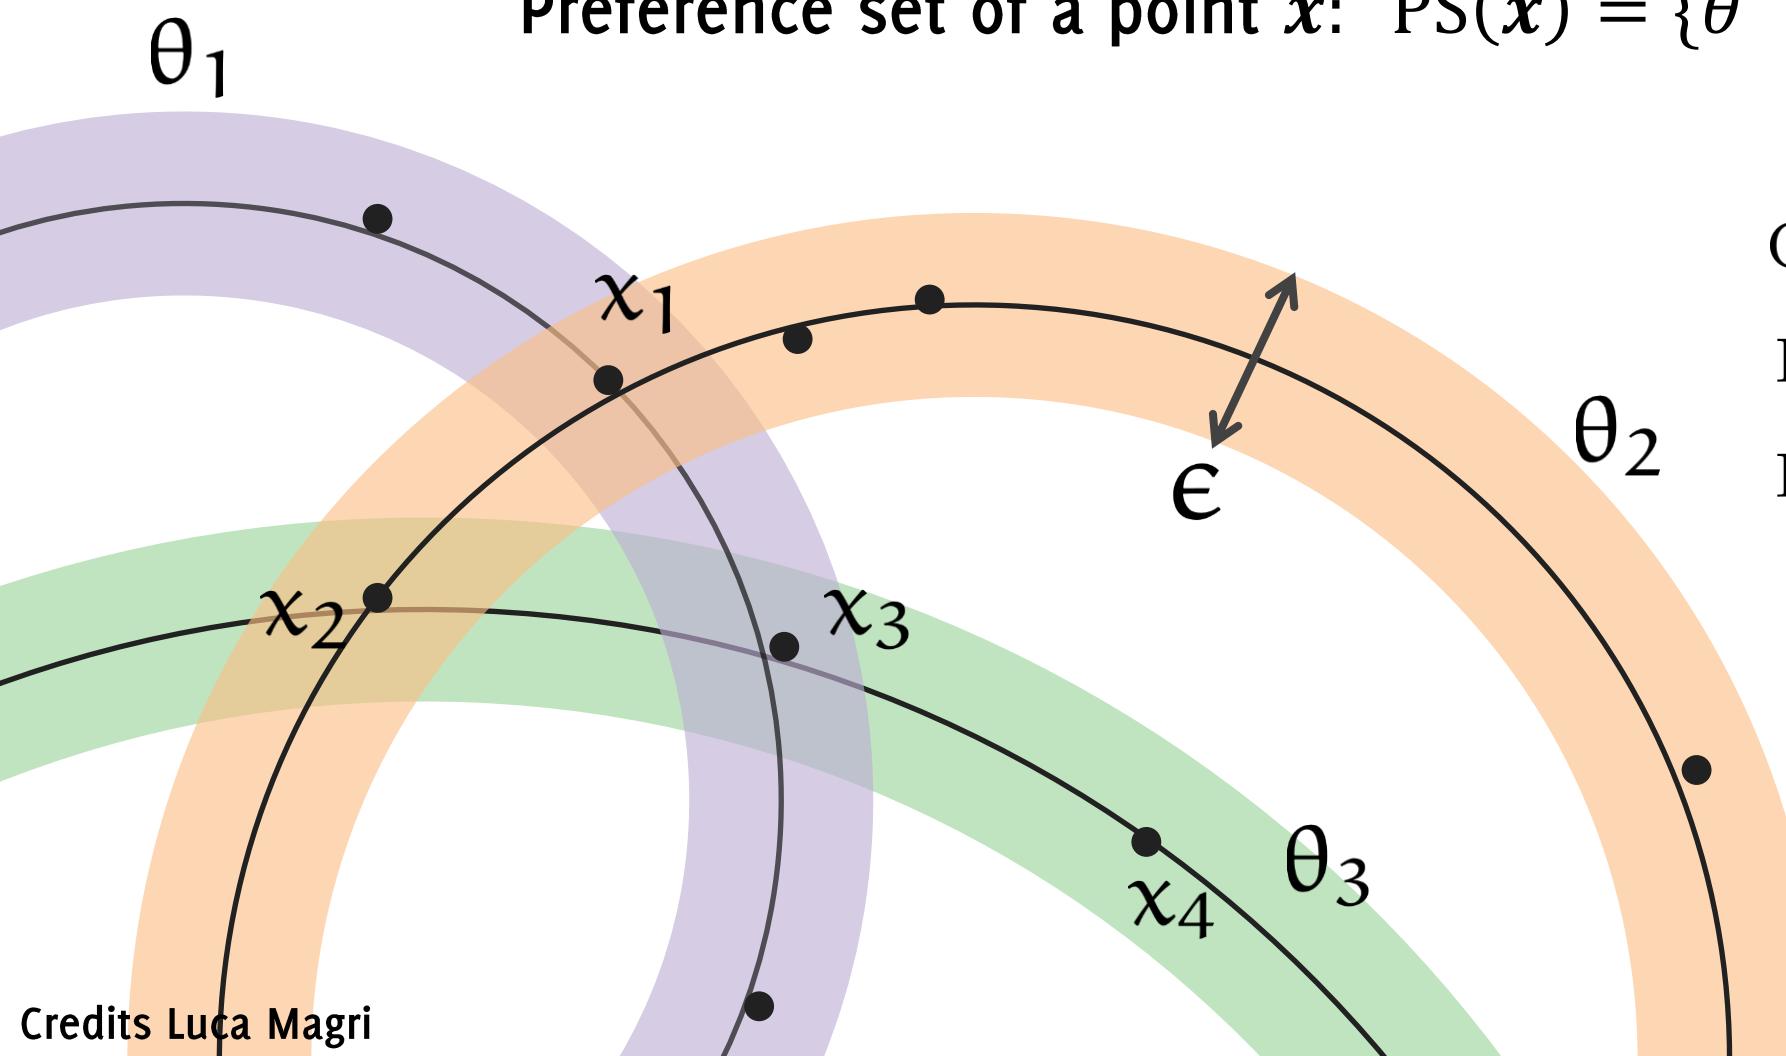
Unfortunately, this does not fit well with the multi-model scenario  
and the problem becomes even more sever in presence of outliers



# Preference-based methods

# From model consensus to point preferences

Consensus set of a model  $\theta$ :  $CS(\theta) = \{x : r(x, \theta) < \epsilon\}$   
Preference set of a point  $x$ :  $PS(x) = \{\theta : r(x, \theta) < \epsilon\}$



$$CS(\theta_3) = \{x_2, x_3, x_4\}$$

$$PS(x_1) = \{\theta_1, \theta_2\}$$

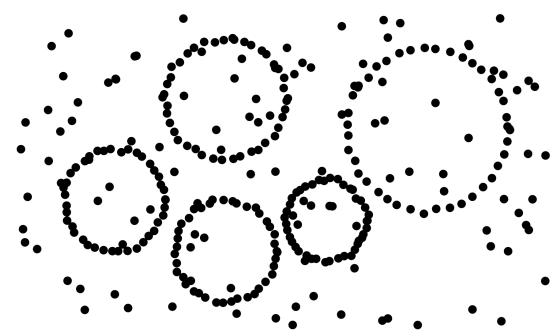
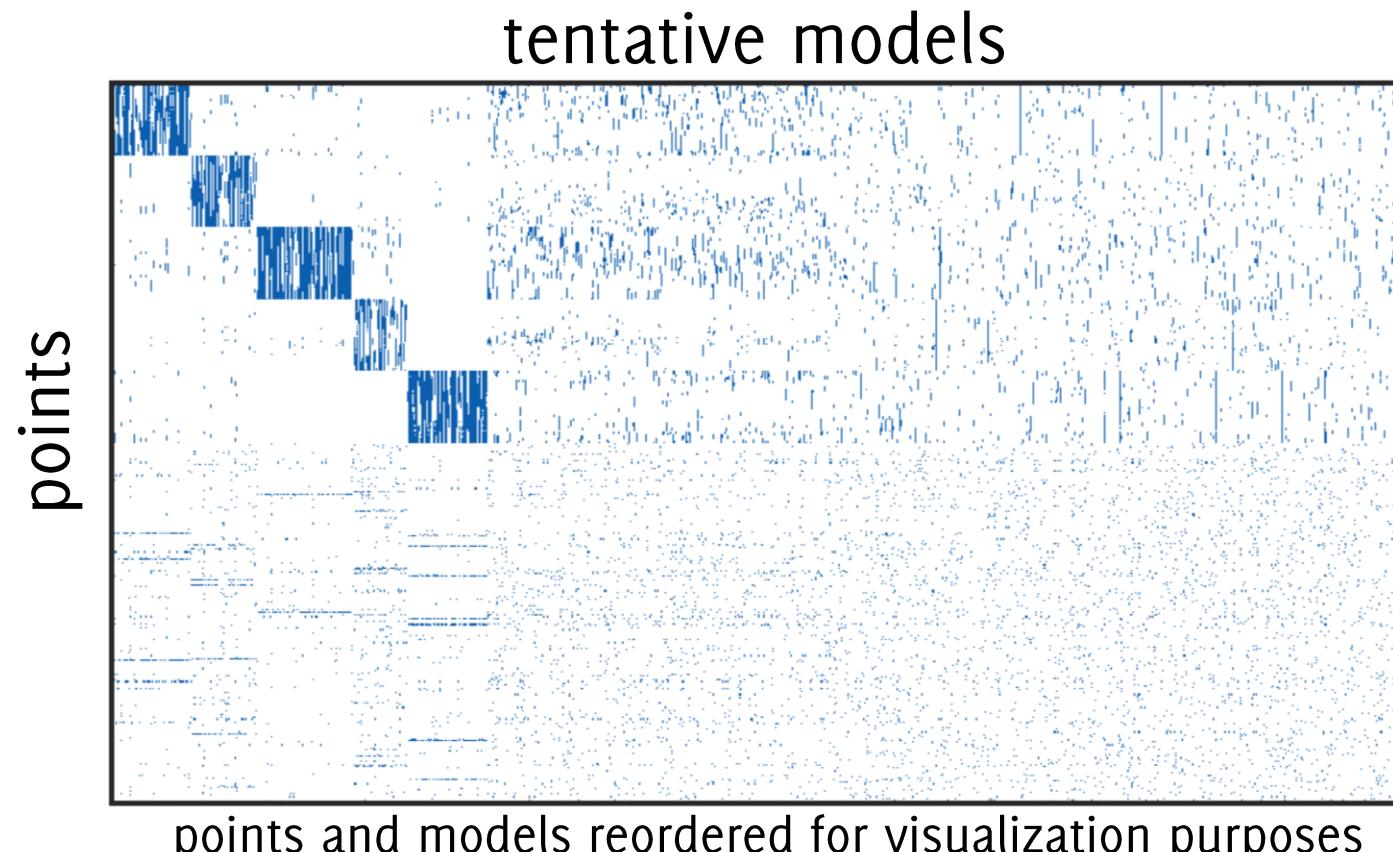
$$PS(x_2) = \{\theta_2, \theta_3\}$$

# Multi model fitting: the preference trick

Generate a pool of  $m$  random models (as in RanSaC)

$$H = \{\theta_1, \theta_2, \dots, \theta_m\}$$

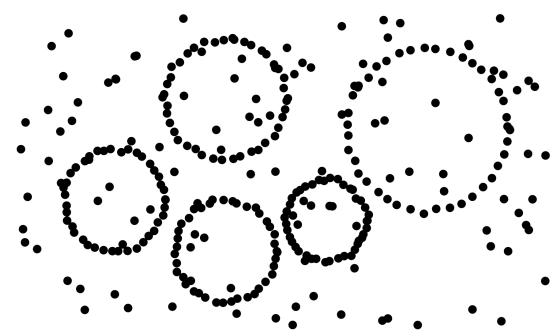
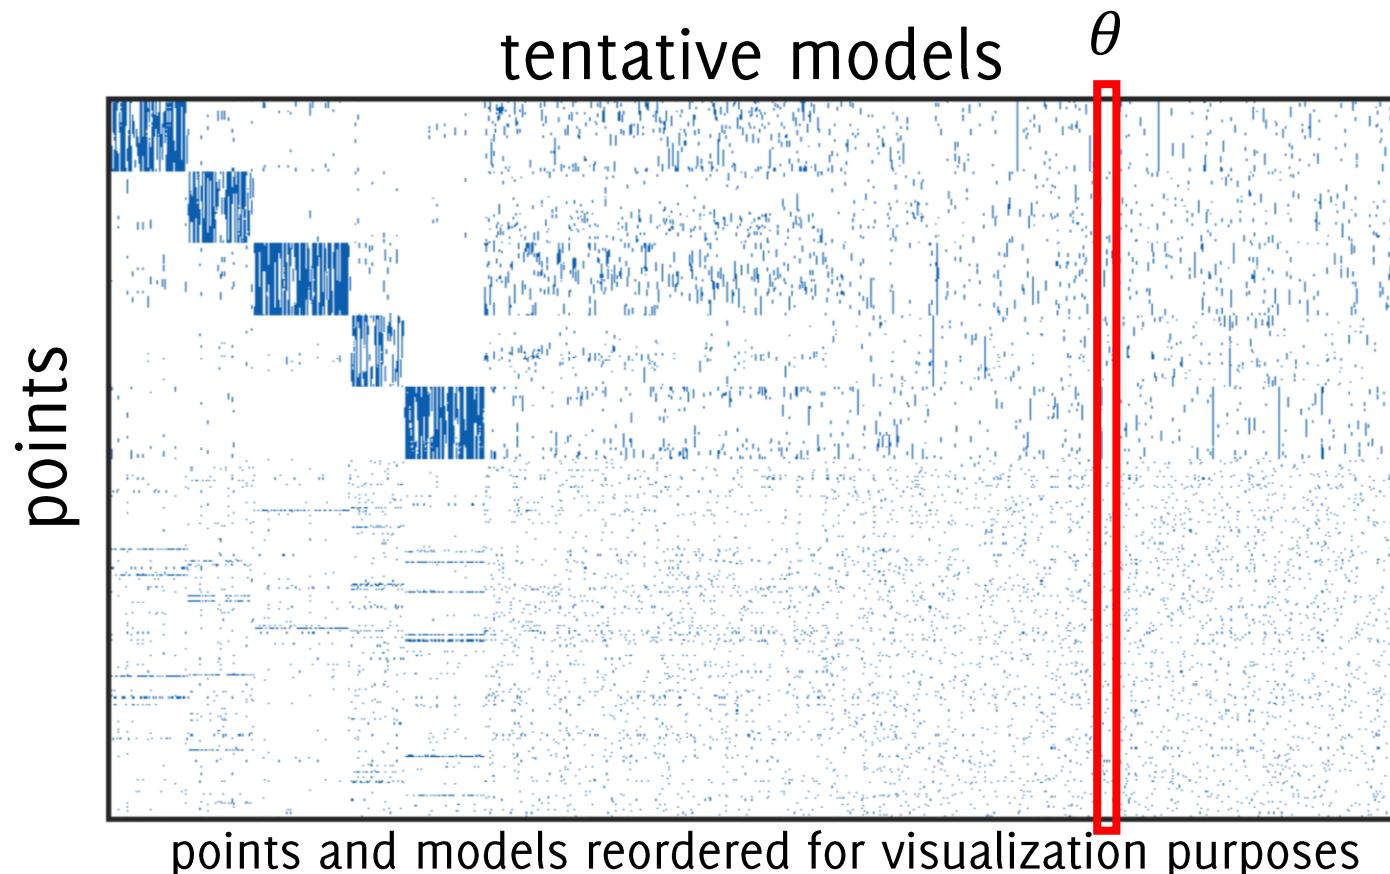
Build the Preference Matrix



# Multi model fitting: the preference trick

Consensus set of a model  $\theta$ :  $\text{CS}(\theta) = \{x : r(x, \theta) < \epsilon\}$

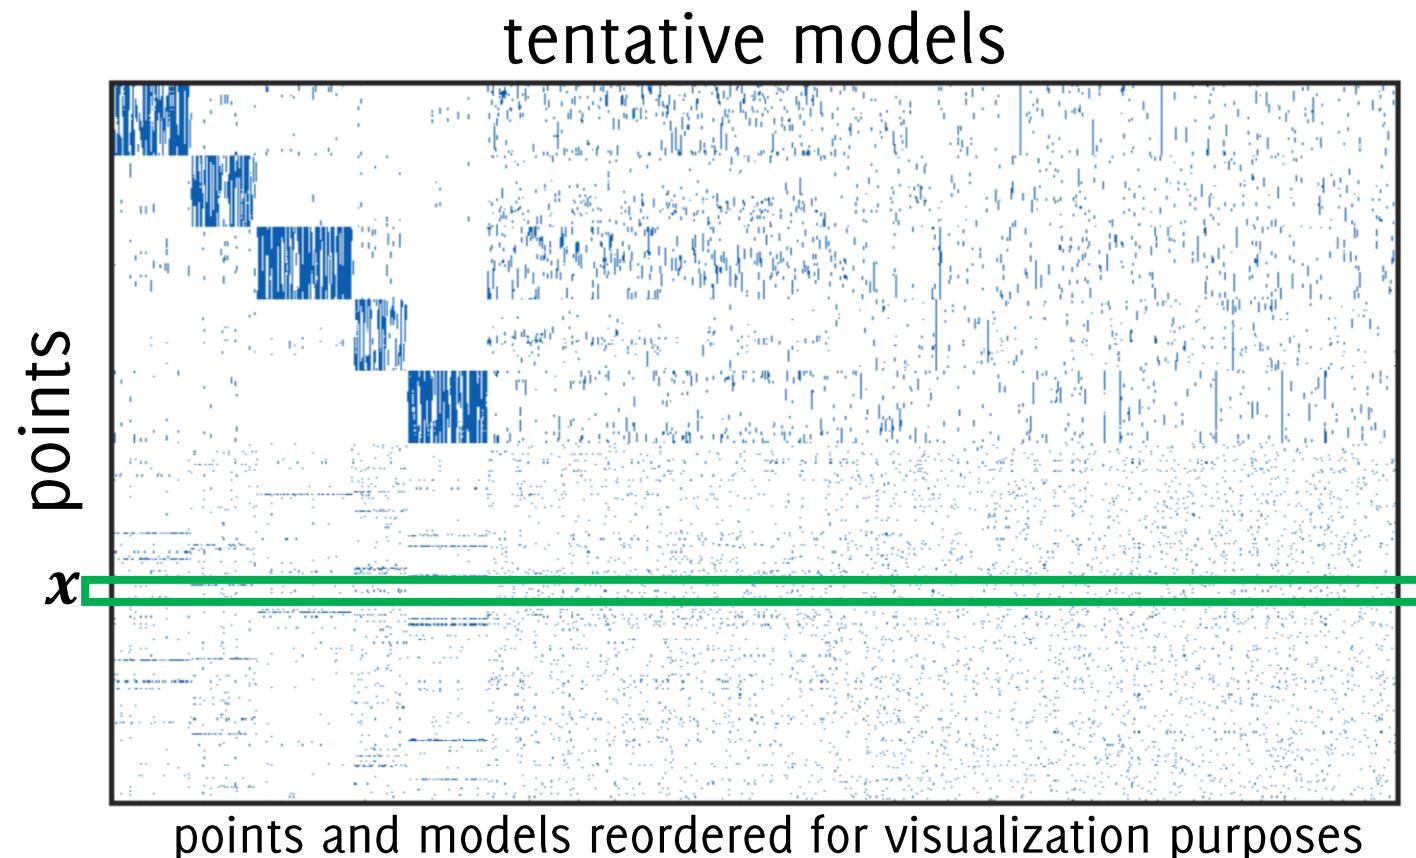
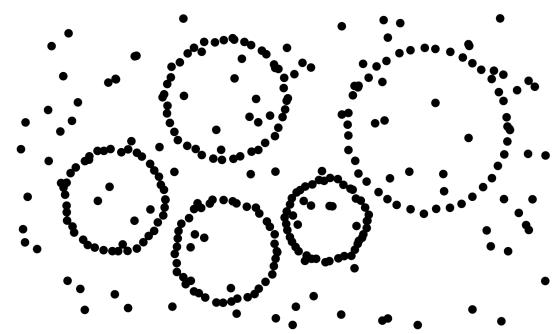
Preference set of a point  $x$ :  $\text{PS}(x) = \{\theta : r(x, \theta) < \epsilon\}$



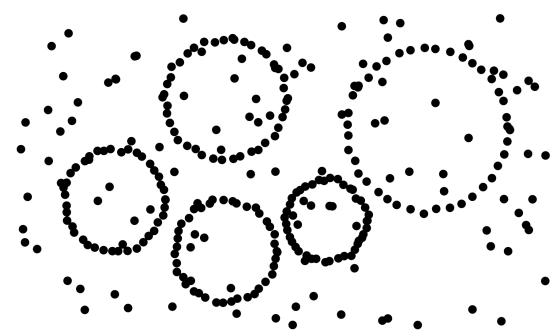
# Multi model fitting: the preference trick

Consensus set of a model  $\theta$ :  $CS(\theta) = \{x : r(x, \theta) < \epsilon\}$

Preference set of a point  $x$ :  $PS(x) = \{\theta : r(x, \theta) < \epsilon\}$



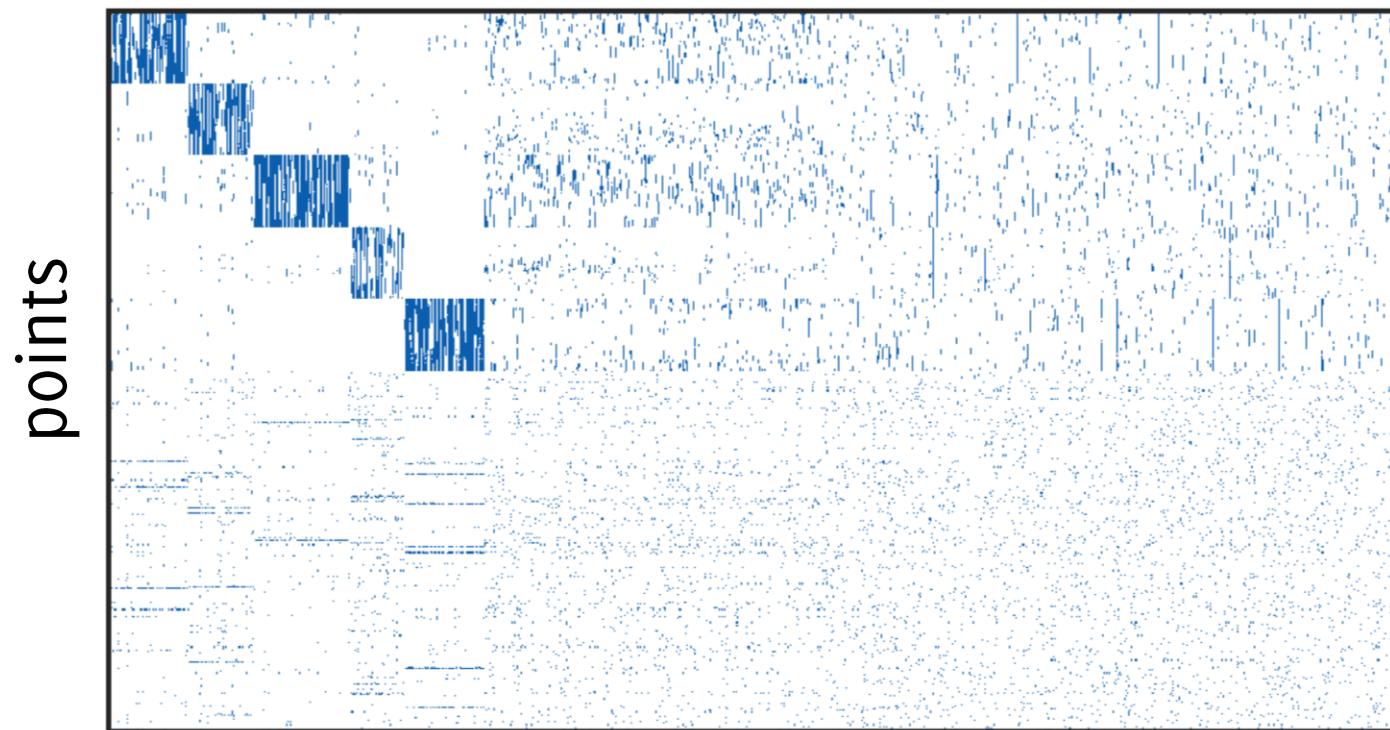
# Multi model fitting: the preference trick



Point  $\leftrightarrow$  subset of preferred sampled models

Block diagonal matrix  $\Rightarrow$  point of the same structure have similar preferences

tentative models

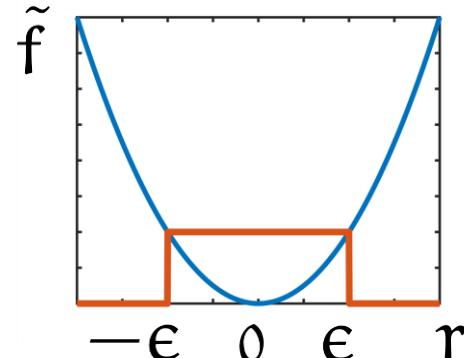


points and models reordered for visualization purposes

# Multi model fitting: a lift to Preference Space

PS( $x$ )

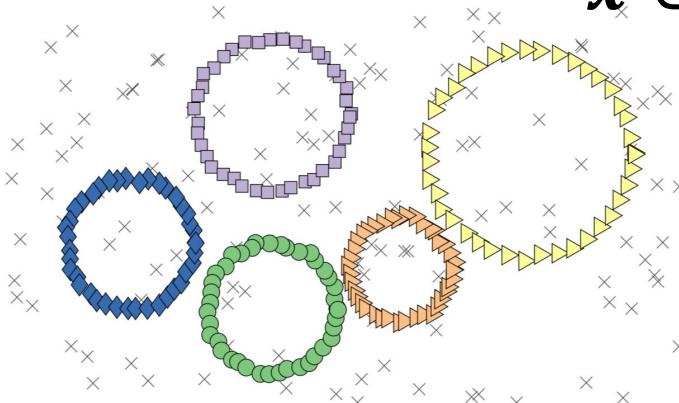
Voting function



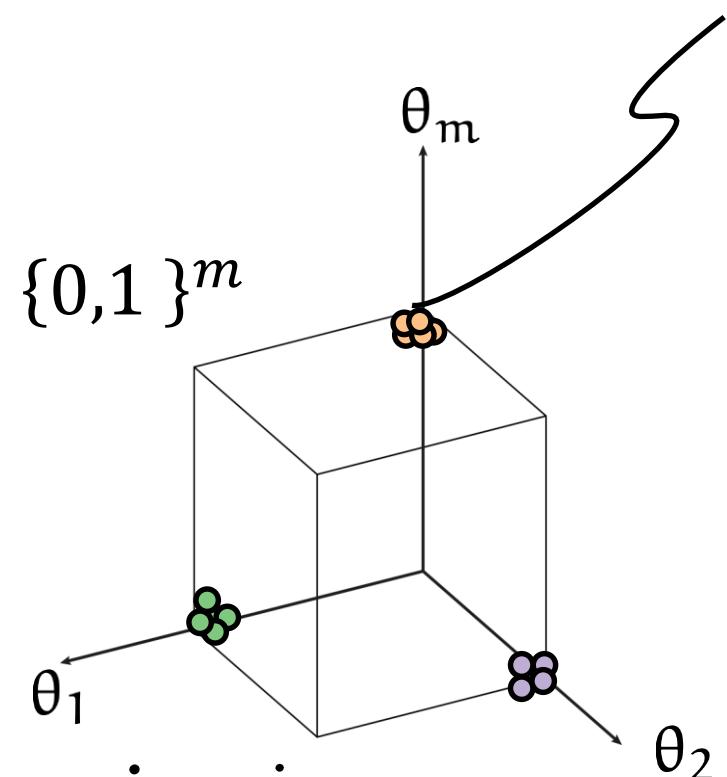
Vector of binary votes to sampled models

$$x \mapsto [\hat{f}(r(x, \theta_1)), \dots, \hat{f}(r(x, \theta_m))] \in \{0, 1\}^m$$

$$x \in \mathbb{R}^d$$



$$\text{PS}(x) \in \{0, 1\}^m$$

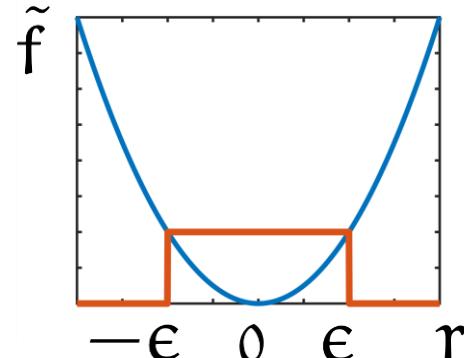


Structures correspond to high-density regions in  
the preference space

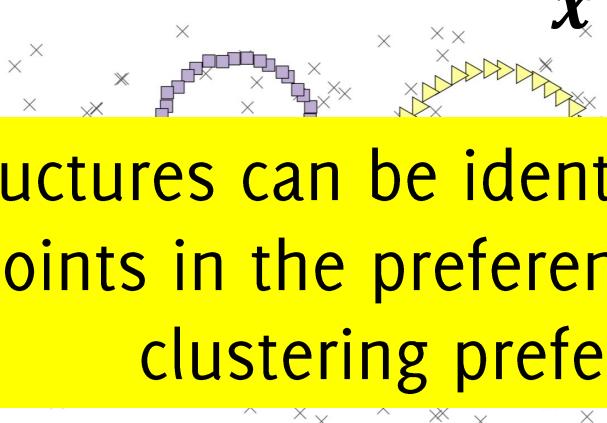
# Multi model fitting: a lift to Preference Space

PS( $x$ )

Voting function



$$x \in \mathbb{R}^d$$

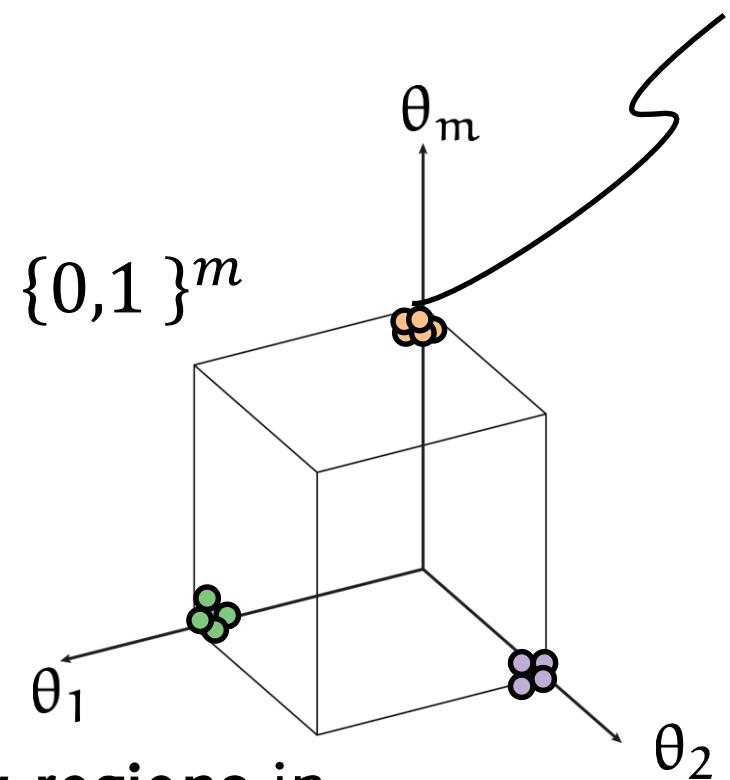


Vector of binary votes to sampled models

$$x \mapsto [\hat{f}(r(x, \theta_1)), \dots, \hat{f}(r(x, \theta_m))] \in \{0, 1\}^m$$



$$PS(x) \in \{0, 1\}^m$$



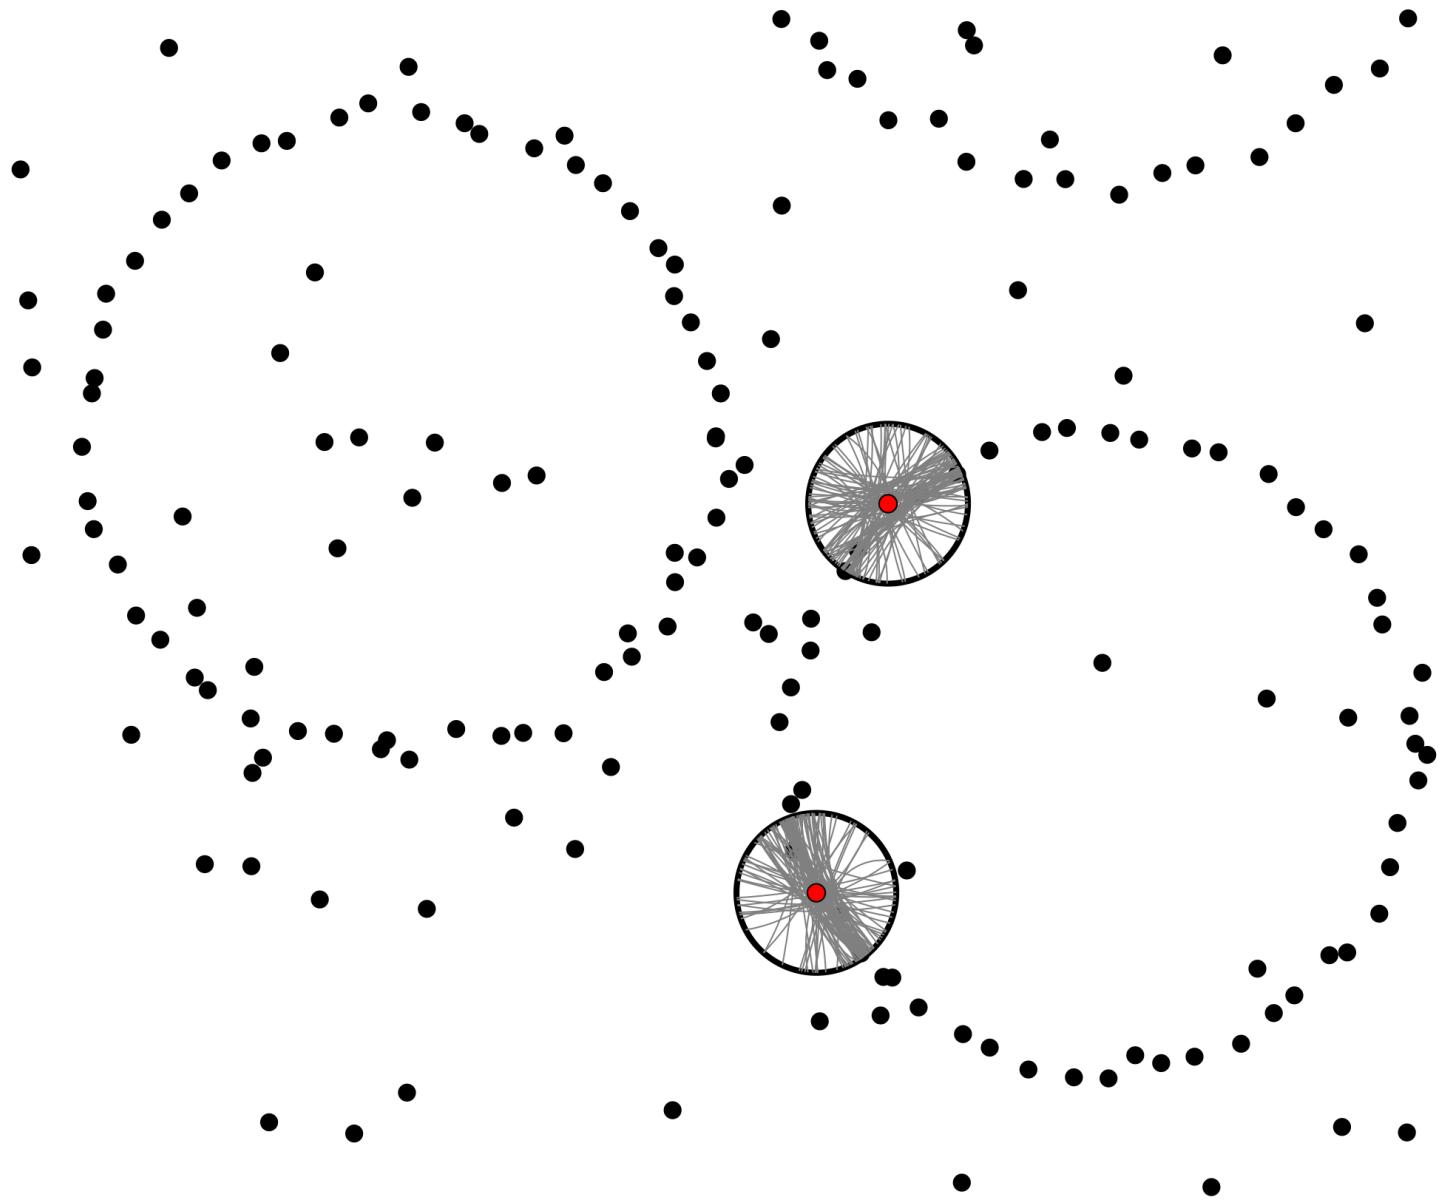
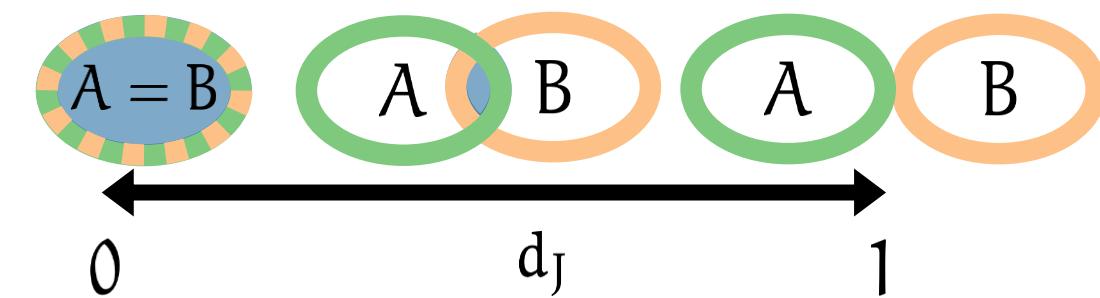
Structures can be identified by clustering points in the preference space, i.e. by clustering preference sets.

Structures correspond to high-density regions in the preference space

# Multi model fitting: a lift to Preference Space

The [Jaccard distance](#) can be used to measure distance between preference sets.

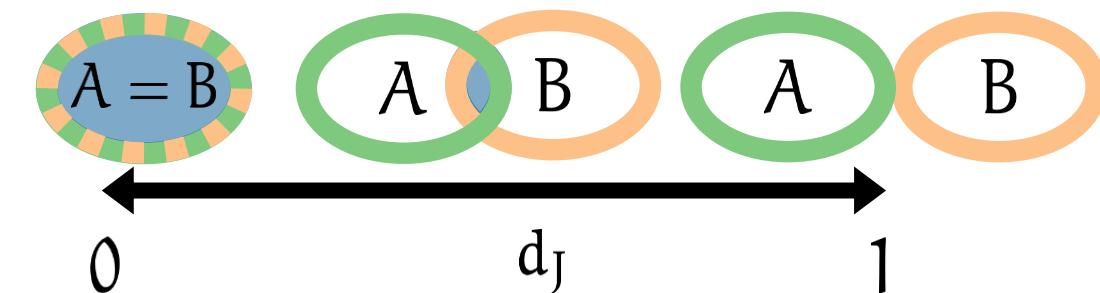
$$d_J(A, B) = 1 - \frac{|A \cap B|}{|A \cup B|}$$



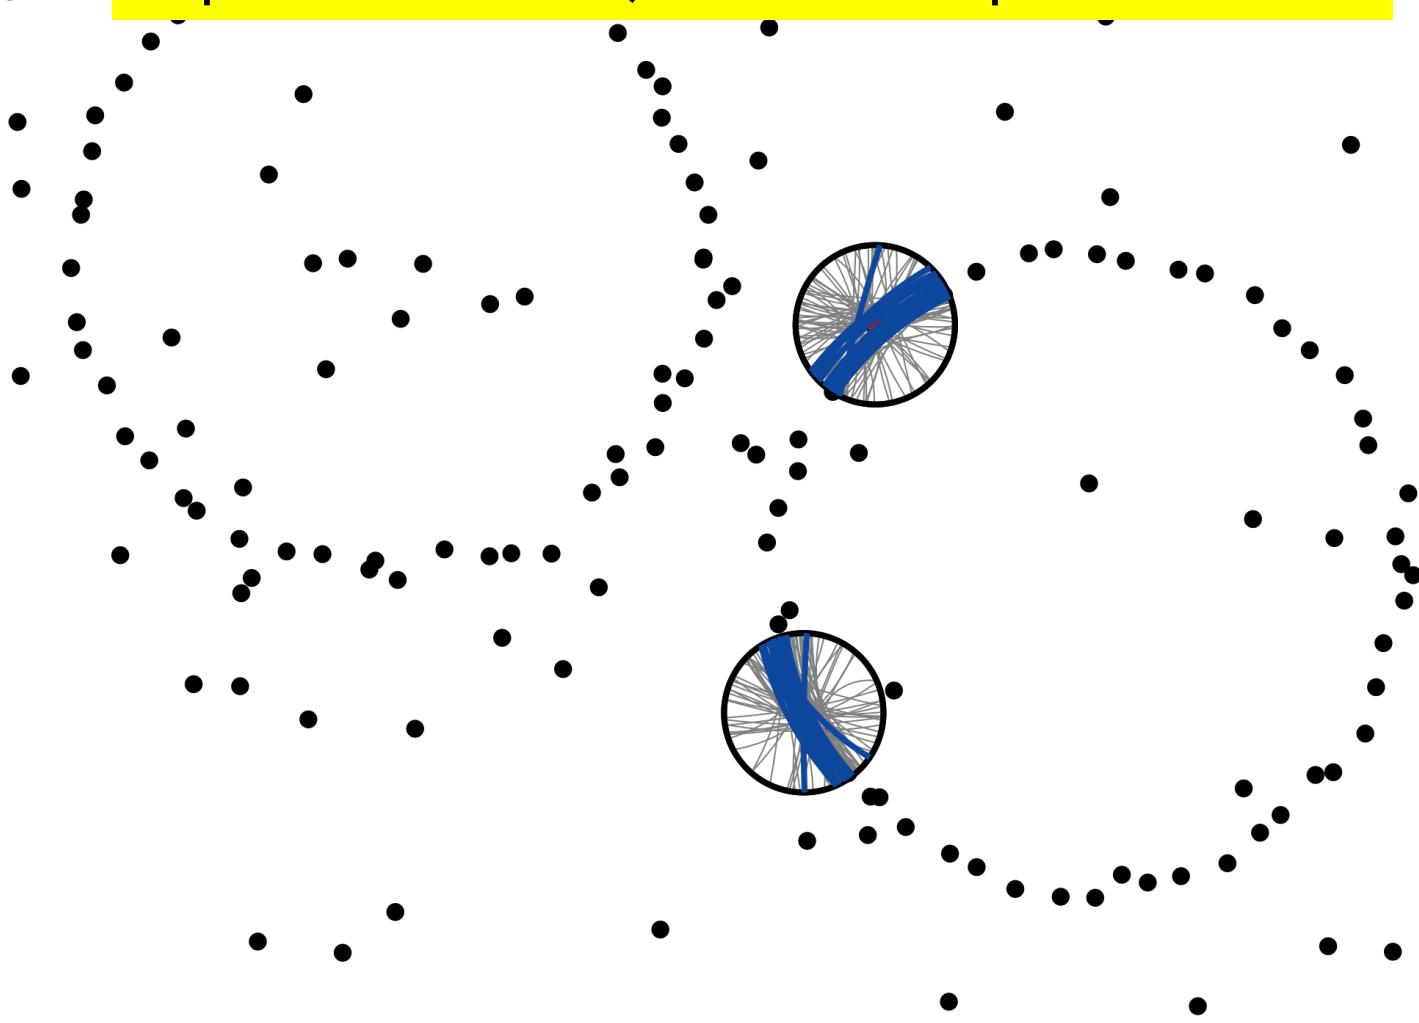
# Multi model fitting: a lift to Preference Space

The [Jaccard distance](#) can be used to measure distance between preference sets.

$$d_J(A, B) = 1 - \frac{|A \cap B|}{|A \cup B|}$$



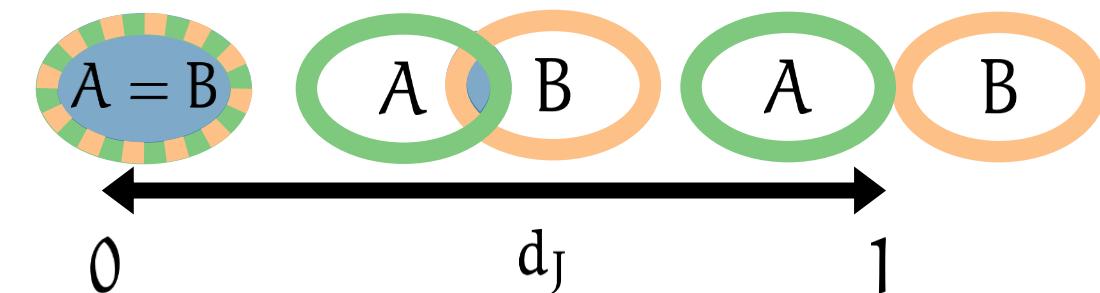
the more models in common in their preference set, the closer points are.



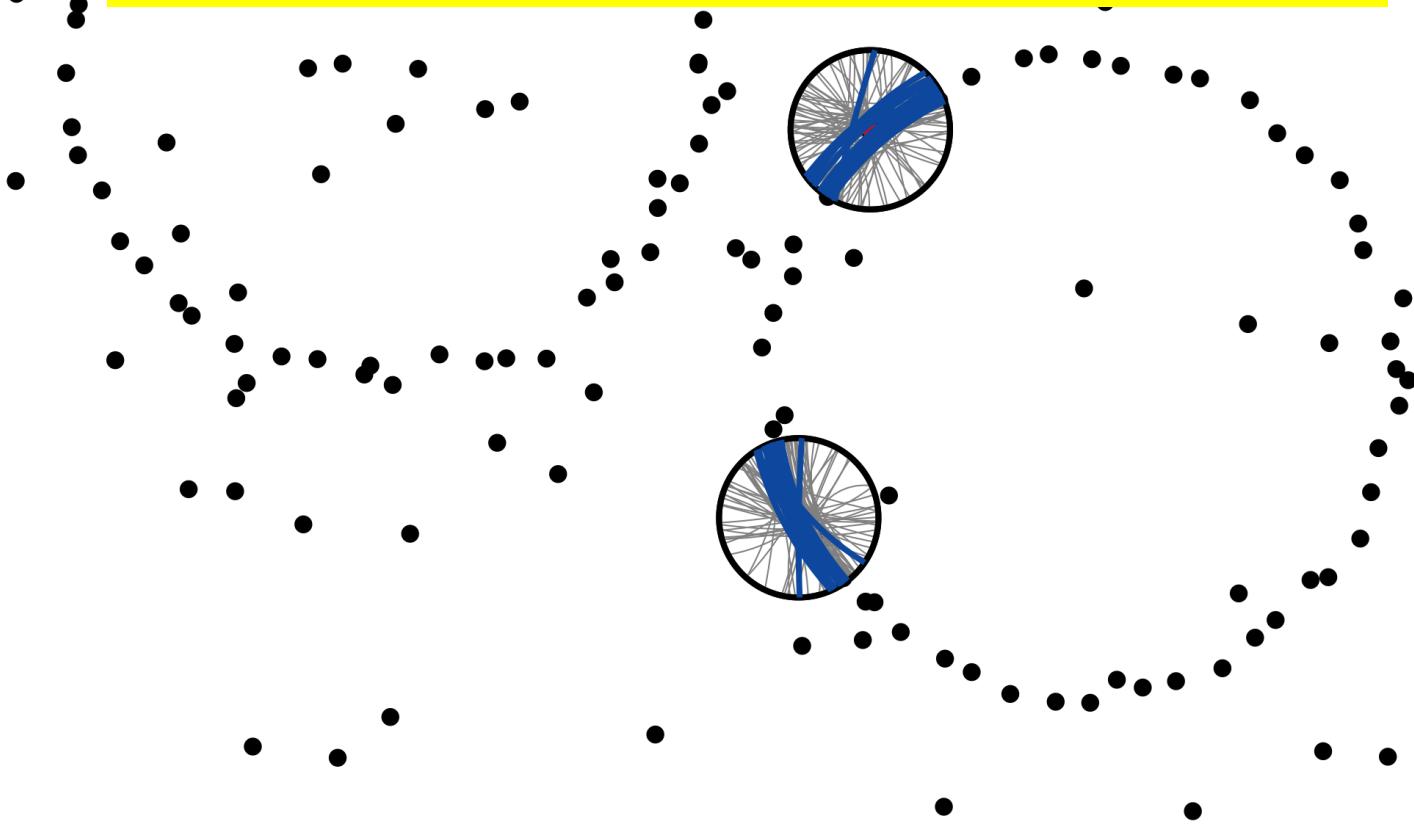
# Identify Structures by Clustering Preferences

The Jaccard distance can be used to measure distance between preference sets.

$$d_J(A, B) = 1 - \frac{|A \cap B|}{|A \cup B|}$$

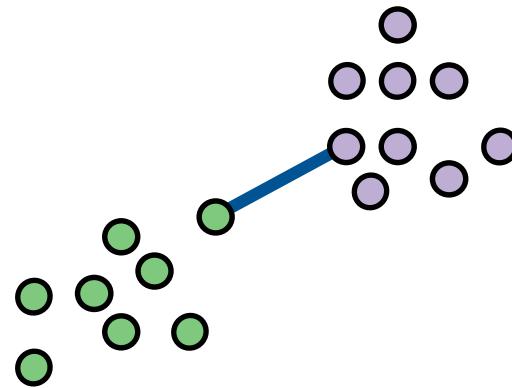


Structures are identified by clustering points, or better preference sets. Each cluster will then correspond to a structure.



# Structure Identification by Clustering in PS

Hierarchical clustering can be used in the Preference Space to recover the structures.



Single linkage

Average linkage

Complete linkage

Centroid linkage

Distances are measured in the Preference Space, and each element of clustering is identified by a preference set.

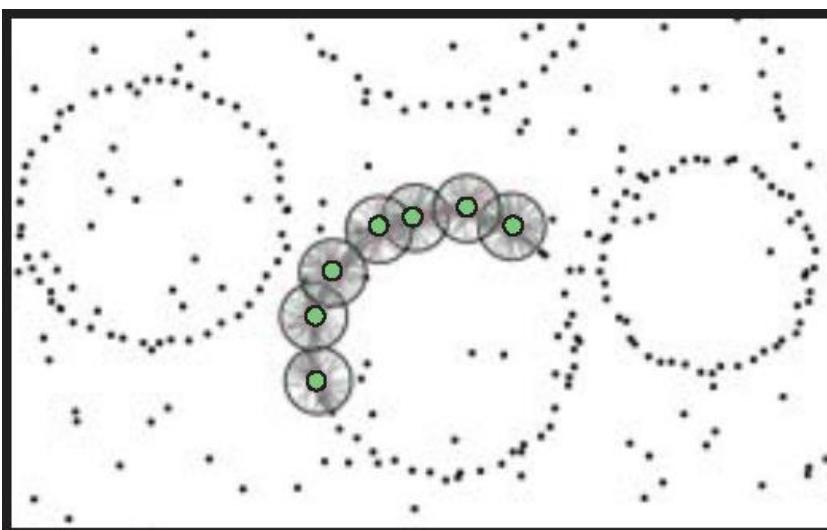
Instead of using centroids, we derive a conceptual representation for each cluster (which also lives in the preference space).

# J-linkage clustering [Toldo and Fusiello, ECCV 08]

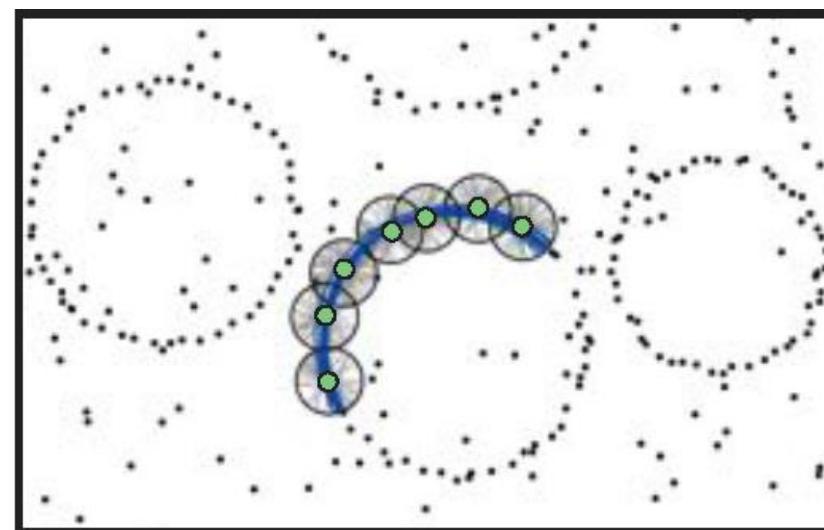
The representation of a cluster  $U \subset X$  in the preference space is the intersection of the PS of its points

$$U \subseteq X, PS(U) = \bigcap_{x \in U} PS(x)$$

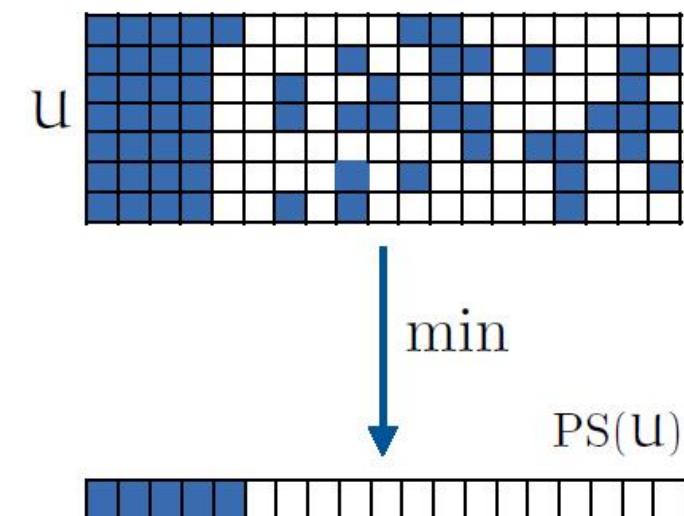
This is the component-wise min of rows in the preference matrix.



$$U \subseteq X$$



$$PS(U)$$



# J-linkage clustering [Toldo and Fusiello, ECCV 08]

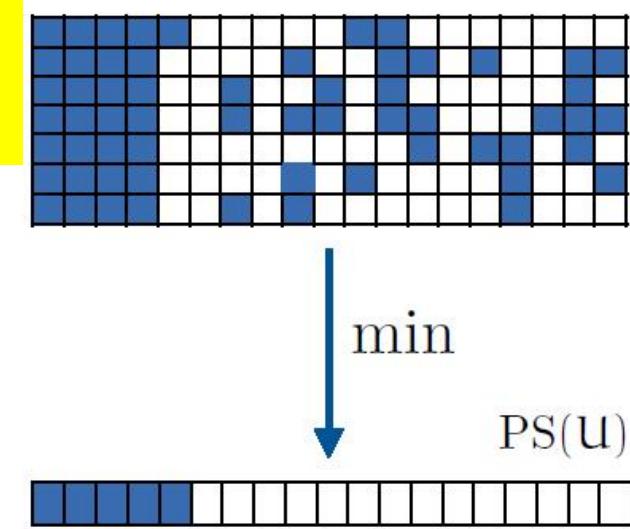
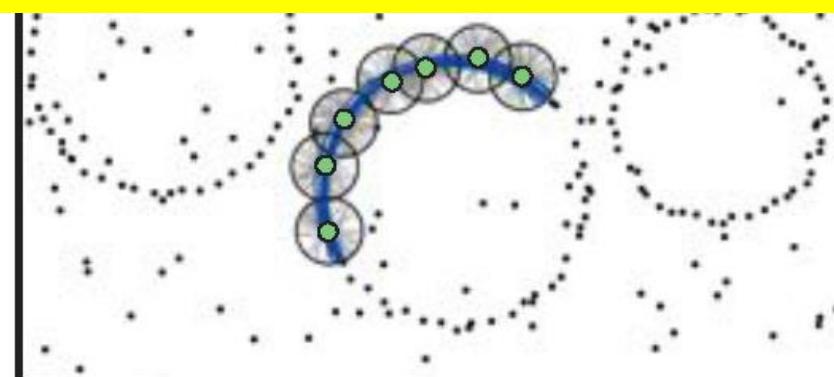
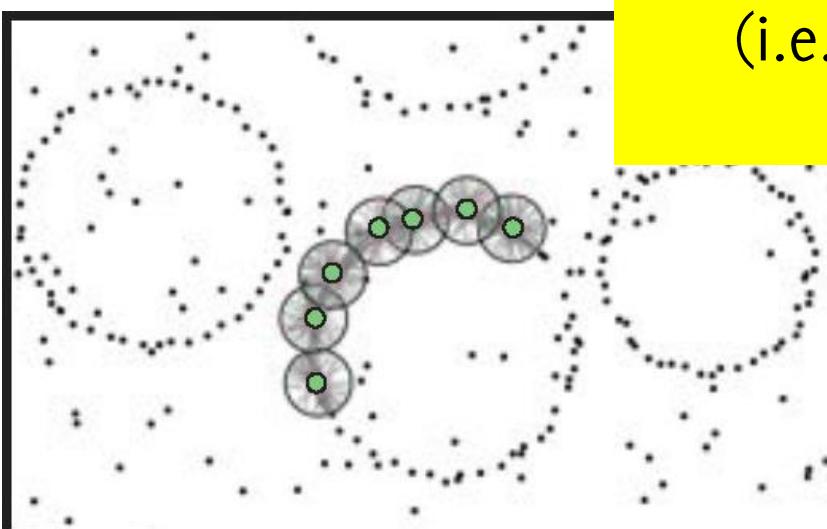
The representation of a cluster  $U \subset X$  in the preference space is the intersection of the PS of its points

$$U \subset X \quad PS(U) = \bigcap PS(x)$$

This is the compone

J-linkage iterates a hierarchical clustering scheme based on this distance until all the representatives of the clusters are disjoint (i.e. Jaccard distance = 1, no models in common)

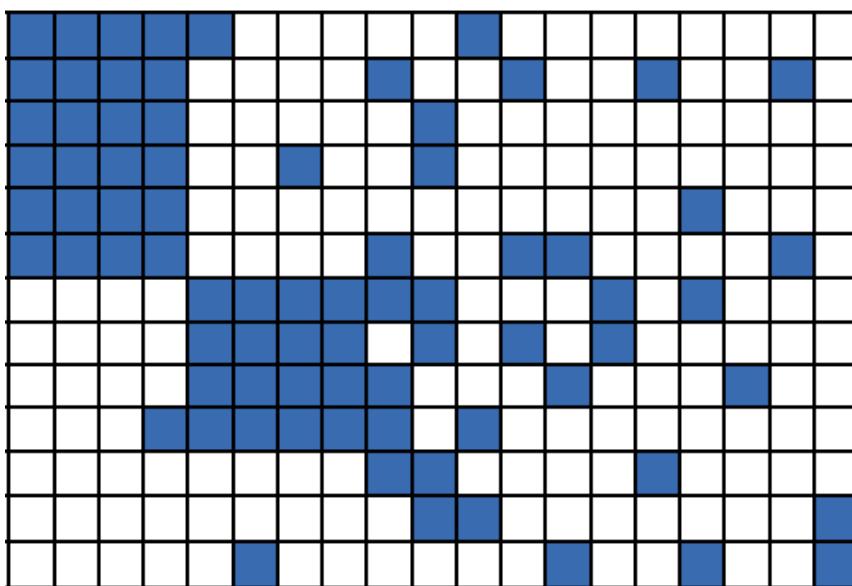
natrix.



# J-linkage clustering [Toldo and Fusiello, ECCV 08]

• •

• •



Preference matrix

**Input:**  $X$  data,  $\epsilon$  inlier threshold

**Output:** Partition in structures and models

Randomly sample model hypotheses  $H \subset \Theta$ ;  
Compute PS;

Put each point in its own cluster  $C_i = \{x_i\}$ ;

Compute  $d_J$  Jaccard distance between PS;

**while**  $\min(d_J) < 1$  **do**

    Find pair  $(C_i, C_j)$  of clusters with the  $\min d_J$ ;

    Replace the clusters with their union;

    Compute the PS of  $C_i \cup C_j$ ;

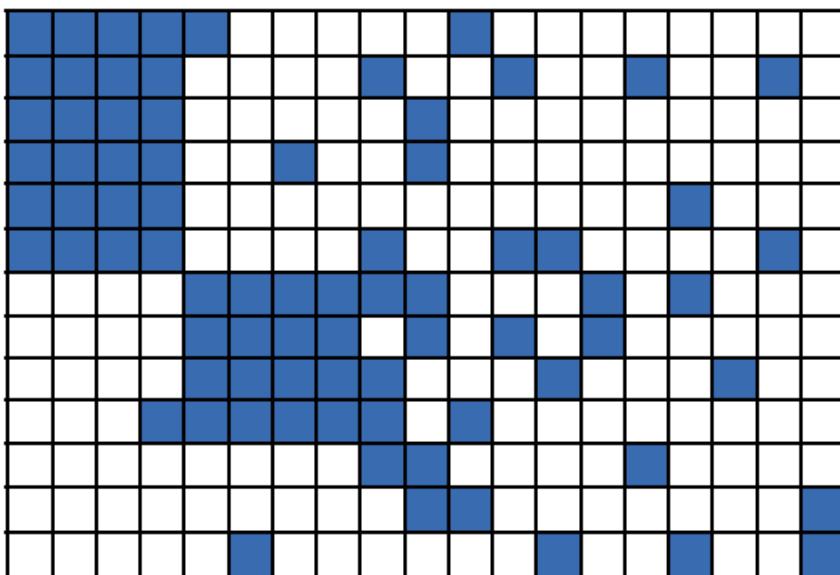
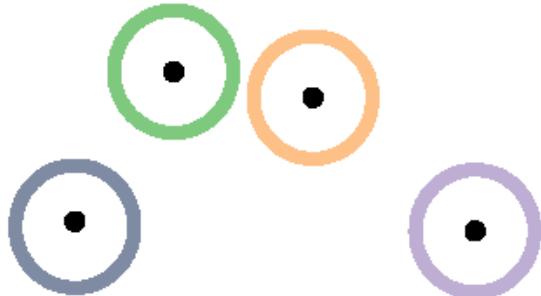
    Update  $d_J$ ;

**end**

Local fit of models to clusters;

# J-linkage clustering [Toldo and Fusiello, ECCV 08]

Put each point in its own cluster



Preference matrix

**Input:**  $X$  data,  $\epsilon$  inlier threshold

**Output:** Partition in structures and models

Randomly sample model hypotheses  $H \subset \Theta$ ;  
Compute PS;

Put each point in its own cluster  $C_i = \{x_i\}$ ;

Compute  $d_J$  Jaccard distance between PS;

**while**  $\min(d_J) < 1$  **do**

    Find pair  $(C_i, C_j)$  of clusters with the min  $d_J$ ;

    Replace the clusters with their union;

    Compute the PS of  $C_i \cup C_j$ ;

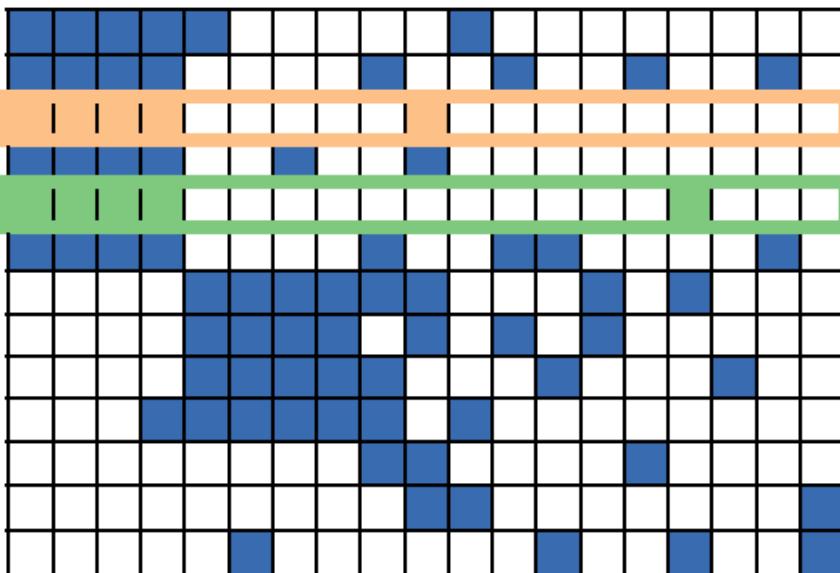
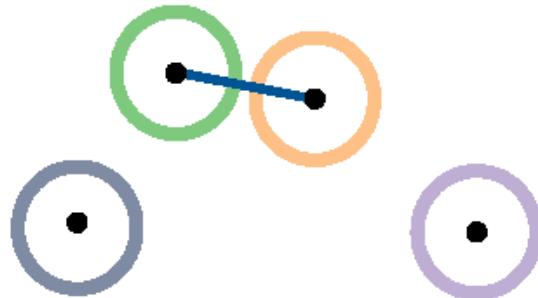
    Update  $d_J$ ;

**end**

Local fit of models to clusters;

# J-linkage clustering [Toldo and Fusiello, ECCV 08]

Find the closest points in Preference Space



Preference matrix

**Input:**  $X$  data,  $\epsilon$  inlier threshold

**Output:** Partition in structures and models

Randomly sample model hypotheses  $H \subset \Theta$ ;  
Compute PS;

Put each point in its own cluster  $C_i = \{x_i\}$ ;

Compute  $d_J$  Jaccard distance between PS;

**while**  $\min(d_J) < 1$  **do**

    Find pair  $(C_i, C_j)$  of clusters with the min  $d_J$

    Replace the clusters with their union;

    Compute the PS of  $C_i \cup C_j$ ;

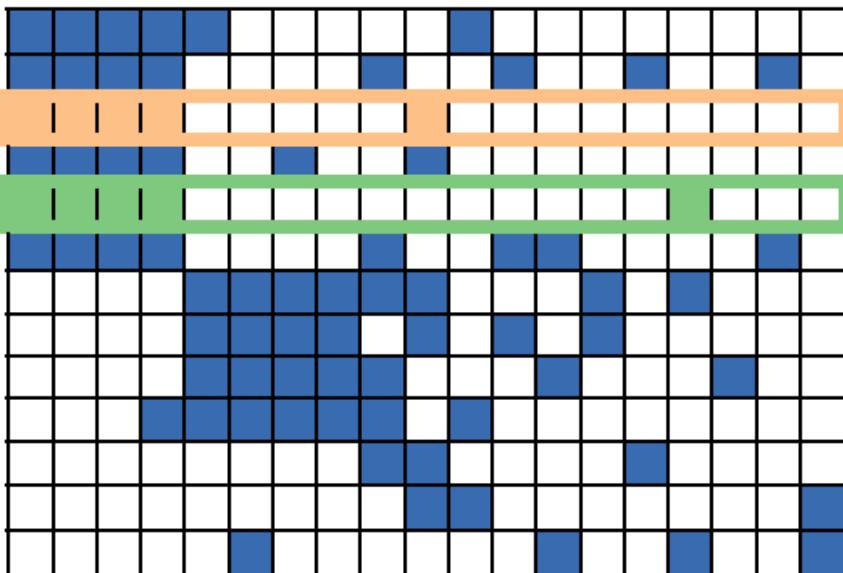
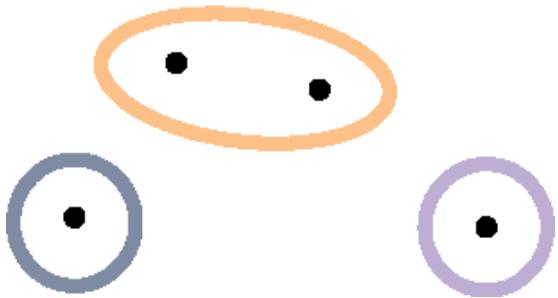
    Update  $d_J$ ;

**end**

Local fit of models to clusters;

# J-linkage clustering [Toldo and Fusiello, ECCV 08]

Merge cluster



Preference matrix

**Input:**  $X$  data,  $\epsilon$  inlier threshold

**Output:** Partition in structures and models

Randomly sample model hypotheses  $H \subset \Theta$ ;  
Compute PS;

Put each point in its own cluster  $C_i = \{x_i\}$ ;

Compute  $d_J$  Jaccard distance between PS;

**while**  $\min(d_J) < 1$  **do**

    Find pair  $(C_i, C_j)$  of clusters with the min  $d_J$ ;

**Replace the clusters with their union;**

    Compute the PS of  $C_i \cup C_j$ ;

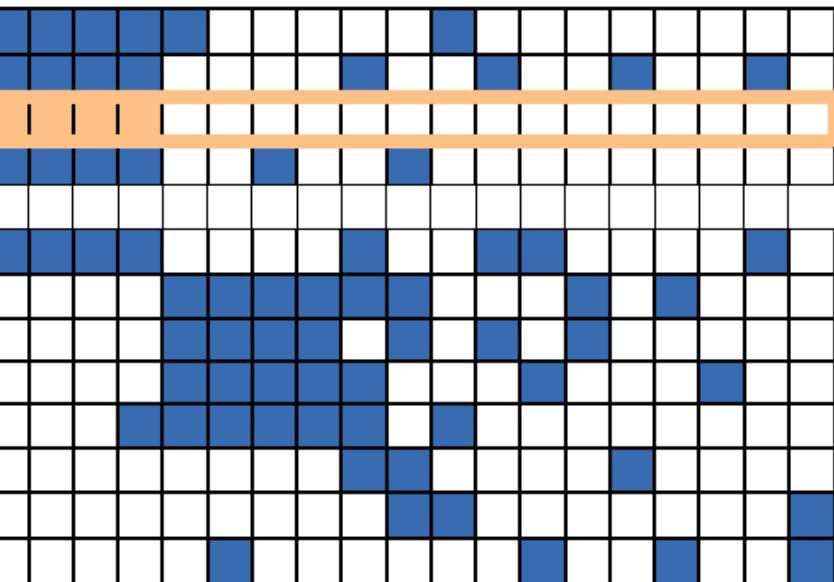
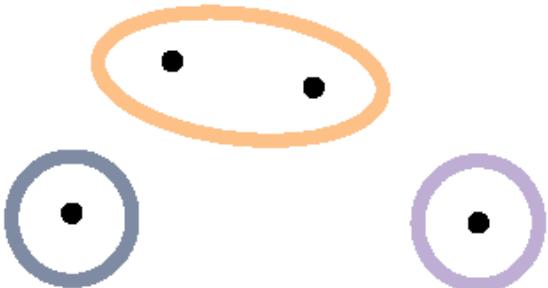
    Update  $d_J$ ;

**end**

Local fit of models to clusters;

# I-linkage clustering [Toldo and Fusiello, ECCV 08]

Update preferences



Preference matrix

**Input:**  $X$  data,  $\epsilon$  inlier threshold

**Output:** Partition in structures and models

Randomly sample model hypotheses  $H \subset \Theta$ ;  
Compute PS;

Put each point in its own cluster  $C_i = \{x_i\}$ ;

Compute  $d_J$  Jaccard distance between PS;

**while**  $\min(d_J) < 1$  **do**

    Find pair  $(C_i, C_j)$  of clusters with the  $\min d_J$ ;

    Replace the clusters with their union;

    Compute the PS of  $C_i \cup C_j$ ;

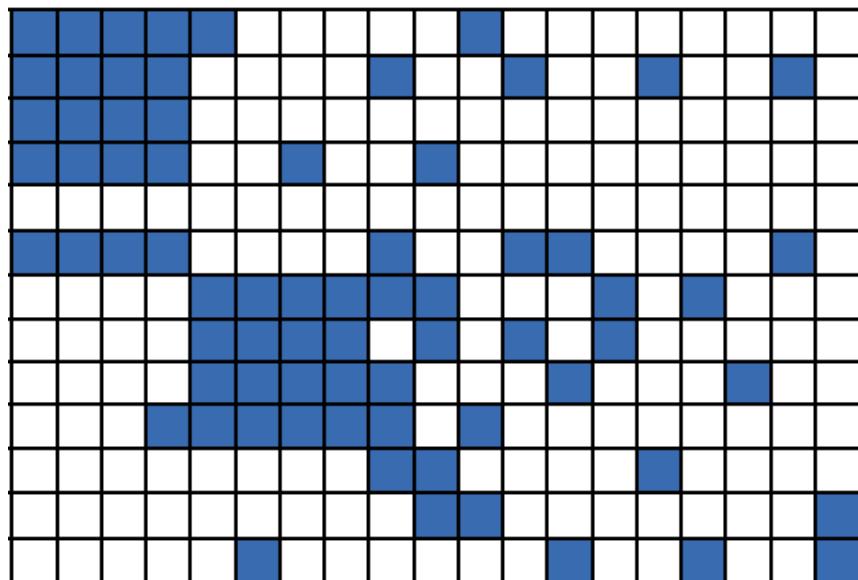
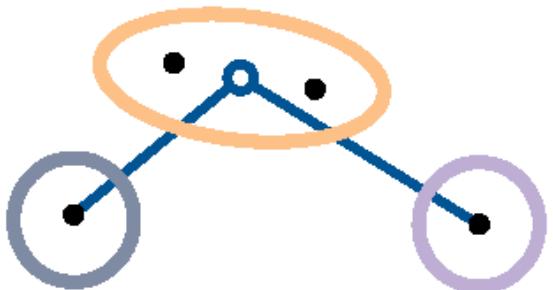
    Update  $d_J$ ;

**end**

Local fit of models to clusters;

# J-linkage clustering [Toldo and Fusiello, ECCV 08]

Update distances



Preference matrix

**Input:**  $X$  data,  $\epsilon$  inlier threshold

**Output:** Partition in structures and models

Randomly sample model hypotheses  $H \subset \Theta$ ;  
Compute PS;

Put each point in its own cluster  $C_i = \{x_i\}$ ;

Compute  $d_J$  Jaccard distance between PS;

**while**  $\min(d_J) < 1$  **do**

    Find pair  $(C_i, C_j)$  of clusters with the  $\min d_J$ ;

    Replace the clusters with their union;

    Compute the PS of  $C_i \cup C_j$ ;

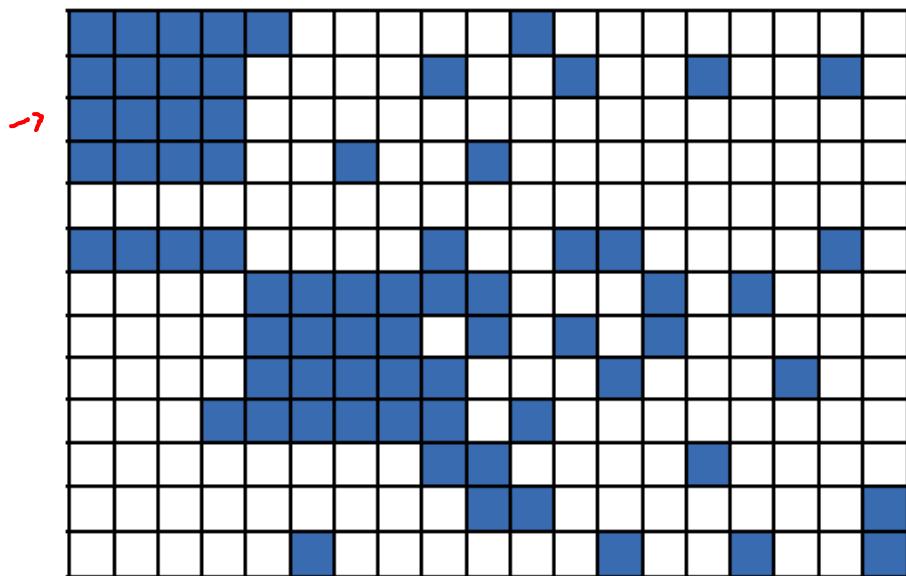
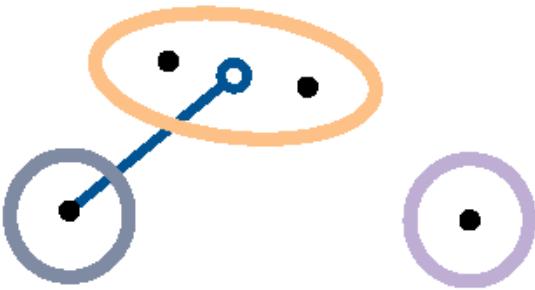
    Update  $d_J$ ;

**end**

Local fit of models to clusters;

# J-linkage clustering [Toldo and Fusiello, ECCV 08]

Continue until all PS are disjoint...



Preference matrix

**Input:**  $X$  data,  $\epsilon$  inlier threshold

**Output:** Partition in structures and models

Randomly sample model hypotheses  $H \subset \Theta$ ;  
Compute PS;

Put each point in its own cluster  $C_i = \{x_i\}$ ;

Compute  $d_J$  Jaccard distance between PS;

**while**  $\min(d_J) < 1$  **do**

    Find pair  $(C_i, C_j)$  of clusters with the min  $d_J$ ;

    Replace the clusters with their union;

    Compute the PS of  $C_i \cup C_j$ ;

    Update  $d_J$ ;

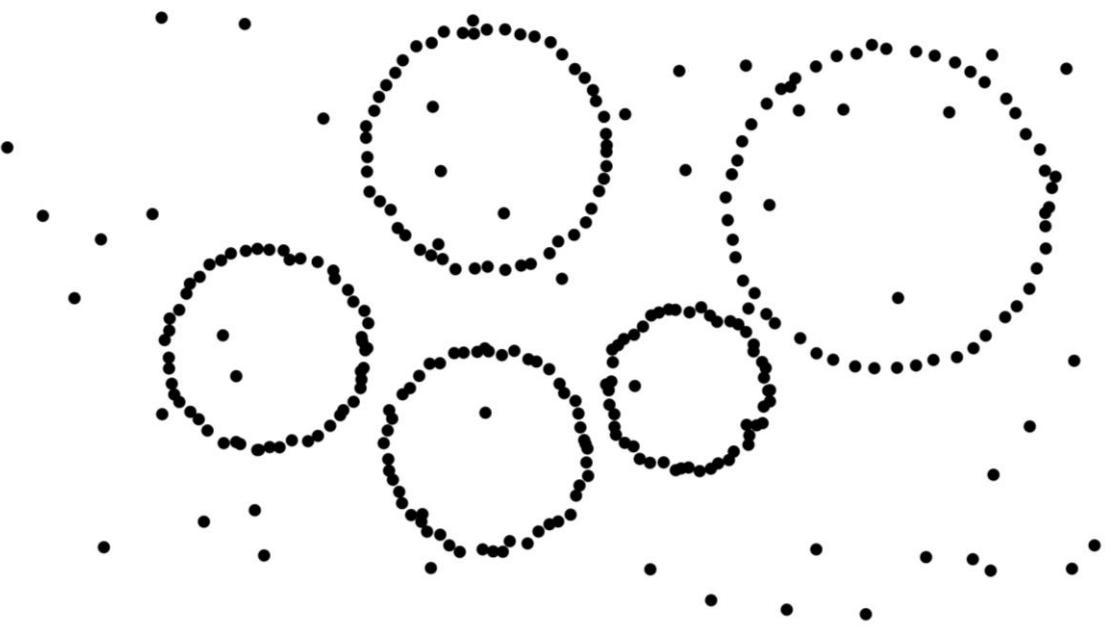
**end**

Local fit of models to clusters;

# J-linkage clustering [Toldo and Fusiello, ECCV 08]

**Input:**  $X$  data,  $\epsilon$  inlier threshold

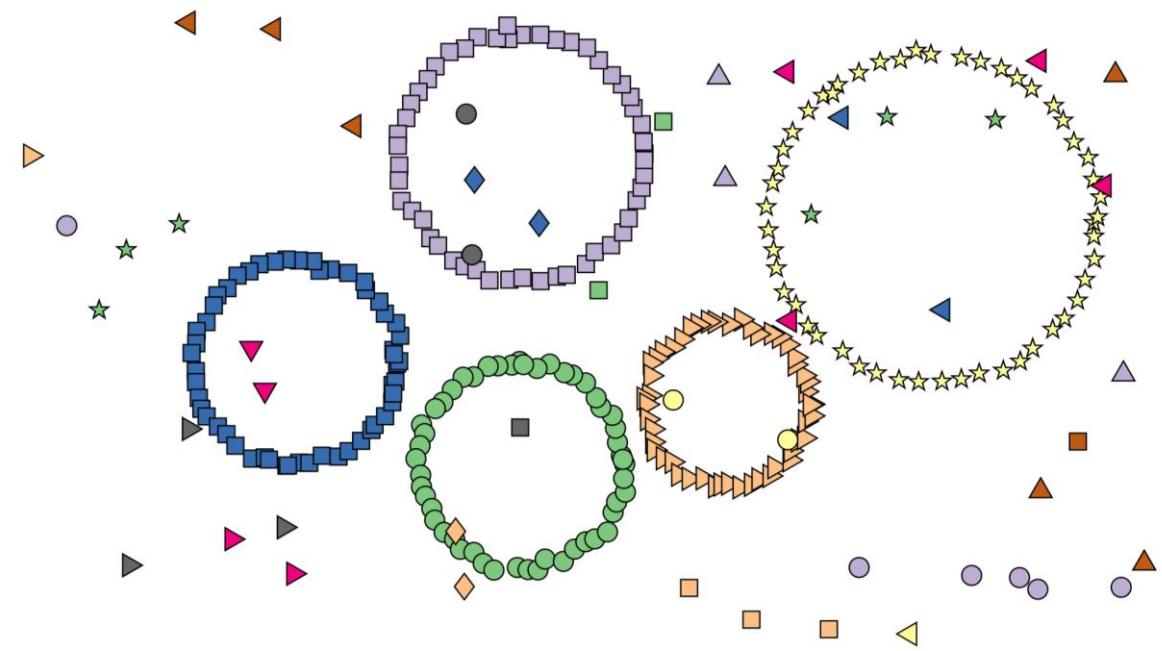
**Output:** Partition in structures and models



- Randomly sample model hypotheses  $H \subset \Theta$ ;  
Compute PS;
- Put each point in its own cluster  $C_i = \{x_i\}$ ;  
Compute  $d_J$  Jaccard distance between PS;  
**while**  $\min(d_J) < 1$  **do**
  - Find pair  $(C_i, C_j)$  of clusters with the  $\min d_J$ ;
  - Replace the clusters with their union;
  - Compute the PS of  $C_i \cup C_j$ ;
  - Update  $d_J$ ;**end**  
Local fit of models to clusters;

# J-linkage clustering [Toldo and Fusiello, ECCV 08]

**Input:**  $X$  data,  $\epsilon$  inlier threshold  
**Output:** Partition in structures and models



```
Randomly sample model hypotheses  $H \subset \Theta$ ;  
Compute PS;  
Put each point in its own cluster  $C_i = \{x_i\}$ ;  
Compute  $d_J$  Jaccard distance between PS;  
while  $\min(d_J) < 1$  do  
    Find pair  $(C_i, C_j)$  of clusters with the min  $d_J$ ;  
    Replace the clusters with their union;  
    Compute the PS of  $C_i \cup C_j$ ;  
    Update  $d_J$ ;  
end  
Local fit of models to clusters;
```

# J-linkage clustering [Toldo and Fusiello, ECCV 08]

## Pro:

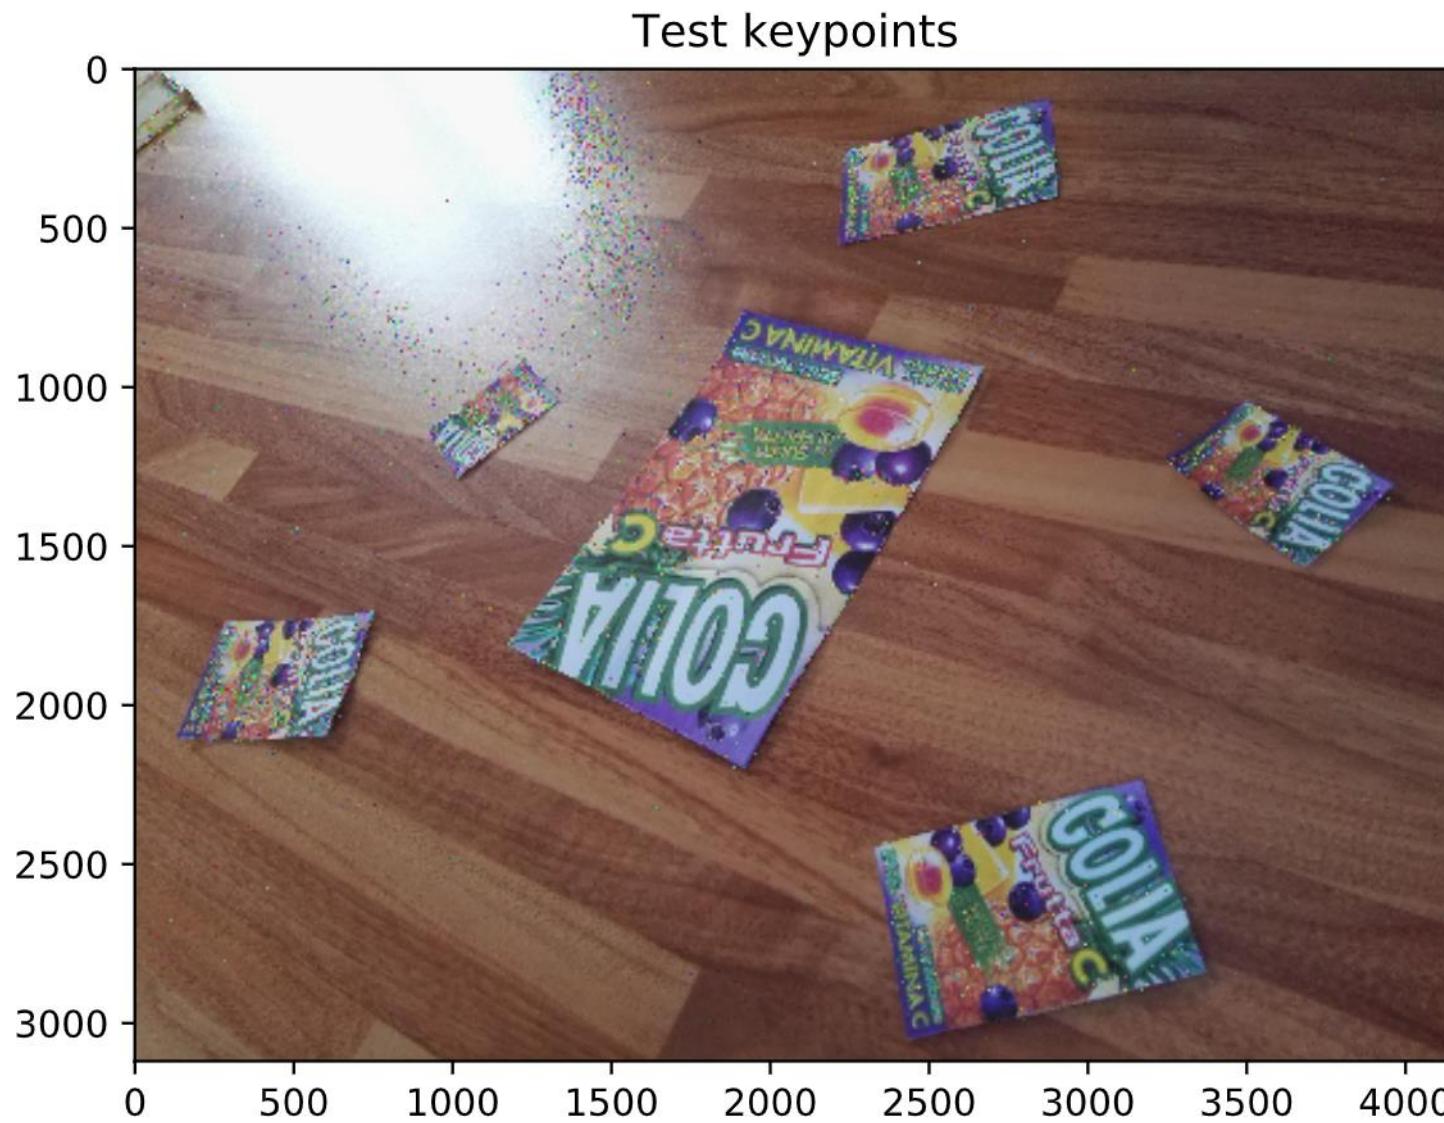
- The number of structures is automatically determined
- For each **cluster** there exists at least one **model** that fits all the points of the cluster
- **Clusters are “maximal”** in the sense that does not exist a model that explain all the points of two distinct clusters

## Cons:

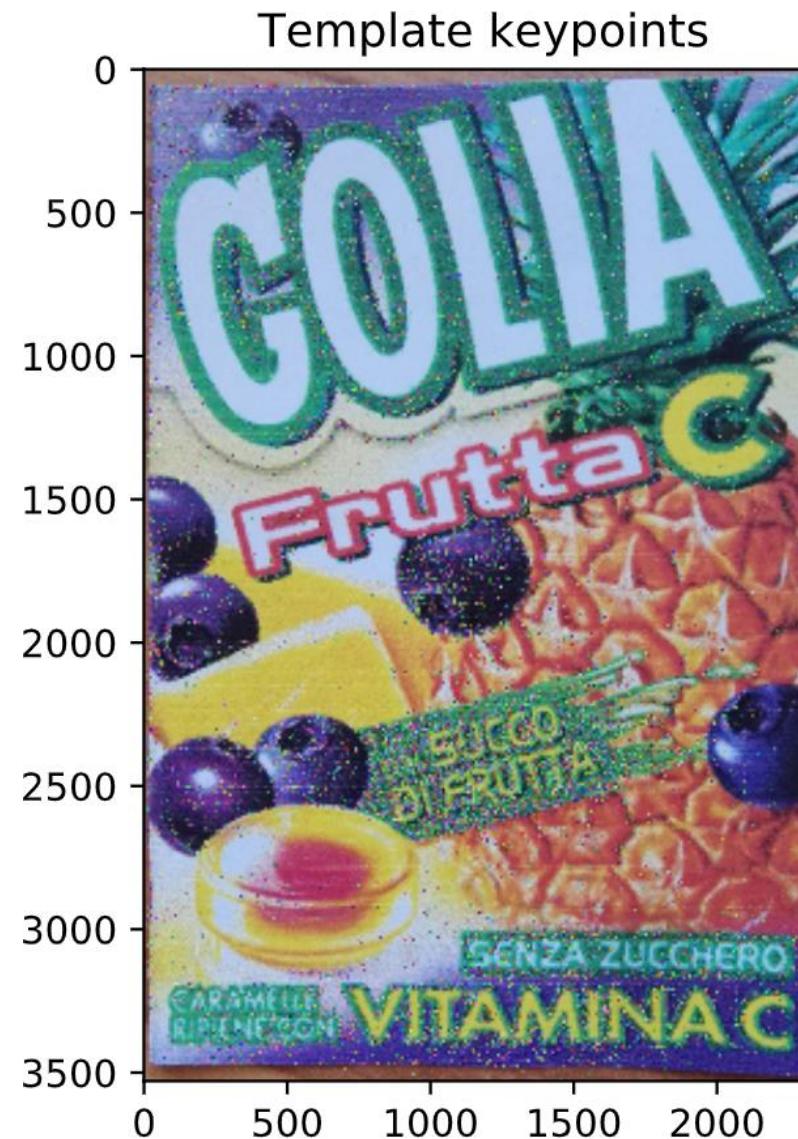
- It still relies on a pre-defined inlier threshold  $\epsilon$
- **Outliers have to be filtered out a posteriori** (they will always end up in small clusters)

# Object Detection by Computer Vision Features

# Object Detection: Keypoint Extraction

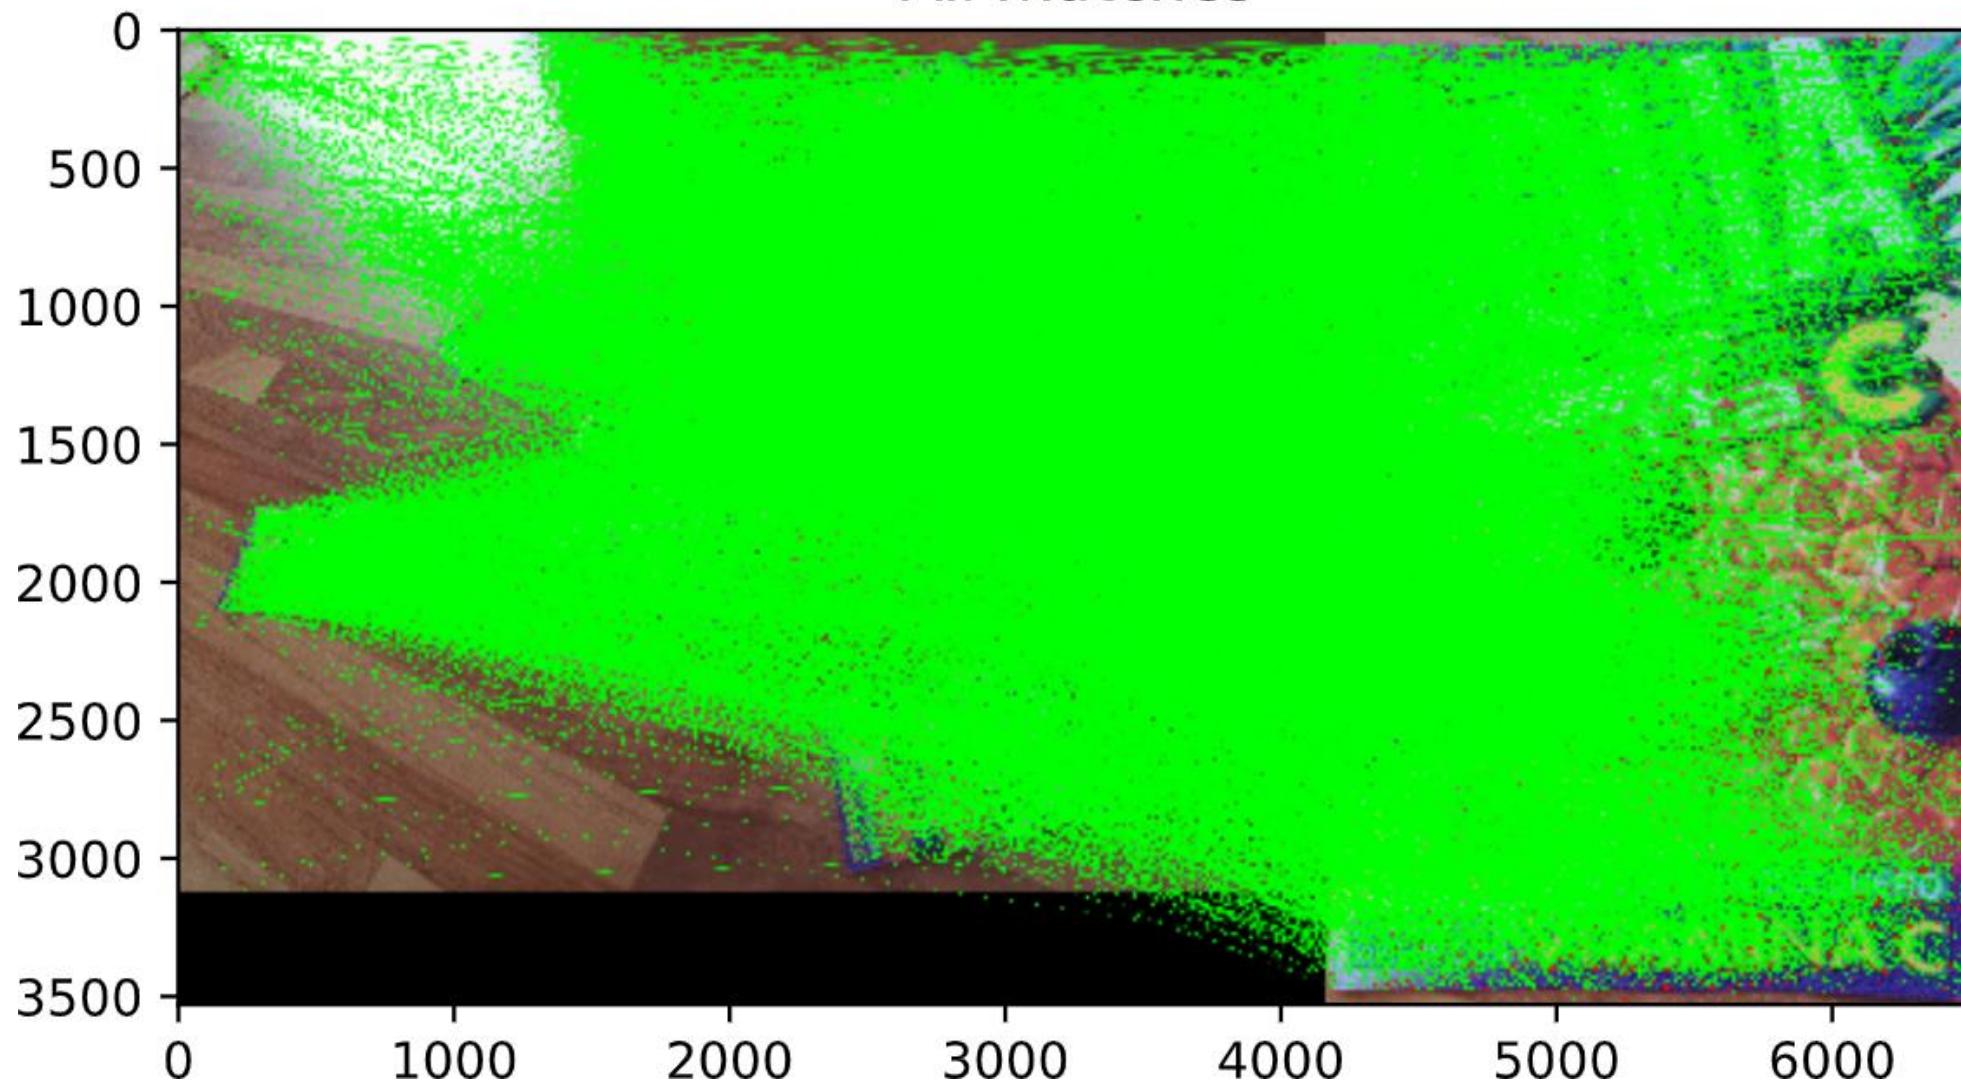


# Object Detection: Keypoint Extraction



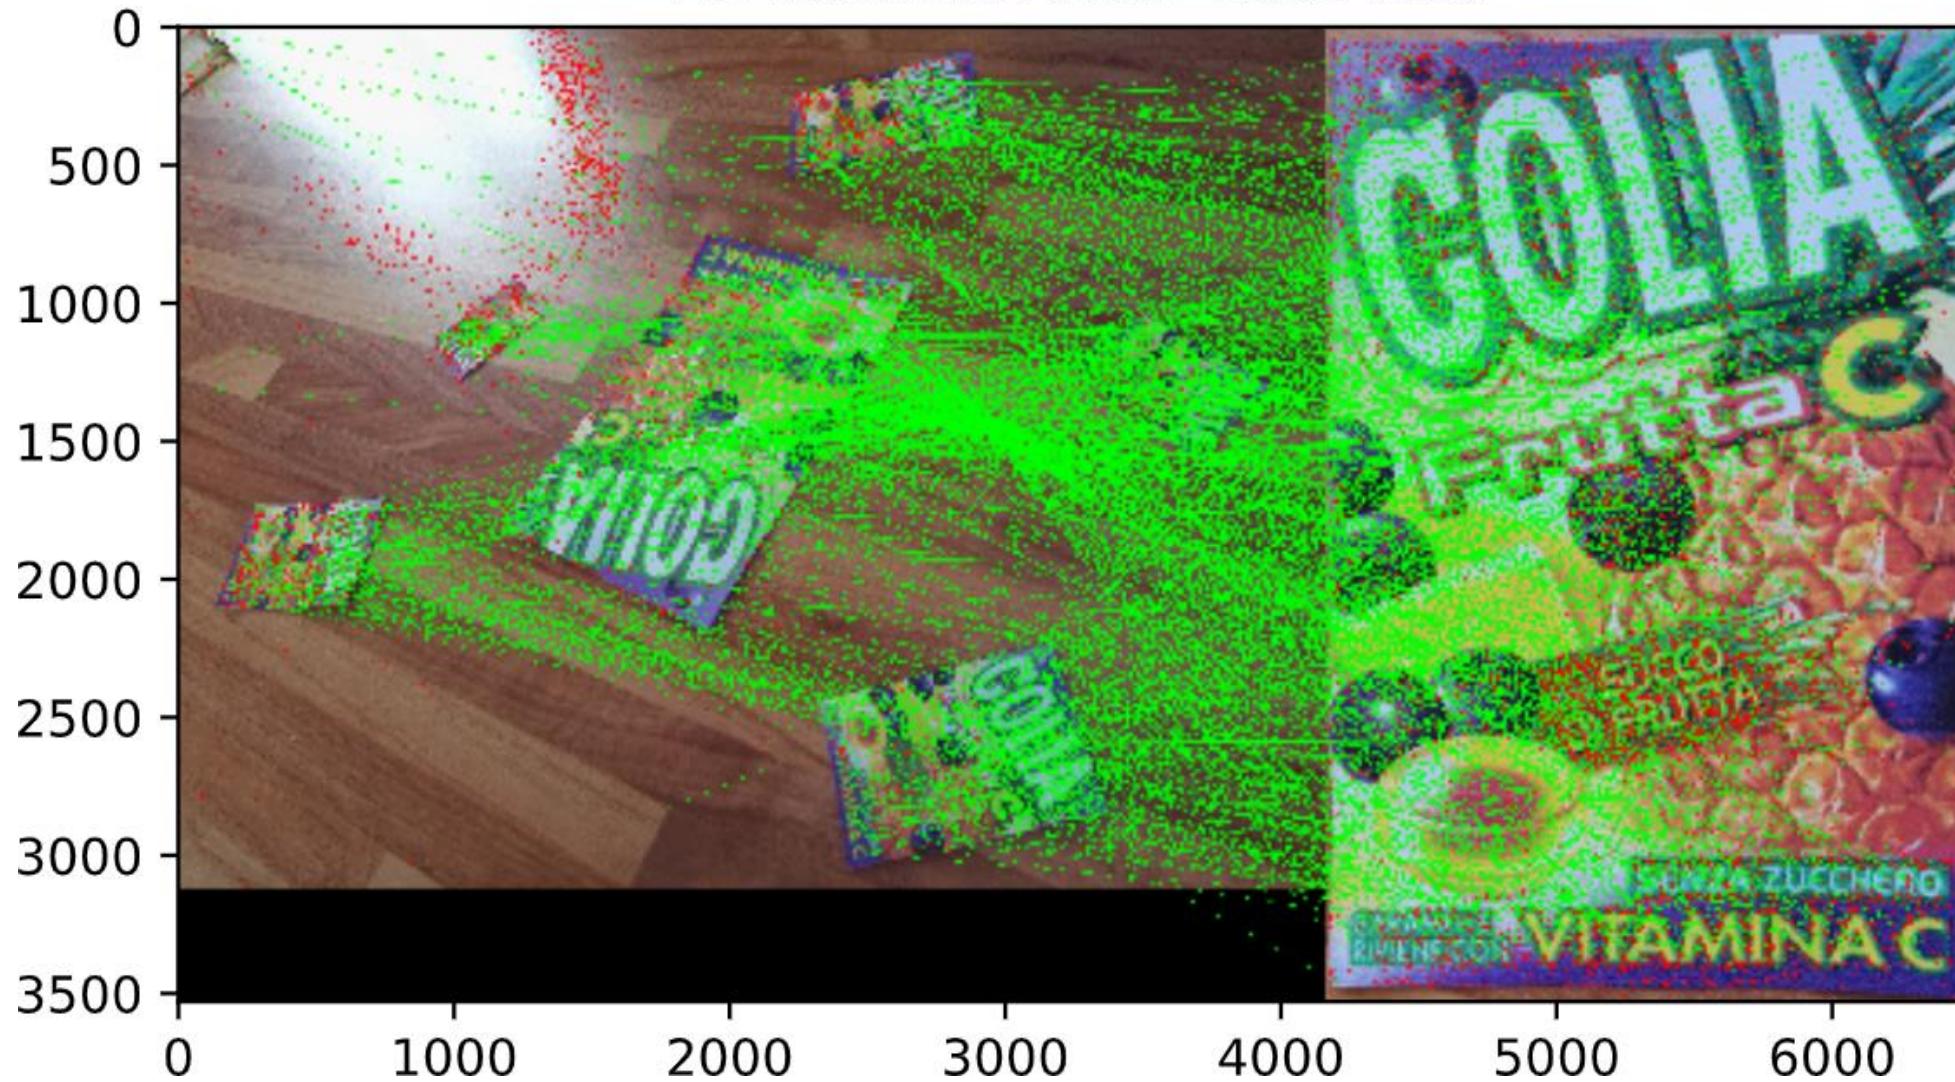
# Keypoint Matching

All matches



# Keypoint Matching: Ratio Test

All matches after ratio test



# Sequential Ransac Iterations



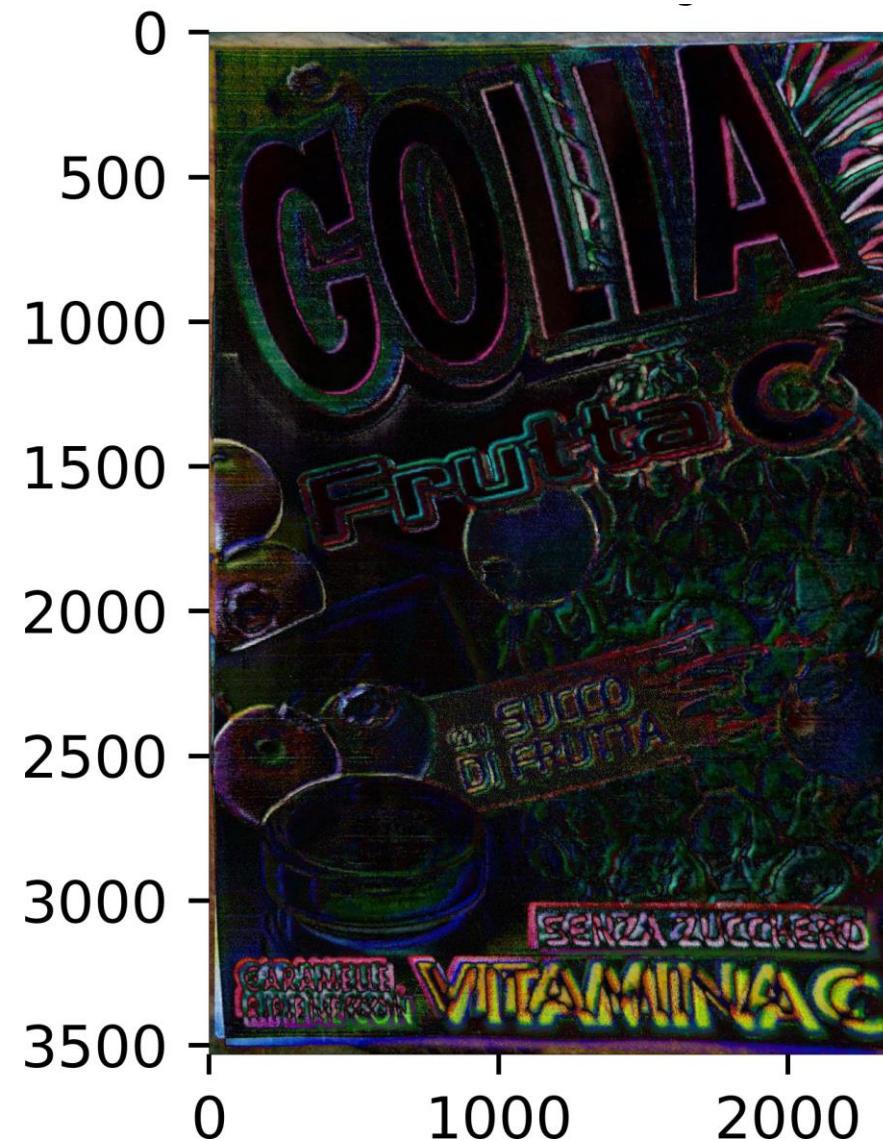
# Sequential Ransac Image Rectification



# Sequential Ransac: Histogram Matching



# Sequential Ransac: Image Differences



# Sequential Ransac

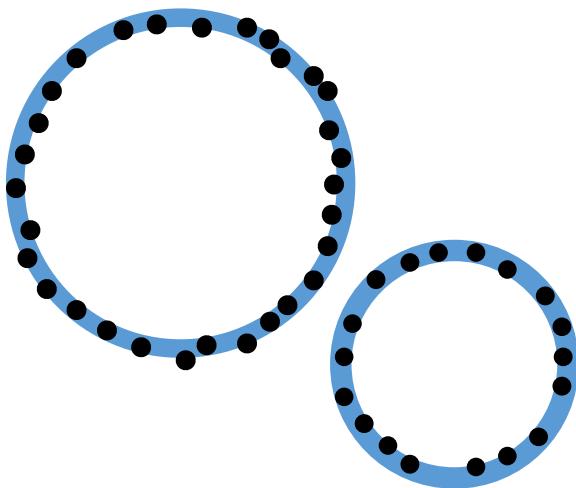


# Tempalte matching limitation (it is not cherry creme)



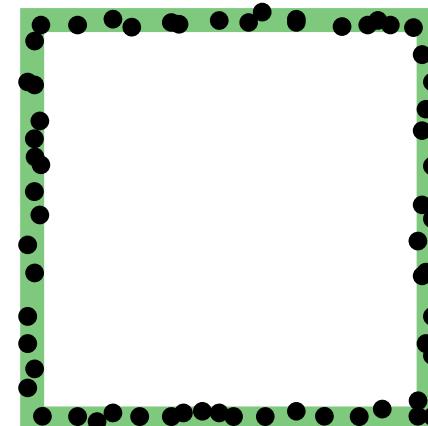
# Towards multi-class model fitting

By using T-linkage we are able to detect separately



Multiple circles

or

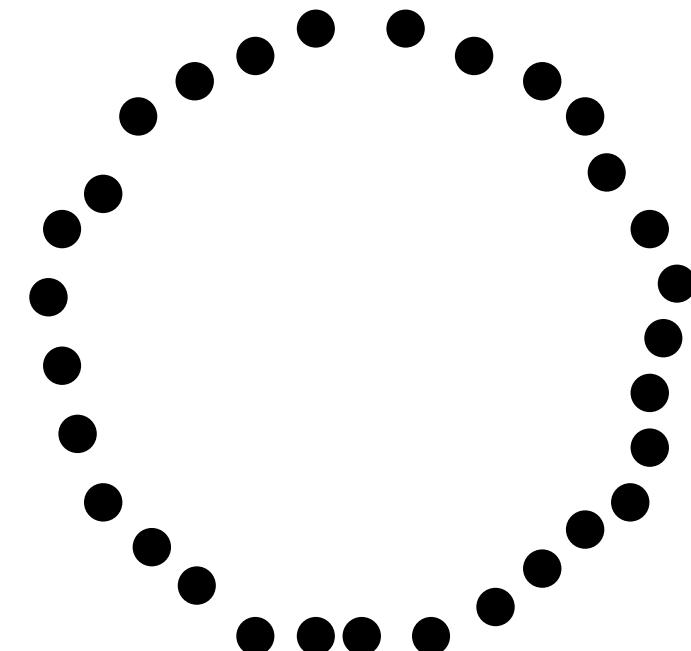


Multiple lines

# Multi-class model fitting

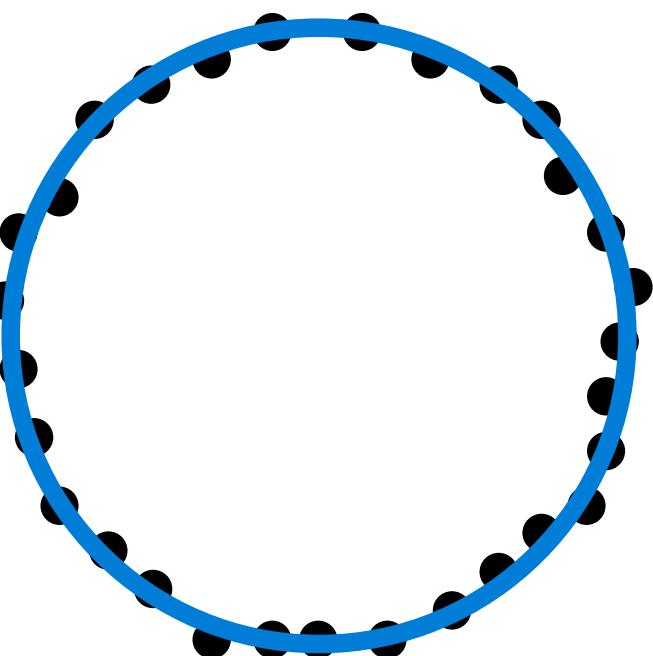
If we want to describes data with multiple classes of models (e.g., lines and circles), at some points we must face a model selection problem.

circle or a polygon?



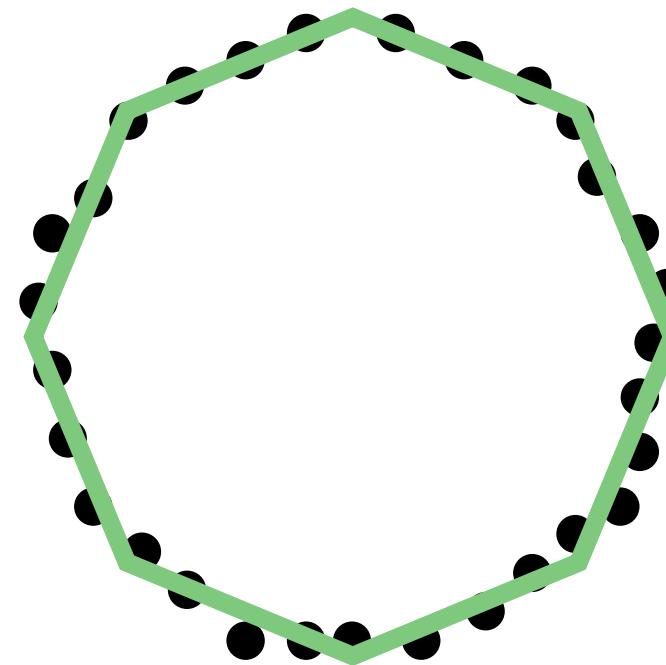
# Multi-class model fitting

Different interpretations of the data are possible



1 complex instance  
(circle/cylinder/fundamental matrix)

vs  
many simpler instances  
(lines/planes/homographies)



# Geometric Robust Information Criterion (GRIC)

**GRIC cost:**

$$-\frac{1}{2} \underbrace{\sum \left( \frac{r_i}{\sigma} \right)^2}_{\$} + \underbrace{\lambda_1 dn + \lambda_2 k}_{\text{€}}$$



Geometric  
residuals

Model  
complexity

$r_i$  residuals  
 $\sigma$  noise estimate

$d$  dim. of model manifold  
 $n$  # points  
 $k$  # model parameters

AIC	$-2L + 2p$
BIC	$-2L + p \log(rn)$
MDL	$-2L + p/2 \log_2(rn)$
GBIC	$-2L + \log(r)dn + \log(rn)k$
GRIC	$-2L + \lambda_1 dn + \lambda_2 k$

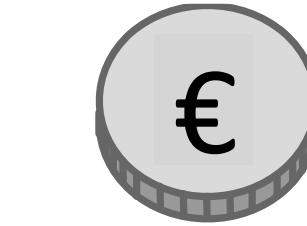


P. H.S. Torr. *Model Selection for Two View Geometry: A Review*. Springer Berlin Heidelberg, Berlin, Heidelberg, 1999

# Geometric Robust Information Criterion (GRIC)

**GRIC cost:**

$$-\frac{1}{2} \underbrace{\sum \left( \frac{r_i}{\sigma} \right)^2}_{\$} + \underbrace{\lambda_1 dn + \lambda_2 k}_{\text{€}}$$



**Geometric residual and model complexity  
are expressed in different “currency”**

AIC	$-2L + 2p$
BIC	$-2L + p \log(rn)$
MDL	$-2L + p/2 \log_2(rn)$
GBIC	$-2L + \log(r)dn + \log(rn)k$
GRIC	$-2L + \lambda_1 dn + \lambda_2 k$

$r_i$  residuals  
 $\sigma$  noise estimate

$d$  dim. of model manifold  
 $n$  # points  
 $k$  # model parameters



P. H.S. Torr. *Model Selection for Two View Geometry: A Review*. Springer Berlin Heidelberg, Berlin, Heidelberg, 1999

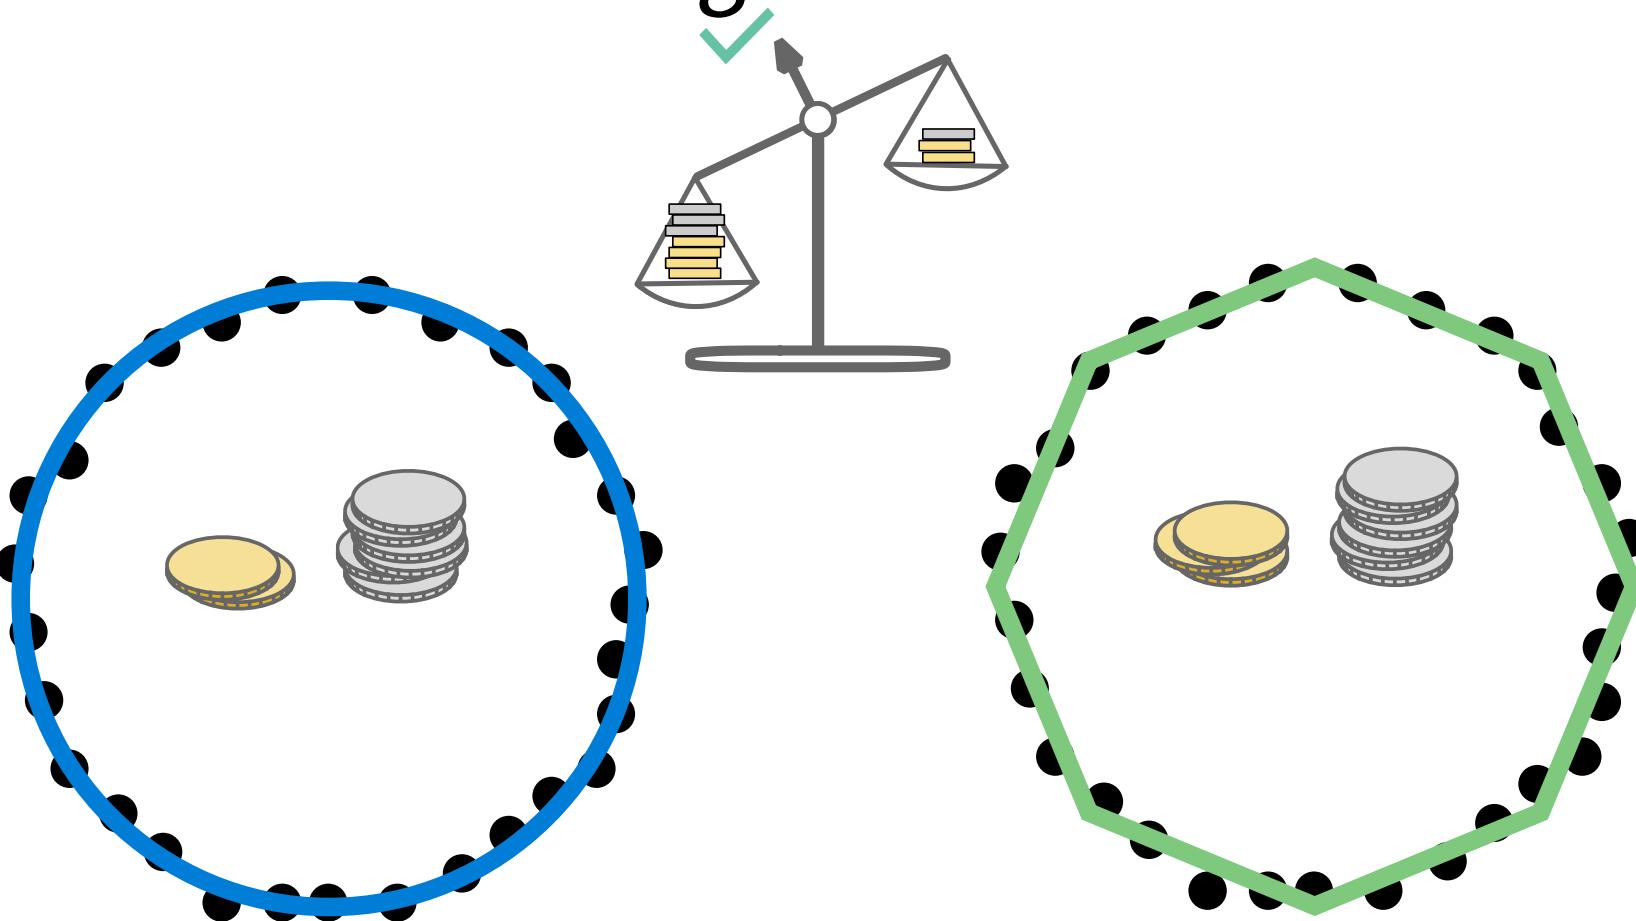
$$-\frac{1}{2} \sum \left( \frac{r_i}{\sigma} \right)^2 + \lambda_1 dn + \lambda_2 k$$



set the “exchange rate”

Models are in the eye of the beholder!

# Multi-class model fitting



1 complex instance

(circle/cylinder/fundamental matrix)

vs

many simpler instances

(lines/planes/homographies)

And if you want to investigate further in a  
research thesis...

# Thesis on Multi model fitting: beyond analytical models

*Classical multi-model fitting methods leverage on **analytical models** with **very few parameters**.*

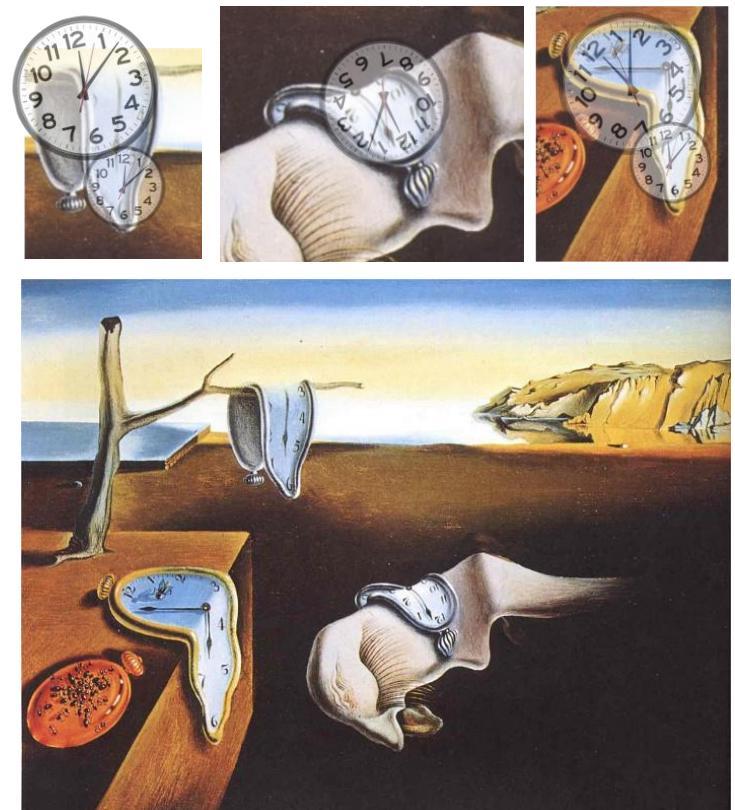
*These can be easily integrated in robust fitting framework to gain resilience against outliers, but, due to their simplicity, they are not able to encode higher level semantic information.*

## Research challenge:

- Go beyond analytical models: use more general models to foster the extraction of higher-level information from data.
- Non-analytical models can be:
  - **Parameter-less** (partition of data)
  - **Neural** models (exorbitant numbers of parameters)
- How non-analytical models can be fitted in a robust way (resampling)?

## Applications:

- Ensamble clustering
- Instance Segmentation Networks



*It is possible to use Preference Analysis to deal with **non-rigid** models ?*

*Luca Magri, Giacomo Boracchi*

# Thesis on Learning for 3D vision

*Combine the metrological accuracy of traditional model-fitting methods with the high-level understanding capabilities of deep learning models.:*

## Research challenges:

- 3D scene are complex, this make the training unstable
- Robustness across viewpoint changes and illumination condition is required

## Aim:

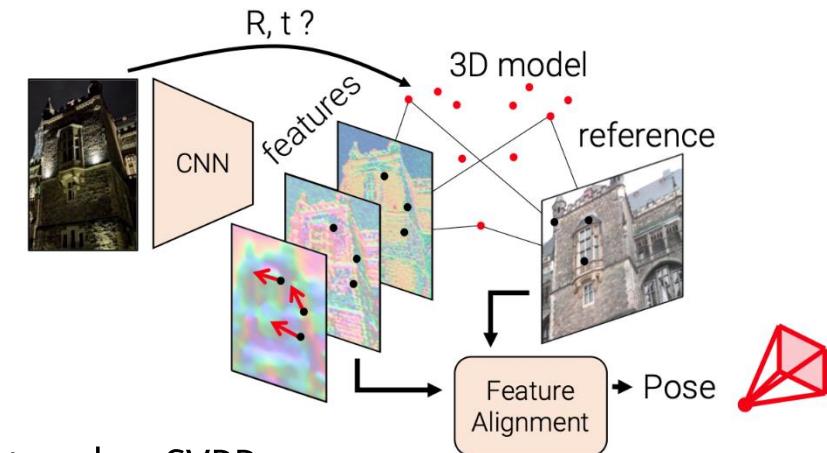
- Conceive new algorithm able to exploit deep networks to learn robust and invariant visual features, and leverage on principled algorithm to geometric estimation.

## Application:

- 3D reconstruction (Structure from Motion)
- Deep trifocal tensor estimation

## References to start with:

- [1] Sarlin et Al, SuperGlue: Learning Feature Matching with Graph Neural Networks. CVPR 20
- [2] Sarlin et Al, Back to the Feature: Learning Robust Camera Localization from Pixels to Pose. CVPR 21
- [3] Lindenberger et Al, Pixel-Perfect Structure-from-Motion with Featuremetric Refinement. ICCV 21
- [4] Zhang et Al. Content Aware Unsupervised Deep Homography Estimation



# Thesis on Neural Sampling for 3D data

*Indoor modelling aims at recovering structural elements such as rooms, doors, and walls, then reconstructs a structured 3D model consisting of such elements. Instance Segmentation architectures can provide higher-level information to enable semantic-aware multi-model fitting*

## Research goals:

- Investigate how Instance Segmentation Networks can be integrated to steer the sampling towards promising model
- Enforce robustness through resampling scheme
- Test on 3D datasets for indoor classification

## Materials and Methods:

- 3D datasets
- Access to a server mounting GPUs will be provided

## Some reference to start:

- [1] Jiang, Li, et al. PointGroup: Dual-Set Point Grouping for 3D Instance Segmentation. In CVPR. 2020.
- [2] Yang, Bo, et al. Learning object bounding boxes for 3d instance segmentation on point clouds. In NIPS. 2019.
- [3] Yi, Li, et al. Gspn: Generative shape proposal network for 3d instance segmentation in point cloud. In CVPR. 2019.



<http://www.scan-net.org>

Or in a research internship...

drop me an email for further information

[giacomo.boracchi@polimi.it](mailto:giacomo.boracchi@polimi.it)

# Collaboration with Antares Vision Group



Antares Vision is currently sponsoring a PhD Scholarship under our supervision

We are planning to expand collaboration with an MSc internship in their research labs

Topics for a research thesis covers:

- Computer-Vision for industrial applications
- Deep Learning models for industrial applications

And in particular for visual inspection, recognition, tracking, anomaly detection,...

# Collaboration with Antares Vision Group



## Antares Vision Group

- 900+ employies worldwide
- 26% of workforce devoted to R&D activities
- Headquarter in Travagliato (Brescia)
- Includes Orobix, an AI service company

## LIFE SCIENCE

WE SERVE MOST OF THE LARGEST  
**LIFE SCIENCE** GLOBAL PLAYERS

10 OF TOP 20 GLOBAL PHARMACEUTICAL  
MANUFACTURERS RAPRESENT  
OUR LOYAL CUSTOMER BASE

## FOOD & BEVERAGE

WE SERVE MOST OF THE LEADING GLOBAL  
**FOOD & BEVERAGE MULTINATIONALS**

THE FIRST 4 CUSTOMERS STARTED THE  
RELATIONSHIP MORE THAN 10 YEARS AGO.

# Collaboration with Antares Vision Group



## INSPECTION

### Camera Based system

Visual Inspection for product and packaging

Smart Camera

Embedded Vision

SWIR camera

Hyperspectral NIR/VIS

### Laser Absorption Spectroscopy

Pressure/Vacuum measurement

Head Space Gas Analysis

Leak detection [CCIT]

Laboratory Instruments

### High Voltage

Leak detection [CCIT]

Laboratory Instruments

### Sensor Based

Vacuum/Pressure Decay

Leak Detection

X-Ray and HF

CheckWeigh

Metal Detector

Event Based

## TRACEABILITY

### Serialization

Aggregation

Mobile Traceability

Real-time IOT



## AUTHENTICATION

### Visible and Invisible Tag

Serial Number Management

Visual Inspection Systems



## DATA MANAGEMENT

### Business Intelligence

Data Analysis

Advanced Analytics

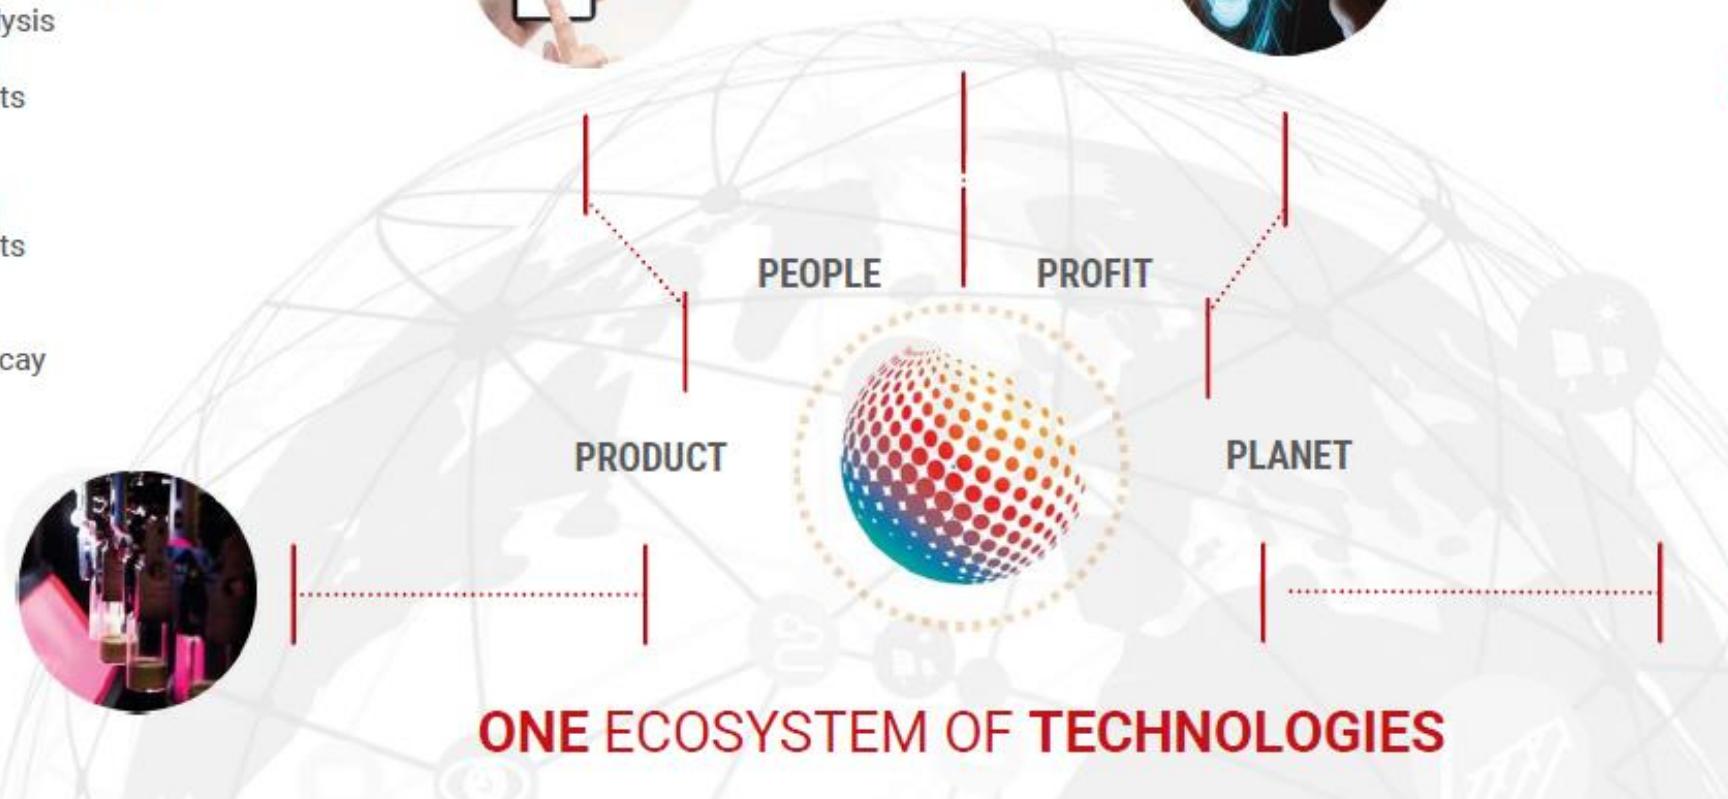
Artificial Intelligence

Blockchain

Big Data

Cloud Computing

Integrated Platforms



# Collaboration with Antares Vision Group



## AI solutions

**Ready-to-use AI solutions with a shorter and consolidated time-to value.**  
Flexible, but focused on solving well-identified problems.

### **detectiv.ai**

Real-time AI solution for anomaly detection

### **AI-go**

Vision inspection platform:  
classification, segmentation, OCR

## AI engineering

**Tailor-made AI models, solving a wider variety of problems.**  
Leverage a team of experts in cross-industry AI applications. Wider range of applications (e.g. Vision inspection and Deep learning, Time series, Bayesian models, Reinforcement learning).

## AI managed services\*

**We take care of all the AI of a company ensuring trust, compliance and reliability once in production.**

### **invariant.ai®**

Platform for the deployment, governance, monitoring and lifecycle management of AI systems in critical processes.

**invariant.ai®  
EDGE**

**invariant.ai®  
CLOUD**

\*Oròbix named in the **2021 Gartner “Market Guide for Artificial Intelligence Service Providers”**